

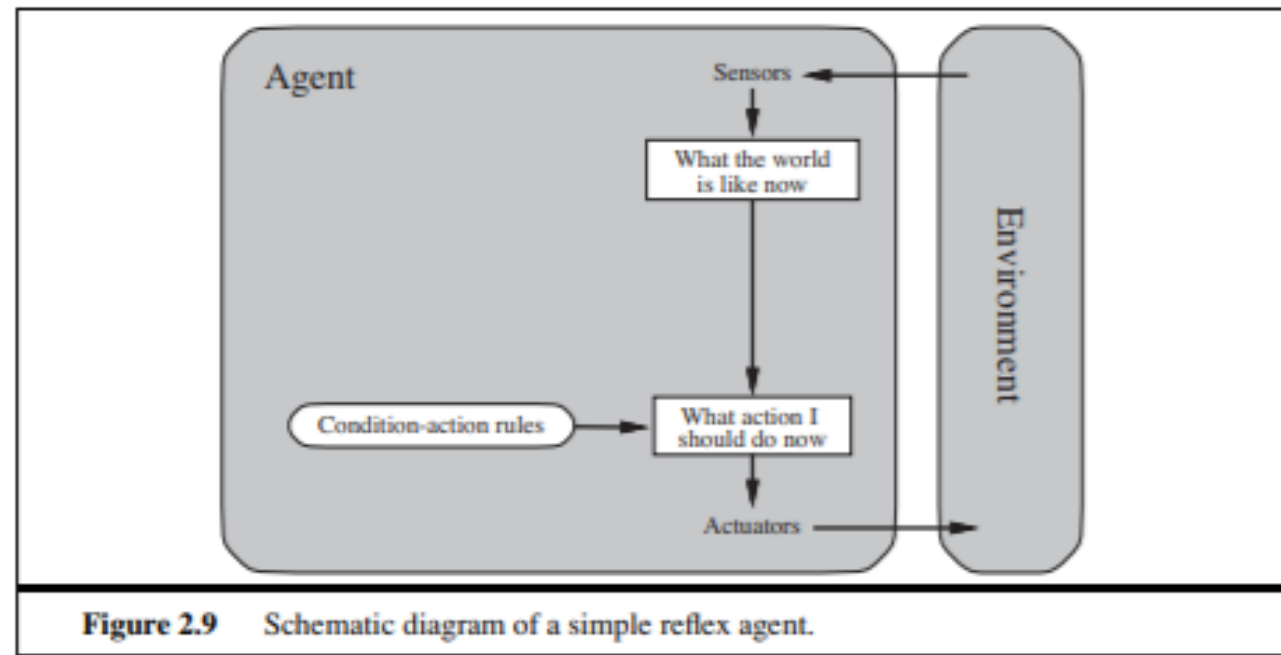


Structure of Agents

- Kinds of agents
 - Simple reflex agents
 - Model based reflex agents
 - Goal-based agents
 - Utility-based agents
 - Learning agents


Simple Reflex Agent

- Use simple “if then” rules also called **condition-action rule**
- Can be short sighted
- These agents select actions on the basis of the current percept, ignoring the rest of the percept history.



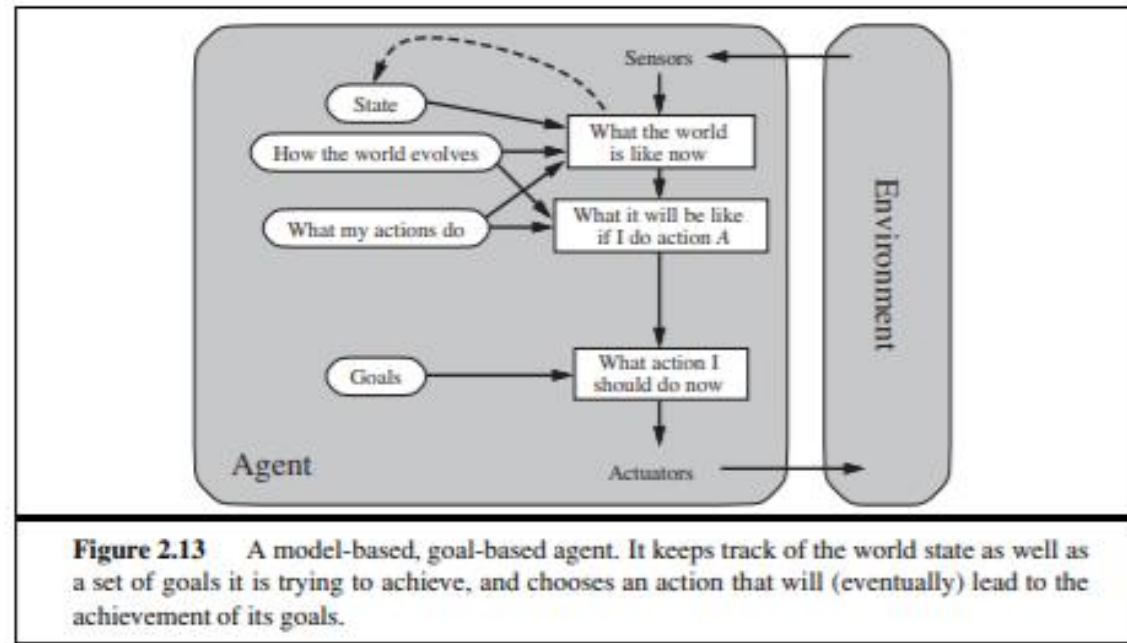
Model based reflex Agent

- Store previously-observed information
- Can reason about unobserved aspects of current state



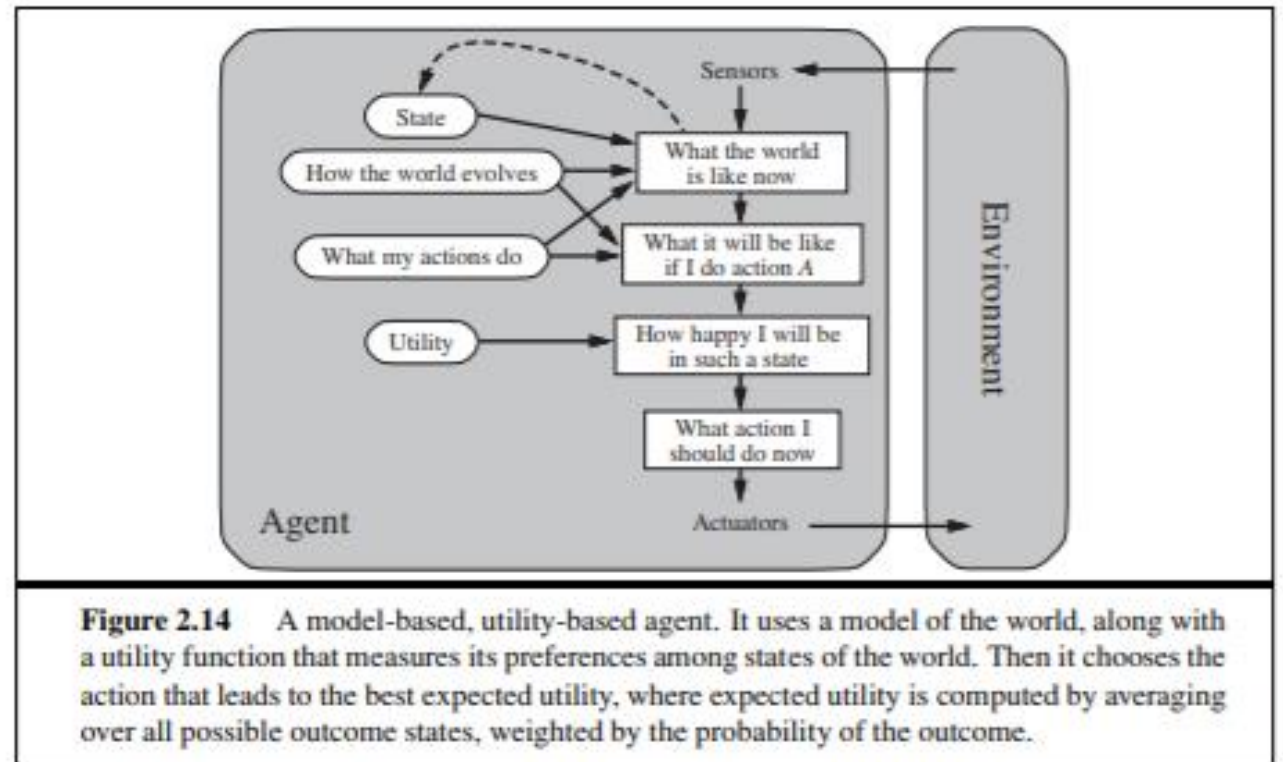
Goal-based Agent

- Goal reflects desires of agents
- May project actions to see if consistent with goals
- Takes time, world may change during reasoning
- They have specific goals or objectives that they try to achieve, and they take actions based on the current percepts and their internal state to reach those goals



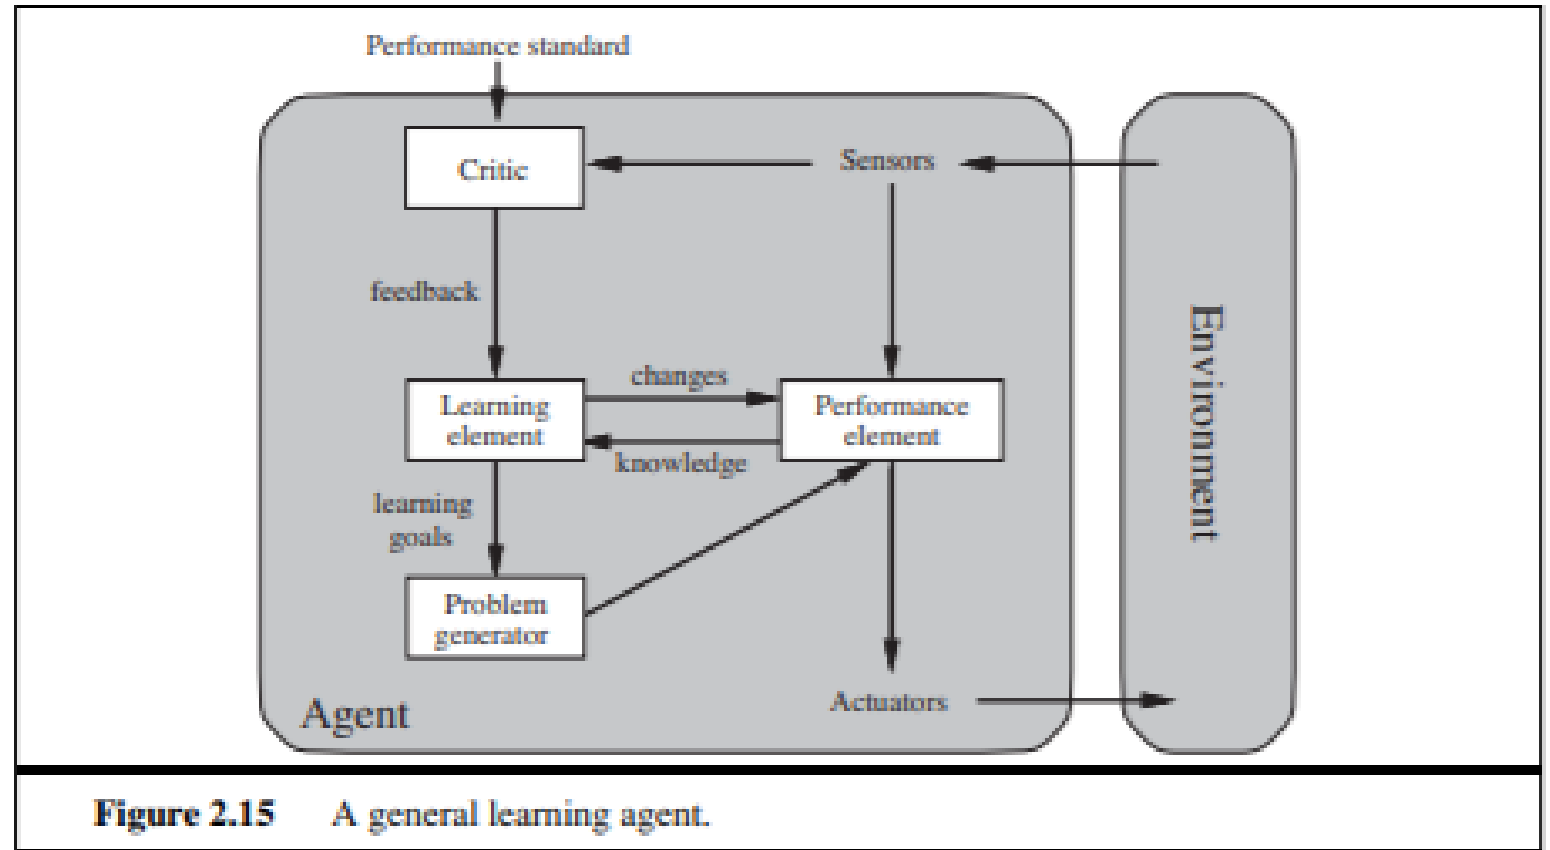
Utility-based Agent

- Goals alone are not enough to generate high-quality behavior in most environments
- Utility is a measure of the value or desirability of a particular state or outcome. The agent uses utility to determine the actions it should take in order to achieve its goals.
- They take into account the long-term consequences of their actions to maximize a specific utility function.



Learning Agent

- They are able to improve their performance over time by learning from their experiences and adjusting their behavior accordingly
- The agent uses past experiences and feedback to continuously improve its decision-making and problem-solving abilities.





1. Simple Reflex Agent:

1. **Example:** A thermostat in a heating system.
2. **Description:** The thermostat senses the current temperature (the percept) and takes a simple action based on a predefined rule: if the temperature is below a certain threshold, turn on the heating; if it's above, turn it off.

2. Model-Based Agent:

1. **Example:** Chess-playing computer program.
2. **Description:** The agent maintains an internal model of the chessboard and the positions of the pieces. It uses this model to simulate and evaluate potential future moves, helping it make decisions based on the expected outcomes.

3. Goal-Based Agent:

1. **Example:** GPS navigation system.
2. **Description:** The agent has a specific goal (the destination) and plans its actions (the route) to achieve that goal. It takes into account the current state (location) and continually updates its plan based on changing conditions, like traffic.

4. Utility-Based Agent:

1. **Example:** Shopping recommendation system.
2. **Description:** The agent recommends products based on user preferences and past behavior. It assigns a utility (or value) to each recommended item, and the recommendation with the highest utility is presented to the user.


5. Learning Agent:

1. **Example:** Email spam filter.
2. **Description:** The agent learns from examples to distinguish between spam and non-spam emails. It continuously updates its knowledge based on user feedback and new examples. Over time, it becomes more accurate in classifying emails.



Homework

- Readings
 - CH 2- Intelligent Agent (Section 2.1 - 2.4)



Artificial Intelligence

CH-3: Solving Problem by Searching



Today's Topic

- Problem Solving Agent
- Problem formulation
 - – What to DO to get a GOAL {What agent type is it?}
- Example problems
- Basic search algorithms



Problem Solving Agent

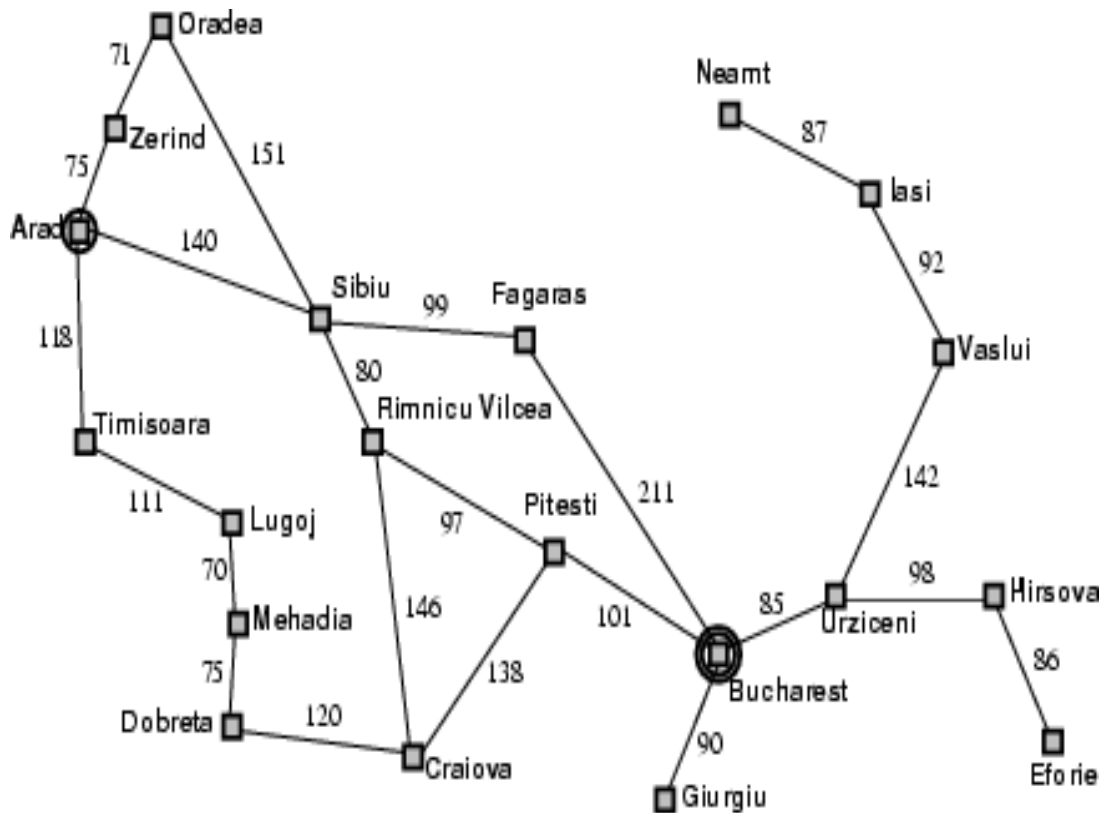
Steps: Goal formulation

Problem formulation

Search

Execute

Example: Romania



- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- **Formulate goal:**
 - be in Bucharest
- Formulate problem:**
 - **states:** various cities
 - **actions:** drive between cities
- **Find solution:**
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest



Problem

A **problem** is defined by four items:

1. initial state e.g., "at Arad"
2. Action, operator or successor function $S(x)$ = set of action-state pairs
 - e.g., $S(Arad) = \{ \langle Arad \rightarrow Zerind, Zerind \rangle, \dots \}$

Initial state + successor function – state space

It defines the **possible actions and resulting states** that can be reached from the state "Arad".

The notation $\langle Arad \rightarrow Zerind, Zerind \rangle$ represents an action-state pair, where "Arad" is the current state, "Zerind" is the next state, and " $\rightarrow Zerind$ " represents the action that leads from "Arad" to "Zerind".



Problem

A **problem** is defined by four items:

3. **goal test**: defined in problem e.g. Bucharest

4. **path cost** (additive) : assigns cost to a path

- e.g., sum of distances, number of actions executed, etc.
- $c(x,a,y)$ is the **step cost**, assumed to be ≥ 0
- More than one solutions... select a preferable solution

A **solution** is a sequence of actions leading from the initial state to a goal state



Formulating Problem

- Formulation of Problem includes: initial state, actions, transition model, goal test and path cost---Model
- Real-World factors:
 - the traveling companions
 - the current radio program
 - the scenery out of the window,
 - the proximity of law enforcement officers,
 - the distance to the next rest stop, the condition of the road, the weather, and so on

The process of removing detail from a representation is called **abstraction**

Problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent: seq, an action sequence, initially empty
               state, some description of the current world state
               goal, a goal, initially null
               problem, a problem formulation

  state  $\leftarrow$  UPDATE-STATE(state, percept)
  if seq is empty then
    goal  $\leftarrow$  FORMULATE-GOAL(state)
    problem  $\leftarrow$  FORMULATE-PROBLEM(state, goal)
    seq  $\leftarrow$  SEARCH(problem)
    if seq = failure then return a null action
  action  $\leftarrow$  FIRST(seq)
  seq  $\leftarrow$  REST(seq)
  return action
```

Figure 3.1 A simple problem-solving agent. It first formulates a goal and a problem, searches for a sequence of actions that would solve the problem, and then executes the actions one at a time. When this is complete, it formulates another goal and starts over.



Example Problems

- Toy Problems:
 - It is intended to illustrate various problem-solving methods
 - Use to compare performance of algorithms
- Real-world Problems:
 - Real-world problems are complex, real-life challenges that require a solution
 - It is the one whose solutions people actually care about

Example: The 8-puzzle

- states?
- actions?
- goal test?
- path cost?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Example: The 8-puzzle

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

[Note: Optimal solution of n-family Puzzle is NP-hard]



- How many moves will be required to reach the goal state?

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State


Vacuum world state space graph

- states?
- actions?
- goal test?
- path cost?



Vacuum world state space graph

- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action





Problem Formulation Examples

- Toy Problem (Vacuum Cleaner)=done
- 8-Puzzle=done
- 8-Queen Problem
- Route-finding Problem
 - Routing in networks; Operations research (military, business etc); air-line travel planning systems
- Touring Problem (TSP=traveling salesperson Problem)
 - State space include current city + set of cities already visited
 - Application to stocking machines on shop floors; automatic circuit drills




Real-life Examples

- Robot navigation (Ch.25),
- VLSI Layout
- Automatic assembly sequencing (Michie, 1972),
- Protein sequencing(Ch.10)
- Scheduling problems

8-Queen Problem

- Place 8 Queens on chessboard such that no queen attack any other
 - no two queens can be placed on the same row, column, or diagonal.
- Problem Formulation
 - **Complete-state:** all queens placed on the board
 - **Incremental:** maintaining a partially filled board and adding one queen at a time, so it does not threaten any of the previously placed queens





Complete State formulation

- Considers the entire state of the board as a single entity
- Starts with all 8 queens on the board and moves them around until a solution is found.
- **goal test:** 8 queens on board, none attacked
- **path cost:** irrelevant
- **states:** any arrangement of 0-8 queens on the board
- **operators:** add or remove a queen to/from any square
- 64^8 possible combinations to investigate??



Complete State formulation

A more sensible choice would use the fact that placing a queen where it is already under attack cannot work because subsequent placings will not undo the attack. So, try the following instead:

- **States:** Arrangement of n (0 to 8) queens on board, one per column in the leftmost n columns, with no queen attacking any other
- **Action:** Add queen in leftmost empty column such that queen is not attacking any other queen



Incremental formulation

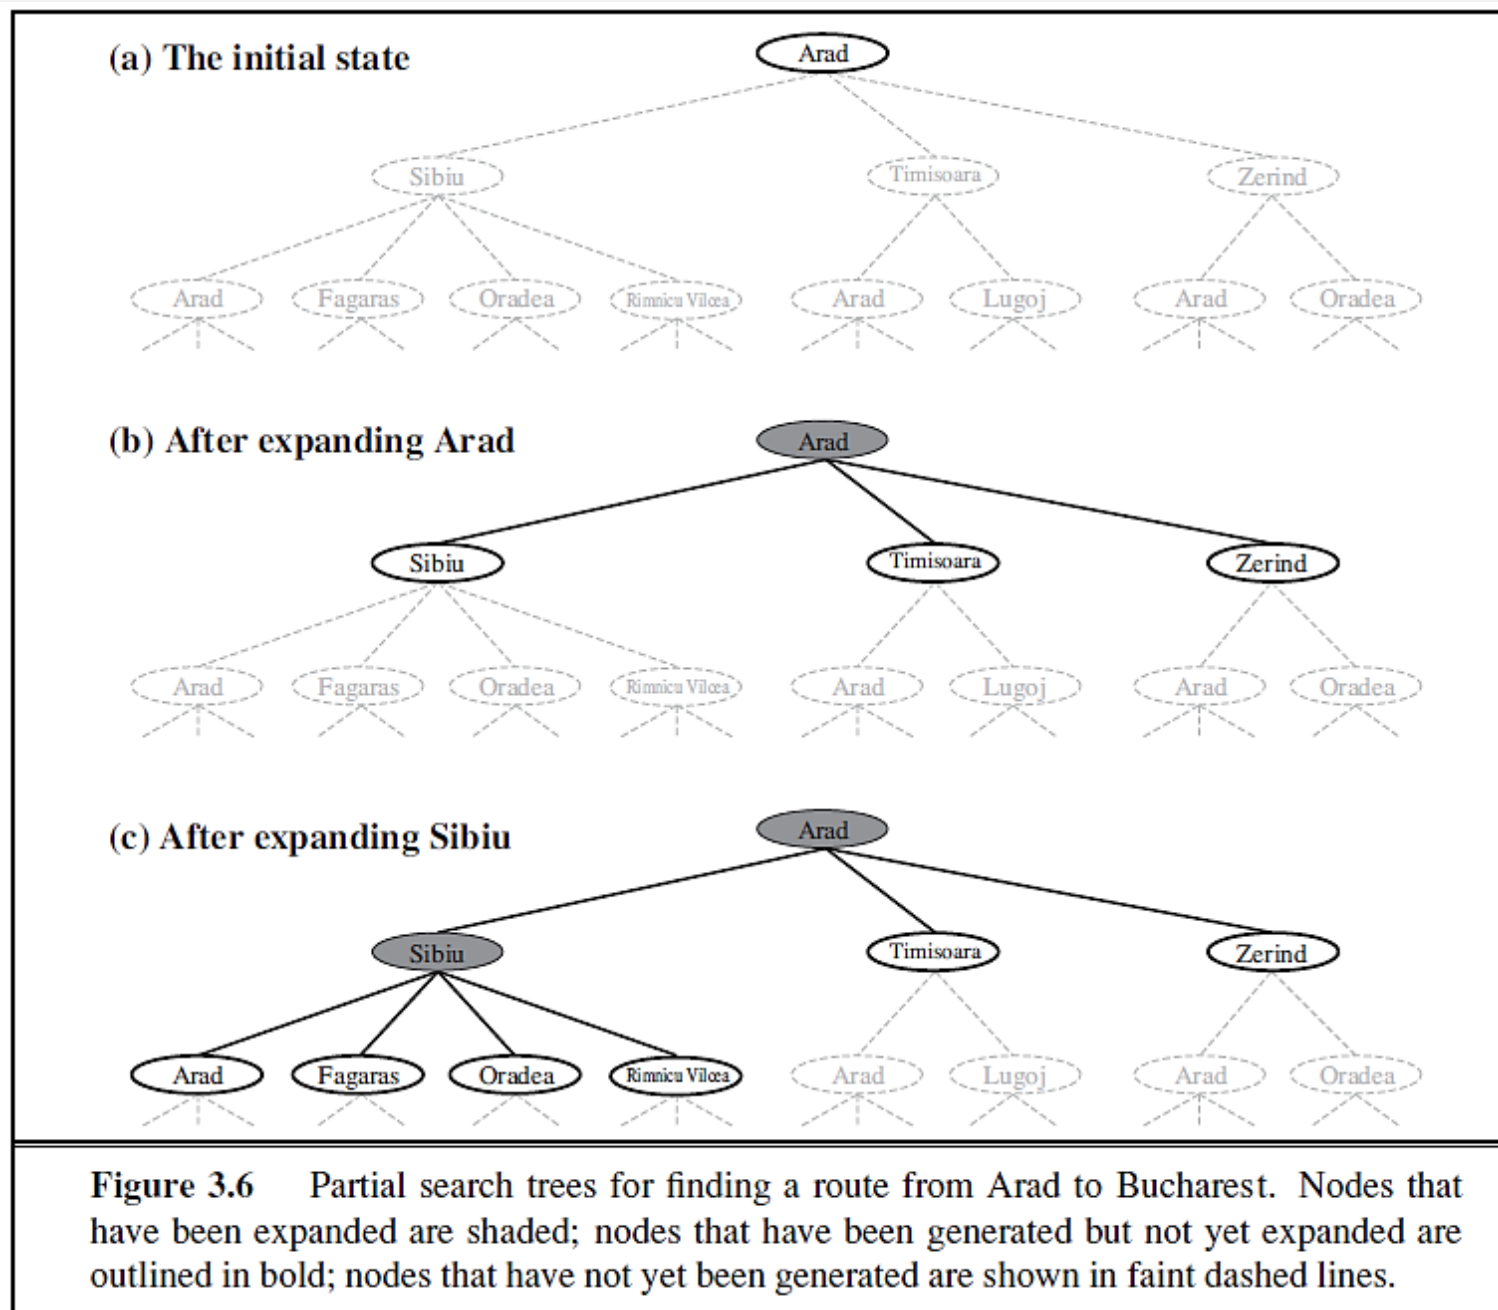
- States: Any arrangement of 0 to 8 queens on the board
- Initial state: No queens on board
- Action: Add queen to any empty cell(square)
- Goal Test: 8 queens on board and no attack
- Path cost: Not Interested
- multiplying the number of possibilities for each queen



Searching for Solutions


- Search the state space:
 - State space can be represented by a **search tree**
 - Root of the tree is the initial state also called **search node**
 - Children generated through successor function
 - The choice of which state to expand is determined by **search strategy**

Tree Search Example




Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost** $g(x)$, **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

State-space to graph formulation & Searching





Search Graph

A **graph** consists of a set N of **nodes** and a set A of ordered pairs of nodes, called **arcs or edges**.

Node n_2 is a **neighbor** of n_1 if there is an arc from n_1 to n_2 . That is, if $\langle n_1, n_2 \rangle \in A$.

A **path** is a sequence of nodes $n_0, n_1, n_2, \dots, n_k$ such that $\langle n_{i-1}, n_i \rangle \in A$.

A **cycle** is a non-empty path such that the start node is the same as the end node

A **directed acyclic graph** (DAG) is a graph with no cycles

Given a start node and goal nodes, a **solution** is a path from a start node to a goal node.



Search strategies & Measuring Problem-solving Algorithm's Performance

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **Completeness**: Does it always find a solution if one exists?
 - **Time complexity**: How long does it take to solve the problem (number of nodes generated)
 - **Space complexity**: How much memory is needed to perform the search (maximum number of nodes in memory)
 - **Optimality**: does it always find a least-cost solution?



Measuring Problem-solving Algorithm's Performance

- In AI, graph is implicitly represented by initial state and successor function and complexities is expressed in terms of three quantities:
 - Branching factor (b) – maximum number of successors of any node
 - Depth (d) – depth of shallowest goal node
 - m – maximum length of any path in the state space
- Time is often measured in terms of no of nodes generated during the search and space in terms of maximum number of nodes stored in memory
- To assess the effectiveness of a search algorithm, we also consider search cost— which typically depends on the time complexity but can also include a term for memory usage—or we can use the total cost, which combines the **search cost** and **the path cost** of the solution found.



Branching Factor

The ***forward branching factor*** of a node is the number of arcs going out of the node

The ***backward branching factor*** of a node is the number of arcs going into the node

