

CS4049

Bioinformatics

Spring 2025

Rushda Muneer

Count and Profile Matrices

- **Count(Motifs)** counts the occurrences of each nucleotide in the corresponding column of the motif matrix.
- Divide each element in the count matrix by **t**, the number of rows in the motif matrix.
- This results in a **profile matrix** $P = \text{Profile}(\text{Motifs})$, where each element $P(i, j)$ is the frequency of nucleotide **i** in column **j**.
- Each column of the profile matrix sums to **1** (representing nucleotide frequencies).

Motifs	T	C	G	G	G	G	g	T	T	T	t	t	
	c	C	G	G	t	G	A	c	T	T	a	C	
	a	C	G	G	G	G	A	T	T	T	t	C	
	T	t	G	G	G	G	A	c	T	T	t	t	
	a	a	G	G	G	G	A	c	T	T	C	C	
	T	t	G	G	G	G	A	c	T	T	C	C	
	T	C	G	G	G	G	A	T	T	c	a	t	
	T	C	G	G	G	G	A	T	T	c	C	t	
	T	a	G	G	G	G	A	a	c	T	a	C	
T	C	G	G	G	t	A	T	a	a	C	C		
(Motifs)	3 + 4 + 0 + 0 + 1 + 1 + 1 + 5 + 2 + 3 + 6 + 4 = 30												
(Motifs)	A:	2	2	0	0	0	0	9	1	1	1	3	0
	C:	1	6	0	0	0	0	0	4	1	2	4	6
	G:	0	0	10	10	9	9	1	0	0	0	0	0
	T:	7	2	0	0	1	1	0	5	8	7	3	4
(Motifs)	A:	.2	.2	0	0	0	0	.9	.1	.1	.1	.3	0
	C:	.1	.6	0	0	0	0	0	.4	.1	.2	.4	.6
	G:	0	0	1	1	.9	.9	.1	0	0	0	0	0
	T:	.7	.2	0	0	.1	.1	0	.5	.8	.7	.3	.4

Using the Profile Matrix to Roll Dice

- Let *Motifs* be a collection of *k*-mers taken from *t* strings *Dna*
- View each column of *Profile*(*Motifs*) as a **four-sided biased die**
- A profile matrix with ***k* columns** can be viewed as a **collection of *k* dice**, which we will **roll to randomly generate a *k*-mer**

Example: The probability $\Pr(\text{"ACGGGGATTACC"} \mid \text{Profile})$ that *Profile* generates "ACGGGGATTACC" is computed by simply multiplying the entries of each nucleotide in respective column of the profile matrix.

	A:	.2	.2	0	0	0	0	.9	.1	.1	.1	.3	0
	C:	.1	.6	.0	0	0	0	0	.4	.1	.2	.4	.6
Profile	G:	0	0	1	1	.9	.9	.1	0	0	0	0	0
	T:	.7	.2	0	0	.1	.1	0	.5	.8	.7	.3	.4

$$\Pr(\text{ACGGGGATTACC}, \text{Profile}) = .2 \cdot .6 \cdot 1 \cdot 1 \cdot .9 \cdot .9 \cdot .9 \cdot .5 \cdot .8 \cdot .1 \cdot .4 \cdot .6$$

$$= 0.000839808$$

The probability to generate "ACGGGGATTACC" based upon the given profile matrix is 0.000839808

Profile - Most Probable k-mer

- K-mer with the highest $\text{Pr}(\text{k-mer} \mid \text{Profile})$ among all k-mers in the sequence

Example: Finding profile-most probable 6-mer in
CTATAAACCTTACAT

A	1/2	7/8	3/8	0	1/8	0
C	1/8	0	1/2	5/8	3/8	0
T	1/8	1/8	0	0	1/4	7/8
G	1/4	0	1/8	3/8	1/4	1/8

6-mer	Pr (6-mer Profile)	
CTATAA ACCTTACAT	$1/8 \times 1/8 \times 3/8 \times 0 \times 1/8 \times 0$	0
CTATA AA CCTTACAT	$1/8 \times 7/8 \times 0 \times 0 \times 1/8 \times 0$	0
CTATA AA CCTTACAT	$1/2 \times 1/8 \times 3/8 \times 0 \times 1/8 \times 0$	0
CTATA AA C C TTACAT	$1/8 \times 7/8 \times 3/8 \times 0 \times 3/8 \times 0$	0
CTAT AA C C TTACAT	$1/2 \times 7/8 \times 3/8 \times 5/8 \times 3/8 \times 7/8$	0.0336
CTATA AA CCTTACAT	$1/2 \times 7/8 \times 1/2 \times 5/8 \times 1/4 \times 7/8$	0.0299
CTATA AA CCTTACAT	$1/2 \times 0 \times 1/2 \times 0 \times 1/4 \times 0$	0
CTATA AA CCTTACAT	$1/8 \times 0 \times 0 \times 0 \times 1/8 \times 0$	0
CTATA AA CCTTACAT	$1/8 \times 1/8 \times 0 \times 0 \times 3/8 \times 0$	0
CTATA AA CCTTACAT	$1/8 \times 1/8 \times 3/8 \times 5/8 \times 1/8 \times 7/8$	0.0004

Greedy Motif Search

1. Select a **k-mer**, Motif₁ in DNA₁.
2. Construct a **profile matrix**, Profile for this lone k-mer.
3. Set Motif₂ as the **Profile-most probable k-mer** in DNA₂.
4. **Update** Profile using Motif₁ and Motif₂.
5. Set Motif₃ as the **Profile-most probable k-mer** in DNA₃.
6. Repeat for each **DNA_i**, selecting the Profile-most probable k-mer.
7. After obtaining a k-mer from each string, **form a collection Motifs**.
8. Compare Motifs with the **current best scoring motif collection**.
9. If the new Motifs collection has a **higher score**, update the **best collection**.
10. Move Motif₁ one symbol over in Dna1 and repeat the process.

Analyzing Greedy Motif Finding

- GreedyMotifSearch() is significantly faster than MedianString().
- **Speed vs. Accuracy Trade-off:**
 - GreedyMotifSearch() sacrifices accuracy for speed.
 - Returns "gtAAAtAgaGatGtG" with a distance of 8.
 - The true implanted motif is "AAAAAAAAAGGGGGGG", showing a clear deviation.
- While fast, GreedyMotifSearch() does not always find the most accurate motif.
- Useful for heuristic solutions but may require refinement or additional strategies for precise motif discovery.

Laplace's Rule of Succession

Zero Probability Issue:

- The fourth symbol of "TCGTGGATTCC" results in $\Pr(\text{"TCGTGGATTCC"}, \text{Profile})$ being zero.
- This causes the entire string to have a probability of zero, despite differing from the consensus string at only one position.

Profile	A:	.2	.2	.0	.0	.0	.0	.9	.1	.1	.1	.3	.0
	C:	.1	.6	.0	0	.0	.0	.0	.4	.1	.2	.4	.6
	G:	.0	.0	1	1	.9	.9	.1	.0	.0	.0	.0	.0
	T:	.7	.2	.0	.0	.1	.1	.0	.5	.8	.7	.3	.4

$$\Pr(\text{"TCGTGGATTCC"}, \text{Profile}) = .7 \cdot .6 \cdot 1 \cdot .0 \cdot .9 \cdot .9 \cdot .9 \cdot .5 \cdot .8 \cdot .7 \cdot .4 \cdot .6 = 0$$

- A simple method for introducing **pseudocounts** (a small number used instead of zero).
- Inspired by Laplace's principle used to estimate rare event probabilities.
- In motif finding, typically involves adding 1 (or another small number) to each element of $\text{Count}(\text{Motifs})$.

Laplace's Rule of Succession

Example: we have the following motif, count, and profile matrices

					T	A	A	C						
					G	T	C	T						
					A	C	T	A						
					A	G	G	T						
Count(Motifs)	A:	2	1	1	1					Profile(Motifs)	2/4	1/4	1/4	1/4
	C:	0	1	1	1						0	1/4	1/4	1/4
	G:	1	1	1	0						1/4	1/4	1/4	0
	T:	1	1	1	2						1/4	1/4	1/4	2/4

Laplace's Rule of Succession adds 1 to each element of *Count(Motifs)*, updates the two matrices

Count(Motifs)	A:	2+1	1+1	1+1	1+1	Profile(Motifs)	3/8	2/8	2/8	2/8
	C:	0+1	1+1	1+1	1+1		1/8	2/8	2/8	2/8
	G:	1+1	1+1	1+1	0+1		2/8	2/8	2/8	1/8
	T:	1+1	1+1	1+1	2+1		2/8	2/8	2/8	3/8

Randomized Motif Search

- Begin from a collection of randomly chosen k-mers *Motifs* in *Dna*
- Construct *Profile*(*Motifs*), and use this profile to generate a new collection of k-mers
 - *Motifs*(*Profile*(*Motifs*), *Dna*)
- Hope is that *Motifs*(*Profile*(*Motifs*), *Dna*) has a better score than the original collection of k-mers *Motifs*
- Form the profile matrix of these k-mers
 - *Profile*(*Motifs*(*Profile*(*Motifs*), *Dna*))
- Use it to form the most probable k-mers
 - *Motifs*(*Profile*(*Motifs*(*Profile*(*Motifs*), *Dna*)), *Dna*)
- Continue to iterate. . .
 - ...*Profile*(*Motifs*(*Profile*(*Motifs*(*Profile*(*Motifs*), *Dna*)), *Dna*))...
- As long as the score of the constructed motifs keeps improving

Entropy

- A measure of the uncertainty of a probability distribution (p_1, \dots, p_N) , and is defined as follows:

$$H(p_1, \dots, p_N) = - \sum_{i=1}^N p_i \cdot \log_2 p_i$$

Profile	A:	.2	.2	0	0	0	0	.9	.1	.1	.1	.3	0
	C:	.1	.6	0	0	0	0	0	.4	.1	.2	.4	.6
	G:	0	0	1	1	.9	.9	.1	0	0	0	0	0
	T:	.7	.2	0	0	.1	.1	0	.5	.8	.7	.3	.4

Examples of Entropy Calculations:

2nd column: $H(0.2, 0.6, 0.0, 0.2) \rightarrow - (0.2 \log_2 0.2 + 0.6 \log_2 0.6 + 0.0 \log_2 0.0 + 0.2$

$\log_2 0.2) \approx 1.371$

12th column: $H(0.0, 0.6, 0.0, 0.4) \rightarrow - (0.0 \log_2 0.0 + 0.6 \log_2 0.6 + 0.0 \log_2 0.0 + 0.4$

$\log_2 0.4) \approx 0.971$

5th column: $H(0.0, 0.0, 0.9, 0.1) \rightarrow - (0.0 \log_2 0.0 + 0.0 \log_2 0.0 + 0.9 \log_2 0.9 + 0.1$

$\log_2 0.1) \approx 0.467$

Entropy

- **STOP and Think:** What are the **maximum** and **minimum** possible values for the entropy of a **probability distribution** containing **four** values?
- **Max** = $-1 * (0.25 \log_2 0.25 + 0.25 \log_2 0.25 + 0.25 \log_2 0.25 + 0.25 \log_2 0.25) = 2.0$
- **Min** = $-1 * \log_2 1 = 0.0$
- **Minimum and Maximum Entropy:**
 - A completely conserved column has entropy = 0 (minimum entropy).
 - A column with equally-likely nucleotides (each probability = 1/4) has maximum entropy = 2.
 - The more conserved a column, the smaller its entropy.

Motif Logo

- A diagram for visualizing motif conservation that consists of a stack of letters at each position
- The relative sizes of letters indicate their frequency in the column
- The total height of the letters in a column is based on the **information content** of the column
 - Defined as $2 - H(p_1, \dots, p_N)$
 - Lower entropy = higher information content.
 - Taller columns in the motif logo represent highly conserved positions.



Gibbs Sampling

- A more cautious iterative algorithm.
- Discards only a single k-mer at each iteration unlike Randomized Motif Search which may change all motifs in a single iteration.
- **GibbsSampler()** randomly selects an integer i between 1 and t , and then randomly changes only a single k -mer $Motif_i$.

ttacctt aac	t tac cttaac	ttacctt aac	ttacctt aac
g ata tctgtc	gat atc tgtc	g ata tctgtc	gatatc tgt c
acg gcgttcg	acggcg ttc g	acg gcgttcg	acg gcgttcg
ccct aaa gag	ccctaa aga g	ccct aaa gag	ccct aaa gag
cgtc aga ggt	cgt cagaggt	cgtc aga ggt	cgtc aga ggt
RandomizedMotifSearch		GibbsSampler	
(may change all k-mers in one step)		(changes one k-mer in one step)	

Proposal Submission

- Sunday, February 9th, 2025 11:59 pm
- 2 pages Proposal
- Project title
- Group members (2-3), contact details
- Abstract (250 words)
- Introduction, Literature Review
- Proposed Methodology, Tentative/algorithms or tools
- Expected outcome or deliverables
- References

Quiz 1

- Wednesday, February 12th, 2025 11:59 pm
- Chapter 1 and 2 (all lecture slides 1.1 to 3.1)

Assignment 2

- Wednesday, February 12th, 2025 11:59 pm
- Programming based with some related questions