


Deep Learning

Syed Irtaza Muzaffar

Training Neural Networks: Forward and Backward Propagation

Neural Networks




Output of a neural network can be visualised graphically as *forward propagation of information*.

Neural Networks

Notation

- ▶ Input layer neurons will be indexed by i .
- ▶ Hidden layer neurons will be indexed by j .
- ▶ Next hidden layer or output layer neurons will be indexed by k .
- ▶ Weights of j -th hidden neuron will be denoted by the vector $\mathbf{w}_j^{(1)} \in \mathbb{R}^D$.
- ▶ Weight between i -th input neuron and j -th hidden neuron is $w_{ji}^{(1)}$.
- ▶ Weights of k -th output neuron will be denoted by the vector $\mathbf{w}_k^{(2)} \in \mathbb{R}^M$.
- ▶ Weight between j -th hidden neuron and k -th output neuron is $w_{kj}^{(2)}$.




Neural Networks

Forward Propagation

- ▶ For input \mathbf{x} , denote output of hidden layer as the vector $\mathbf{z}(\mathbf{x}) \in \mathbb{R}^M$.
- ▶ Model $z_j(\mathbf{x})$ as a non-linear function $h(a_j)$ where *pre-activation* $a_j = \mathbf{w}_j^{(1)T} \mathbf{x}$ with adjustable parameters $\mathbf{w}_j^{(1)}$.
- ▶ So the k -th output can be written as

$$\begin{aligned}
 y_k(\mathbf{x}) &= f(a_k) = f(\mathbf{w}_k^{(2)T} \mathbf{z}(\mathbf{x})) \\
 &= f\left(\sum_{j=1}^M w_{kj}^{(2)} z_j(\mathbf{x}) + w_{k0}^{(2)}\right) = f\left(\sum_{j=1}^M w_{kj}^{(2)} h\left(\sum_{i=0}^D w_{ji}^{(1)} x_i\right) + w_{k0}^{(2)}\right)
 \end{aligned}$$

where we have prepended $x_0 = 1$ to absorb bias input and $w_{j0}^{(1)}$ and $w_{k0}^{(2)}$ represent biases.



Neural Networks

Forward Propagation

► The computation

$$y_k(\mathbf{x}, \mathbf{W}) = f \left(\sum_{j=1}^M w_{kj}^{(2)} h \left(\sum_{i=0}^D w_{ji}^{(1)} x_i \right) + w_{k0}^{(2)} \right)$$

can be viewed in two stages:

1. $z_j = h(\mathbf{w}_j^{(1)T} \mathbf{x})$ for $j = 1, \dots, M$.
2. $y_k = f(\mathbf{w}_k^{(2)T} \mathbf{z})$.

Neural Networks

Forward Propagation

- If we define the matrices

$$\mathbf{W}^{(1)} = \underbrace{\begin{bmatrix} \leftarrow \mathbf{w}_1^{(1)T} \rightarrow \\ \leftarrow \mathbf{w}_2^{(1)T} \rightarrow \\ \vdots \\ \leftarrow \mathbf{w}_M^{(1)T} \rightarrow \end{bmatrix}}_{M \times (D+1)} \quad \text{and} \quad \mathbf{W}^{(2)} = \underbrace{\begin{bmatrix} \leftarrow \mathbf{w}_1^{(2)T} \rightarrow \\ \leftarrow \mathbf{w}_2^{(2)T} \rightarrow \\ \vdots \\ \leftarrow \mathbf{w}_K^{(2)T} \rightarrow \end{bmatrix}}_{K \times (M+1)}$$

then forward propagation constitutes

1. $\mathbf{z} = h(\mathbf{W}^{(1)}\mathbf{x})$.
2. Prepend 1 to \mathbf{z} .
3. $\mathbf{y} = f(\mathbf{W}^{(2)}\mathbf{z})$.

Neural Networks for Regression


Gradients

- ▶ Regression requires continuous output $y_k \in \mathbb{R}$.
- ▶ So use *identity* activation function $y_k = f(a_k) = a_k$.
- ▶ Loss can be written as

$$L(\mathbf{W}^{(1)}, \mathbf{W}^{(2)}) = \frac{1}{2} \sum_{n=1}^N \underbrace{\|\mathbf{y}_n - \mathbf{t}_n\|^2}_{L_n} = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_{nk} - t_{nk})^2$$


- ▶ Loss L depends on sum of individual losses L_n .
- ▶ In the following, we will focus on loss L_n for the n -th training sample.
- ▶ We will drop n for notational clarity and refer to L_n simply as L .

How do weights influence loss?



- ▶ $w_{kj}^{(2)}$ influences $a_k^{(2)}$ which influences y_k which influences L .
- ▶ For scalar dependencies, use chain rule.
- ▶ $w_{ji}^{(1)}$ influences $a_j^{(1)}$ which influences z_j which influences $a_1^{(2)}, a_2^{(2)}, a_3^{(2)}$ which influence y_1, y_2, y_3 which influence L .
- ▶ For vector/multivariate dependencies, use multivariate chain rule.

How do weights influence loss?



- Layer 2: $L \leftarrow y_k \leftarrow a_k^{(2)} \leftarrow w_{kj}^{(2)}$.

$$L(y_k(a_k^{(2)}(w_{kj}^{(2)})))$$

- Layer 1: $L \leftarrow \mathbf{y} \leftarrow \mathbf{a}^{(2)} \leftarrow \mathbf{z}_j \leftarrow a_j^{(1)} \leftarrow w_{ji}^{(1)}$.

$$L(\underbrace{y_1(a_1^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)}))))}_{y_1(w_{ji}^{(1)})}, \underbrace{y_2(a_2^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)}))))}_{y_2(w_{ji}^{(1)})}, \dots, \underbrace{y_k(a_k^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)}))))}_{y_k(w_{ji}^{(1)})})$$

Multivariate Chain Rule


- ▶ The chain rule of differentiation states

$$\frac{df(u(x))}{dx} = \frac{df}{du} \frac{du}{dx}$$


- ▶ The *multivariate* chain rule of differentiation states

$$\frac{df(u(x), v(x))}{dx} = \frac{\partial f}{\partial u} \frac{du}{dx} + \frac{\partial f}{\partial v} \frac{dv}{dx}$$

- ▶ The multivariate chain rule applied to compute derivatives w.r.t weights of hidden layers has a special name – *backpropagation*.




Backpropagation



- For the output layer weights

$$\frac{\partial L(y_k(a_k^{(2)}(w_{kj}^{(2)})))}{\partial w_{kj}^{(2)}} = \frac{\partial L}{\partial a_k^{(2)}} \frac{\partial a_k^{(2)}}{\partial w_{kj}^{(2)}} = \delta_k z_j$$

Backpropagation



- For the hidden layer weights, using the multivariate chain rule

$$\begin{aligned}
 & \frac{\partial}{\partial w_{ji}^{(1)}} L(y_1(a_1^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)}))))), y_2(a_2^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)}))))), \dots, y_k(a_k^{(2)}(z_j(a_j^{(1)}(w_{ji}^{(1)})))))) \\
 &= \frac{\partial L}{\partial a_j^{(1)}} \frac{\partial a_j^{(1)}}{\partial w_{ji}^{(1)}} = \underbrace{\sum_{k=1}^K \underbrace{\frac{\partial L}{\partial a_k^{(2)}}}_{\delta_k} \underbrace{\frac{\partial a_k^{(2)}}{\partial z_j}}_{w_{kj}^{(2)}} \underbrace{\frac{\partial z_j}{\partial a_j^{(1)}}}_{h'(a_j^{(1)})} \underbrace{\frac{\partial a_j^{(1)}}{\partial w_{ji}^{(1)}}}_{x_i}}_{\frac{\partial L}{\partial a_j^{(1)}} = \delta_j} = \delta_j x_i
 \end{aligned}$$

- Notice that gradient = $\delta \times \text{input}$ *Deep Learning*

Backpropagation

- It is important to note that

$$\delta_j = h'(a_j) \sum_{k=1}^K \delta_k w_{kj}$$

yields the error δ_j at hidden neuron j by *backpropagating* the errors δ_k from all output neurons that use the output of neuron j .

- More generally, compute error δ_j at a layer by *backpropagating* the errors δ_k from next layer.
- Hence the names *error backpropagation*, *backpropagation*, or simply *backprop*.
- Very useful machine learning technique that is *not limited to neural networks*.

Backpropagation

$$\delta_j^{(1)} = h'(a_j) \sum_{k=1}^K \delta_k^{(2)} w_{kj}$$




Figure: Visual representation of backpropagation of delta values of layer $l + 1$ to compute delta values of layer l .

Backpropagation

Learning Algorithm

1. Forward propagate the input vector \mathbf{x}_n to compute *and store* activations and outputs of every neuron in every layer.
2. Evaluate $\delta_k = \frac{\partial L_n}{\partial a_k}$ for every neuron in output layer.
3. Evaluate $\delta_j = \frac{\partial L_n}{\partial a_j}$ for every neuron in *every* hidden layer via backpropagation.

$$\delta_j = h'(a_j) \sum_{k=1}^K \delta_k w_{kj}$$

4. Compute derivative of each weight $\frac{\partial L_n}{\partial w}$ via $\delta \times \text{input}$.
5. Update each weight via gradient descent $w^{\tau+1} = w^{\tau} - \eta \frac{\partial L_n}{\partial w}$.

Tanh

$A(-1, 1)$ *sigmoidal function*

- ▶ Since range of logistic sigmoid $\sigma(a)$ is $(0, 1)$, we can obtain a function with $(-1, 1)$ range as $2\sigma(a) - 1$.
- ▶ Another related function with $(-1, 1)$ range is the **tanh** function.

$$\tanh(a) = 2\sigma(2a) - 1 = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

where σ is applied on $2a$.

- ▶ Preferred¹ over logistic sigmoid as activation function $h(a)$ of hidden neurons.
- ▶ Just like the logistic sigmoid, derivative of $\tanh(a)$ is simple: $1 - \tanh^2(a)$. (Prove it.)

¹LeCun et al., 'Efficient backprop'.

A Simple Example

- ▶ Two-layer MLP for multivariate regression from $\mathbb{R}^D \rightarrow \mathbb{R}^K$.
- ▶ Linear outputs $y_k = a_k$ with half-SSE $L = \frac{1}{2} \sum_{k=1}^K (y_k - t_k)^2$.
- ▶ M hidden neurons with $\tanh(\cdot)$ activation functions.

Forward propagation

$$a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i$$

$$z_j = \tanh(a_j)$$

$$z_0 = 1$$

$$y_k = \sum_{j=0}^M w_{kj}^{(2)} z_j$$

$$\delta_k = y_k - t_k$$

- ▶ Compute derivatives $\frac{\partial L}{\partial w_{ji}^{(1)}} = \delta_j x_i$ and $\frac{\partial L}{\partial w_{kj}^{(2)}} = \delta_k z_j$.

Backpropagation

$$\delta_j = (1 - z_j^2) \sum_{k=1}^K w_{kj}^{(2)} \delta_k$$

Backpropagation

Verifying Correctness

- ▶ *Numerical derivatives* can be computed via finite *central differences*

$$\frac{\partial L_n}{\partial w_{ji}} = \frac{L_n(w_{ji} + \epsilon) - L_n(w_{ji} - \epsilon)}{2\epsilon} + O(\epsilon^2)$$

- ▶ *Analytical derivatives* computed via backpropagation **must be compared** with numerical derivatives for a few examples to verify correctness.
- ▶ Any implementation of analytical derivatives (not just backpropagation) must be compared with numerical derivatives.
- ▶ Notice that we could have avoided backpropagation and computed all required derivatives numerically.
 - ▶ But cost of numerical differentiation is $O(W^2)$ while that of backpropagation is $O(W)$ where W is the total number of weights (and biases) in the network. (Why?)


Neural Network training finds local minimum

- ▶ For optimisation, we notice that \mathbf{w}^* must be a *stationary point* of $E(\mathbf{w})$.
 - ▶ Minimum, maximum, or saddle point.
 - ▶ A saddle point is where gradient vanishes but point is not an extremum (Example).
- ▶ The goal in neural network minimisation is to find a local minimum.
- ▶ A global minimum, *even if found*, cannot be verified as globally minimum.
- ▶ Due to symmetry, there are multiple equivalent local minima. Reaching *any suitable* local minimum is the goal of neural network optimisation.
- ▶ Since there are no analytical solutions for \mathbf{w}^* , we use iterative, numerical procedures.

Optimisation Options

- ▶ Options for iterative optimisation
 - ▶ Online methods
 - ▶ Stochastic gradient descent
 - ▶ Stochastic gradient descent using mini-batches
 - ▶ Batch methods
 - ▶ Batch gradient descent
 - ▶ Conjugate gradient descent
 - ▶ Quasi-Newton methods
- ▶ Online methods
 - ▶ converge faster since parameter updates are more frequent, and
 - ▶ have greater chance of escaping local minima because stationary point w.r.t to whole data set will generally not be a stationary point w.r.t an individual data point.
- ▶ Batch methods: Conjugate gradient descent and quasi-Newton methods
 - ▶ are more robust and faster than batch gradient descent, and
 - ▶ decrease the error function at each iteration until arriving at a minimum.

Problems with sigmoidal neurons



- ▶ For large $|a|$, sigmoid value approaches either 0 or 1. This is called *saturation*.
- ▶ When the sigmoid saturates, the gradient approaches zero.
- ▶ Neurons with sigmoidal activations stop learning when they saturate.
- ▶ When they are not saturated, they are **almost linear**.
- ▶ There is another reason for the gradient to approach zero during backpropagation.

Vanishing Gradients


- ▶ Notice that gradient of the sigmoid is always between 0 and $\frac{1}{4}$.
- ▶ Now consider the backpropagation equation.

$$\delta_j = \underbrace{h'(a_j)}_{\leq \frac{1}{4}} \sum_{k=1}^K w_{kj} \delta_k$$

where δ_k will also contain *at least* one factor of $\leq \frac{1}{4}$.

- ▶ This means that values of δ_j keep getting smaller as we backpropagate towards the early layers.
- ▶ Since gradient = $\delta \times \text{input}$, the gradients also keep getting smaller for the earlier layers. Known as the *vanishing gradients* problem.
- ▶ *Therefore, while the network might be deep, learning will not be deep.*

Logistic Sigmoid



Activation function

$$y(a) = \frac{1}{1+e^{-a}}$$

Derivative

$$y'(a) = y(a)(1 - y(a))$$


Maximum magnitude of derivative

$$\frac{1}{4}$$

Problem

Cause vanishing gradients

Hyperbolic Tangent



Activation function

$$y(a) = \tanh(a)$$

Derivative

$$y'(a) = 1 - y^2(a)$$


Maximum magnitude of derivative

1

Problem

Cause vanishing gradients

Rectified Linear Unit (ReLU)



Activation function

$$y(a) = \begin{cases} a & \text{if } a > 0 \\ 0 & \text{if } a \leq 0 \end{cases}$$

Derivative


$$y'(a) = \begin{cases} 1 & \text{if } a > 0 \\ 0 & \text{if } a \leq 0 \end{cases}$$

Advantage
Problem

Avoids vanishing gradients
Dead neurons²

²This can be an advantage as well since death implies fewer neurons.

Leaky ReLU




Activation function $y(a) = \begin{cases} a & \text{if } a > 0 \\ ka & \text{if } a \leq 0 \end{cases}$
where $0 \leq k \leq 1$

Derivative $y'(a) = \begin{cases} 1 & \text{if } a > 0 \\ k & \text{if } a \leq 0 \end{cases}$

Advantage Neuron is always learning

Exponential Linear Unit (ELU)



Activation function

$$y(a) = \begin{cases} a & \text{if } a > 0 \\ k(e^a - 1) & \text{if } a \leq 0 \end{cases}$$

where $k > 0$

Derivative

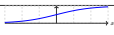




$$y'(a) = \begin{cases} 1 & \text{if } a > 0 \\ y(a) + k & \text{if } a \leq 0 \end{cases}$$

Maximum magnitude of derivative 1

Advantage

Neuron is mostly learning

Better Activation Functions

Name	$f(a)$	Plot	Derivative	Comments
Logistic sigmoid	$\frac{1}{1+e^{-a}}$		$f(a)(1 - f(a))$	Vanishing gradients
Hyperbolic tangent	$\tanh(a)$		$1 - \tanh^2(a)$	Vanishing gradients
Rectified Linear Unit (ReLU)	$\begin{cases} a & \text{if } a > 0 \\ 0 & \text{if } a \leq 0 \end{cases}$		$\begin{cases} 1 \\ 0 \end{cases}$	Dead neurons. Sparsity.
Leaky ReLU	$\begin{cases} a & \text{if } a > 0 \\ ka & \text{if } a \leq 0 \end{cases}$		$\begin{cases} 1 \\ k \end{cases}$	$0 < k < 1$
Exponential Linear Unit (ELU)	$\begin{cases} a & \text{if } a > 0 \\ k(e^a - 1) & \text{if } a \leq 0 \end{cases}$		$\begin{cases} 1 \\ f(a) - k \end{cases}$	$k > 0.$

- ▶ Saturated sigmoidal neurons stop learning. Piecewise-linear units keep learning by avoiding saturation.
- ▶ ELU leads to better accuracy and faster training.
- ▶ *Take home message:* For hidden neurons, use a member of the LU family. They avoid *i)* saturation and *ii)* the vanishing gradient problem.