


Deep Learning

Long Short-Term Memory (LSTM)

Syed Irtaza Muzaffar

Weakness of standard RNN

- ▶ We have already seen that RNNs do not possess long-term memory.
- ▶ Input at time t is soon forgotten because of the recurrent weights W^{11} .
- ▶ Would be nice to decide what and how much to forget/remember based on the input itself.



RNN Cell: Operations at the hidden layer.

Long Short-Term Memory (LSTM)

Building blocks

Let $\mathbf{v}^{(t)} = \begin{bmatrix} \mathbf{h}^{(t-1)} \\ \mathbf{x}^{(t)} \end{bmatrix} \in \mathbb{R}^{(M+D) \times 1}$

Perform 4 affine transformations of $\mathbf{v}^{(t)}$ followed by non-linearities.

$$\mathbf{f}^{(t)} = \sigma \left(W_f \mathbf{v}^{(t)} + \mathbf{b}_f \right) \quad (1)$$

$$\mathbf{i}^{(t)} = \sigma \left(W_i \mathbf{v}^{(t)} + \mathbf{b}_i \right) \quad (2)$$

$$\mathbf{o}^{(t)} = \sigma \left(W_o \mathbf{v}^{(t)} + \mathbf{b}_o \right) \quad (3)$$


$$\tilde{\mathbf{c}}^{(t)} = \tanh \left(W_c \mathbf{v}^{(t)} + \mathbf{b}_c \right) \quad (4)$$

All 4 matrices of size $M \times (M + D)$ and therefore all 4 transformations produce M -dimensional vectors.

Vectors $\mathbf{f}^{(t)}, \mathbf{i}^{(t)}, \mathbf{o}^{(t)}$ contain values in $(0, 1)$ and $\tilde{\mathbf{c}}^{(t)}$ in $(-1, 1)$.

LSTM

Putting everything together



LSTM Cell: Operations at the hidden layer.

- ▶ Vector $\mathbf{c}^{(t)}$ is *interpreted* as the *cell state*.
- ▶ Cell state is recurrent as well.
- ▶ Notice that $\mathbf{c}^{(t)}$ is not forced to contain values in $(0, 1)$ or $(-1, 1)$.

Role of the Gates

$f^{(t)}$: Forget Gate

$$\mathbf{f}^{(t)} = \sigma \left(W_f \mathbf{v}^{(t)} + \mathbf{b}_f \right)$$
$$c_j^{(t)} = f_j^{(t)} c_j^{(t-1)} + i_j^{(t)} \tilde{c}_j^{(t)}$$

- ▶ $f_j^{(t)} \in (0, 1)$ due to the logistic sigmoid.
- ▶ If $f_j^{(t)} = 0$, then $c_j^{(t-1)}$ is *forgotten* in the next state $c_j^{(t)}$.
- ▶ If $f_j^{(t)} = 1$, then $c_j^{(t-1)}$ is *retained completely* in the next state $c_j^{(t)}$.

$\mathbf{f}^{(t)}$ acts as a forget gate on the previous cell state \mathbf{c}^{t-1} .

Role of the Gates

$i^{(t)}$: Input Gate

$$i^{(t)} = \sigma \left(W_i \mathbf{v}^{(t)} + \mathbf{b}_i \right)$$
$$c_j^{(t)} = f_j^{(t)} c_j^{(t-1)} + i_j^{(t)} \tilde{c}_j^{(t)}$$

- ▶ $i_j^{(t)} \in (0, 1)$ due to the logistic sigmoid.
- ▶ If $i_j^{(t)} = 0$, then no new information will be added to $c_j^{(t)}$.
- ▶ If $i_j^{(t)} = 1$, then the potential cell state $\tilde{c}_j^{(t)}$ is *added completely* in the next state $c_j^{(t)}$ *irrespective of forget level*.

$i^{(t)}$ acts as an input gate on the potential cell state $\tilde{c}^{(t)}$.

Role of the Gates

$\mathbf{o}^{(t)}$: Output Gate


$$\mathbf{o}^{(t)} = \sigma \left(W_o \mathbf{v}^{(t)} + \mathbf{b}_o \right)$$

$$h_j^{(t)} = o_j^{(t)} \tanh(c_j^{(t)})$$

- ▶ $o_j^{(t)} \in (0, 1)$ due to the logistic sigmoid.
- ▶ If $o_j^{(t)} = 0$, then cell state $c_j^{(t)}$ will be *completely hidden*.
- ▶ If $o_j^{(t)} = 1$, then cell state $c_j^{(t)}$ is *completely exposed* in both space (\uparrow) and time (\rightarrow).

$\mathbf{o}^{(t)}$ acts as an output gate on $\mathbf{c}^{(t)}$.

LSTM



LSTM Cell: Operations at the hidden layer in detail.

Information flow

- ▶ Depending on $f^{(t)}$ and $i^{(t)}$, an LSTM cell has the ability to push through its cell state $c^{(t-1)}$ exactly or almost unchanged into the next time step $c^{(t)}$.
- ▶ This ensures *flow of the cell state (memory) through time*. Hence long-term memory.
- ▶ This is similar to how other deep learning techniques ensure flow of information in space.
 - ▶ ReLU
 - ▶ Weight initialization
 - ▶ Batchnorm
 - ▶ Residual block

Remembering the past

- ▶ Consider a sentence containing brackets.

England (last year's winners) are expected to put up a good fight.

- ▶ The LSTM cell can learn to set $c_j = 1$ if an opening bracket is seen at time t .
- ▶ It can also learn to keep $c_j = 1$ for a long time until a closing bracket is seen *in the input*.
- ▶ Some other c_k can similarly be used to handle nested brackets and so on.
- ▶ Even the value of c_j itself can be used to signify the level of nesting. It all depends on how and what the LSTM learns.

Peephole Connections

- ▶ Allow gates to look at the cell state as well before deciding what to forget, what to add, and what to output.

$$\mathbf{v}_{f,i}^{(t)} = \begin{bmatrix} \mathbf{c}^{(t-1)} \\ \mathbf{h}^{(t-1)} \\ \mathbf{x}^{(t)} \end{bmatrix} \in \mathbb{R}^{(2M+D) \times 1}$$

$$\mathbf{v}_o^{(t)} = \begin{bmatrix} \mathbf{c}^{(t)} \\ \mathbf{h}^{(t-1)} \\ \mathbf{x}^{(t)} \end{bmatrix} \in \mathbb{R}^{(2M+D) \times 1}$$

Coupled forget and input

- ▶ Use a single forget gate for interpolation.

$$\mathbf{c}^{(t)} = \mathbf{f}^{(t)} \odot \mathbf{c}^{(t-1)} + (1 - \mathbf{f}^{(t)}) \odot \tilde{\mathbf{c}}^{(t)}$$

- ▶ Fewer parameters due to removal of input gate.

Gated Recurrence Unit (GRU)

- ▶ Coupled forget and input gates.
- ▶ Merged hidden and cell state.

$$\mathbf{z}^{(t)} = \sigma \left(W_z \mathbf{v}^{(t)} + \mathbf{b}_z \right)$$

$$\mathbf{r}^{(t)} = \sigma \left(W_r \mathbf{v}^{(t)} + \mathbf{b}_r \right)$$

$$\tilde{\mathbf{h}}^{(t)} = \tanh \left(W_h [\mathbf{r}^{(t)} \odot \mathbf{h}^{(t-1)}; \mathbf{x}^{(t)}] + \mathbf{b}_h \right)$$

$$\mathbf{h}^{(t)} = \left(1 - \mathbf{z}^{(t)} \right) \odot \mathbf{h}^{(t-1)} + \mathbf{z}^{(t)} \odot \tilde{\mathbf{h}}^{(t)}$$

- ▶ Always expose the hidden state.
- ▶ In some variants, the weight matrices can be set to 0.
- ▶ In other variants, the bias vectors can be set to 0.
- ▶ Fewer parameters, faster training, learning from lesser data.

Summary

- ▶ RNNs do not retain long-term context in practice.
- ▶ An LSTM cell has the ability to push through its cell state exactly or almost unchanged into the next time step.
- ▶ This ensures flow of the cell state (memory) through time. Hence long-term memory.