


Deep Learning

Recurrent Neural Networks

Syed Irtaza Muzaffar



*Everything should be made as simple as possible,
but no simpler.*

Albert Einstein

Understanding Recurrent Neural Networks requires some effort and a correct perspective. Do not expect them to be as simple as linear regression.

Static vs. Dynamic Inputs

- ▶ *Static* signals, such as an image, do not change over time.
 - ▶ Ordered with respect to space.
 - ▶ Output depends on current input.
- ▶ *Dynamic* signals, such as text, audio, video or stock price change over time.
 - ▶ Ordered with respect to time.
 - ▶ Output depends on current input as well as past (or even future) inputs.
 - ▶ Also called *temporal*, *sequential* or *time-series* data.


Context in Text

*The Taj _____ was commissioned by Shah Jahan in 1631, to be built in the memory of _____ wife Mumtaz Mahal, who died on 17 June **that** year, giving birth to **their** 14th child, Gauhara Begum. **Construction** started in 1632, and the mausoleum was completed _____ 1643.*

Context in Video




Context in Audio




Time-series Data

- A *single* input will be a *series of vectors* $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^T$.



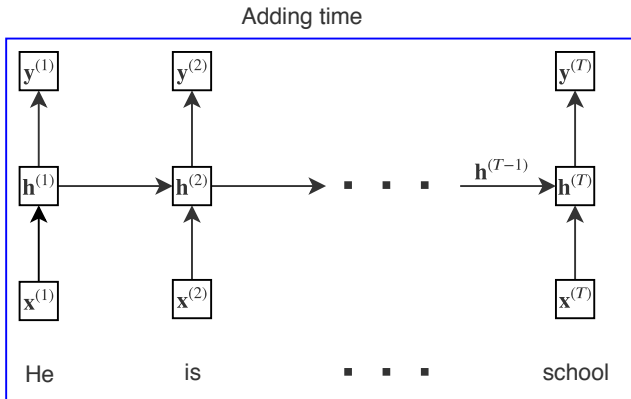
Input component at time t forward propagated through a network.

Representational Shortcut 1 – Space Folding




Each box represents a layer of neurons.

Recurrent Neural Networks




- ▶ A recurrent neural network (RNN) makes hidden state at time t directly dependent on the hidden state at time $t - 1$ and therefore indirectly *on all previous times*.
- ▶ Output y_t depends on all that the network has already seen so far.

Representational Shortcut 2 – Time Folding




Folding
 \Rightarrow
in time



Recurrent Neural Networks

3 sets of weights




$$y^{(t)} = f(\overbrace{W^1 h^{(t)} + b_1}^{a^{1(t)}})$$

$$h^{(t)} = \tanh(\underbrace{W^0 x^{(t)} + W^{11} h^{(t-1)} + b_0}_{a^{0(t)}})$$

Sequence Mappings

One-to-many



Many-to-one







Image caption generation




Sentiment classification

Sequence Mappings


Many-to-many



Many-to-many delayed




Pronoun Verb Adjective



POS tagging

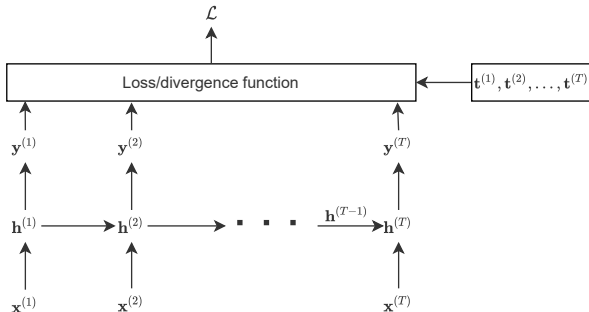
Está loco



Language translation

Loss Functions for Sequences

- For recurrent nets, loss is between *series* of output and target vectors. That is $\mathcal{L}(\{\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(T)}\}, \{\mathbf{t}^{(1)}, \dots, \mathbf{t}^{(T)}\})$.

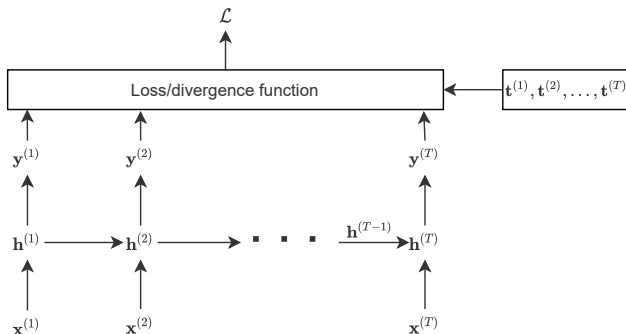


Forward propagation in an RNN unfolded in time.

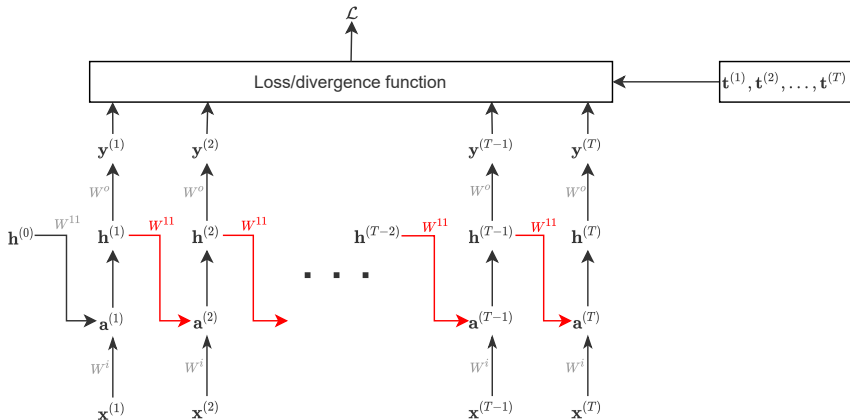
- Notice that loss \mathcal{L} can be computed only after $\mathbf{y}^{(T)}$ has been computed.

Loss Functions for Sequences

- ▶ Loss is *not necessarily* decomposable.
- ▶ In the following, we will assume decomposable loss $\mathcal{L} = \sum_{t=1}^T \mathcal{L}(\mathbf{y}^{(t)}, \mathbf{t}^{(t)})$.
- ▶ In both cases, as long as $\frac{\partial \mathcal{L}}{\partial \mathbf{y}^{(t)}}$ has been computed, backpropagation can proceed.

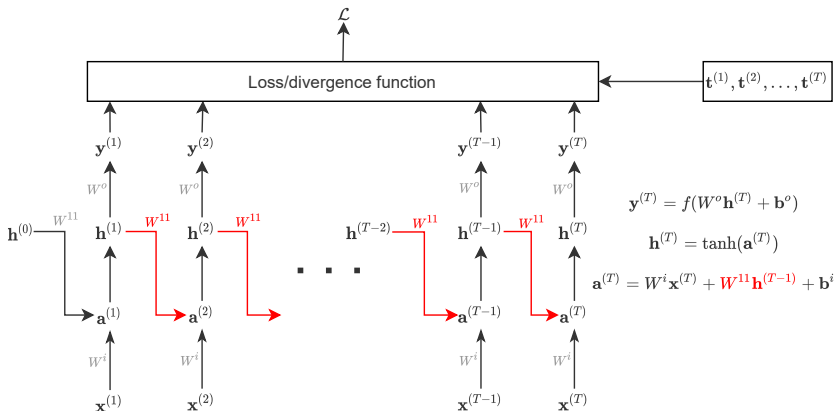


Forward Propagation Through Time



Forward propagation in an RNN unfolded in time. Recurrence between hidden states through pre-activation $a^{(t)}$ is shown in red.

Forward Propagation Through Time



Forward propagation in an RNN unfolded in time. Recurrence between hidden states through pre-activation $\mathbf{a}^{(t)}$ is shown in red.

Notational Clarity


- ▶ At layer l , we will denote the pre-activation by \mathbf{a}^l and activation by \mathbf{h}^l .
- ▶ So output layer \mathbf{y} will be denoted by \mathbf{h}^L in an L -layer network.
- ▶ Input will be denoted by \mathbf{h}^0 .
- ▶ So forward propagation entails $\mathbf{h}^0 \rightarrow \underbrace{\mathbf{a}^1 \rightarrow \mathbf{h}^1}_{\text{layer 1}} \cdots \rightarrow \underbrace{\mathbf{a}^{L-1} \rightarrow \mathbf{h}^{L-1}}_{\text{layer } L-1} \rightarrow \underbrace{\mathbf{a}^L \rightarrow \mathbf{h}^L}_{\text{layer } L}$.
- ▶ For 2 layer network

$$\mathbf{h}^{2,T} = f(\mathbf{a}^{2,T})$$

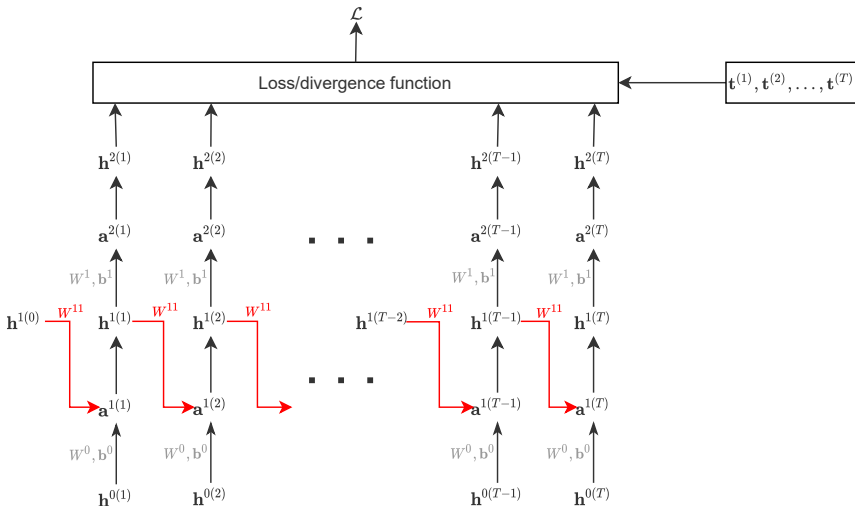
$$\mathbf{a}^{2,T} = W^1 \mathbf{h}^{1,T} + \mathbf{b}^1$$

$$\mathbf{h}^{1,T} = \tanh(\mathbf{a}^{1,T})$$


$$\mathbf{a}^{1,T} = W^0 \mathbf{h}^{0,T} + W^{11} \mathbf{h}^{1,T-1} + \mathbf{b}^0$$




Notational Clarity




RNN Variations




1 hidden state



2 hidden states




Skip connections



Bidirectional

Bidirectional RNN


Step 1: Forward propagation into the future



$$fprop \left(\underbrace{\mathbf{h}^0(1), \mathbf{h}^0(2), \dots, \mathbf{h}^0(T)}_{\text{input sequence}}; \underbrace{\mathbf{h}_f^1(0)}_{\text{init}}, \underbrace{W_f^0, \mathbf{b}_f^0, W_f^{11}, W_f^1, \mathbf{b}_f^1}_{\text{parameters}} \right)$$

Bidirectional RNN


Step 2: Forward propagation into the past



$$fprop \left(\underbrace{h^0(T), h^0(T-1), \dots, h^0(1)}_{\text{input sequence}}; \underbrace{h_b^1(T+1)}_{\text{init}}, \underbrace{W_b^0, b_b^0, W_b^{11}, W_b^1, b_b^1}_{\text{parameters}} \right)$$

Bidirectional RNN

Step 3: Fusion of forward and backward hidden states





$$h^{2(t)} = \tanh(a^{2(t)})$$

$$a^{2(t)} = W_f^1 h_f^{1(t)} + b_f^1 + W_b^1 h_b^{1(t)} + b_b^1$$

Benefit of Recurrent Architectures


n-bit addition




- ▶ Only for n -bit numbers
- ▶ Training set exponential in n
- ▶ Training errors
- ▶ Iterative application to numbers of arbitrary size
- ▶ Exact answers

Benefit of Recurrent Architectures

n-bit XOR



- ▶ Only for n -bit numbers
- ▶ Training set exponential in n
- ▶ Training errors



- ▶ Iterative application to numbers of arbitrary size
- ▶ Exact answer
- ▶ Will be 1 for odd number of ones in input.

Stability issues

- ▶ Even a 1-hidden layer RNN is a very deep network.
- ▶ Viewed in time, an RNN is as deep as the number of time steps.
- ▶ Suffers from vanishing gradients.
- ▶ Also suffers from *exploding gradients*.

Exploding Gradients

- ▶ Consider input $\mathbf{x}^{(1)}$ at time 1 and *assume linear* hidden layer.
- ▶ At time t , the RNN carries a term of the form

$$W^{11} \dots W^{11} W^0 \mathbf{x}^{(1)} = (W^{11})^{t-1} W^0 \mathbf{x}^{(1)}$$

which is an M -dimensional vector.

- ▶ Magnitude of this vector depends on largest eigenvalue λ_{\max} of W^{11} .
 - ▶ $\lambda_{\max} > 1 \implies$ magnitude of $(W^{11})^{t-1} W^0 \mathbf{x}^{(1)}$ keeps increasing.
 - ▶ $\lambda_{\max} < 1 \implies$ magnitude of $(W^{11})^{t-1} W^0 \mathbf{x}^{(1)}$ keeps decreasing.

Exploding Gradients

- ▶ Even during forward propagation, depending on the largest eigenvalue of the recurrent weight matrix W^{11} , input at time t
 - ▶ is either forgotten very soon,
 - ▶ or explodes to very large values.
- ▶ Similar case for backpropagation.
- ▶ Notice that this has nothing to do with the choice of activation function.
- ▶ *Information will explode or vanish through time.*
- ▶ Similar behaviour for non-linear neurons.
- ▶ So, in practice, RNNs do not have long-term memory. Solution: LSTM.

Summary

- ▶ Dynamic signals that change over time can be modeled by RNNs.
- ▶ RNNs can represent mappings of one-to-many, many-to-one, many-to-many, and many-to-many delayed sequences.
- ▶ Previous hidden layer output influences subsequent output.
- ▶ Forward propagation is through space as well as time.
- ▶ Forward propagation can be into the future and/or into the past.
- ▶ Recurrent connection implies really deep network.
- ▶ Information can vanish or even explode through time.