



# Parallel and Distributed Computing

## CS3006

Lecture 9

**OpenMP-II**

4th April 2022

Dr. Rana Asif Rehman

# Review of OpenMP Library Functions

## ➤ Controlling Number of Threads and Processors

- `void omp_set_num_threads (int num_threads);`
- `int omp_get_num_threads ( );`
- `int omp_get_max_threads ( );`
- `int omp_get_thread_num ( );`
- `int omp_get_num_procs ( );`
- `int omp_in_parallel ( );`

## ➤ Controlling and Monitoring Thread Creation

- `void omp_set_dynamic (int dynamic_threads);`
- `int omp_get_dynamic ( );`
- `void omp_set_nested (int nested);`
- `int omp_get_nested ( );`



# OpenMP

**#pragma omp directive [clause list]**

# OpenMP Directives



- ➡ `#pragma omp parallel`
- ➡ `#pragma omp for`

# One more thing to note

## **Difference between *omp for* and *omp parallel***

```
1 #pragma omp parallel
2 {
3   #pragma omp for
4   for (i = 0 < i < n; i++) {
        //omp for schedules/distributes iterations between the threads
5       /* body of parallel for loop */
6   }
7 }
```

**is same as**

```
1 #pragma omp parallel for
3   for (i = 0 < i < n; i++) {
4       /* body of parallel for loop */
6   }
```

# Some Useful Clauses in OpenMP


- A clause is an optional, additional component to a pragma
- **Private:** The private clause directs the compiler to make one or more variables private

```
int k=3;
#pragma omp parallel for default(shared) private(j) shared(k)
for (i = 0; i < n; i++)
    for (j = 0; j < n; j++)
        a[i][j] = MIN(a[i][j], a[i][k]+tmp) ;
```

## Comments:

- Here the private variable *j* is undefined -
  - when this parallel construct is entered
  - when this parallel construct is exited

# Some Useful Clauses in OpenMP

- 
- firstprivate:** It directs the compiler to create private variables having initial values identical to the value of the variable controlled by the master thread as the loop is entered.

```
s = complex_function();  
#pragma omp parallel for firstprivate(s) num_threads(2)  
for (i = 0 ; i < n ; i++) (  
    s = s*omp_get_thread_num();  
    printf("S is %d at thread#%d\n", s,omp_get_thread_num());  
}
```

# Some useful clauses



## lastprivate:

consider the following code

```
s = complex_function();  
#pragma omp parallel for private(j) firstprivate(s)  
for (i = 0 ; i < n ; i ++ ) (  
    s += 1  
}  
printf("s after join:%d\n",s); //undefined value
```

---



# Some Useful Clauses in OpenMP

- ➡ **lastprivate:** used to copy back to the master thread's copy of the variable, the private copy of the variable from the thread that executed the last iteration.

```
s = complex_function();
#pragma omp parallel for private(j) firstprivate(s) lastprivate(s)
for (i = 0 ; i < n ; i++) (
    s += 1;
}
printf("s after join:%d\n",s); //value of s as it was for last iteration of the loop
```


# Reduction clause

- Reductions are so common that OpenMP provides support for them
- May add reduction clause to **parallel for** pragma
- Specify reduction operation and reduction variable
- OpenMP takes care of storing partial results in private variables and combining partial results after the loop
- The reduction clause has this syntax:  
**reduction (<op> :<variable>)**

## ➤ Operators

- + Sum
- \* Product
- & Bitwise and
- | Bitwise or
- ^ Bitwise exclusive or
- && Logical and
- || Logical or

# Reduction clause



```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) reduction(+:area)
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```

# Conditional Parallelism Clause

- ➡ **if Clause:** The if clause gives us the ability to direct the compiler to insert code that determines at run-time whether the loop should be executed in parallel or not.
- ➡ **The clause has this syntax:** `if (<scalar expression> )`

```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) reduction(+:area) if(n>5000)
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```

# Scheduling Loops (a clause)

- **Scheduling** the loops means dividing number of iterations between the processes.
- Syntax of schedule clause  
**schedule** (*<type>*[ , *<chunk>* ] )
- Schedule type is required but, chunk size optional
- A chunk is a contiguous range of iterations
  - Increasing chunk size reduces scheduling overhead and may increase cache hit rate [due to operations on contiguous memory locations]
  - Decreasing chunk size allows finer balancing of workloads

# Scheduling Loops

## 1. **Static:** `schedule(static[, chunk-size])`

- Splits the iteration space into equal chunks of size `chunk-size` and assigns them to threads in a round-robin fashion.
- When no `chunk-size` is specified, the iteration space is split into as many chunks as there are threads (i.e., size of each is  $n/\text{tot.threads}$ ) and one chunk is assigned to each thread.
- Decision about work division is done before actually executing the code.
- Results in lower scheduling overhead. But, can cause load-imbalance if all processors are not of same compute-capability.

# Scheduling Loops

## 1. **Static:** `schedule(static[, chunk-size])`

Example when reducing chunk size improves load-balancing

```
#pragma omp parallel for private (j) schedule(static, 1)
for (i = 0; i < n; i++)
    for (j = i; j < n; j++)
        a[i][j] = complex_func(i, j);
```

# Scheduling Loops

## 2. **Dynamic:** `schedule(dynamic[, chunk-size])`

- The iteration space is partitioned into chunks given by chunk-size
- Initially every thread is assigned single chunk. The decision for remaining iteration chunks is done on run-time
- This means chunk is assigned to threads as they become idle.
- This takes care of the temporal imbalances resulting from static scheduling.
- If no chunk-size is specified, it defaults to a single iteration per chunk



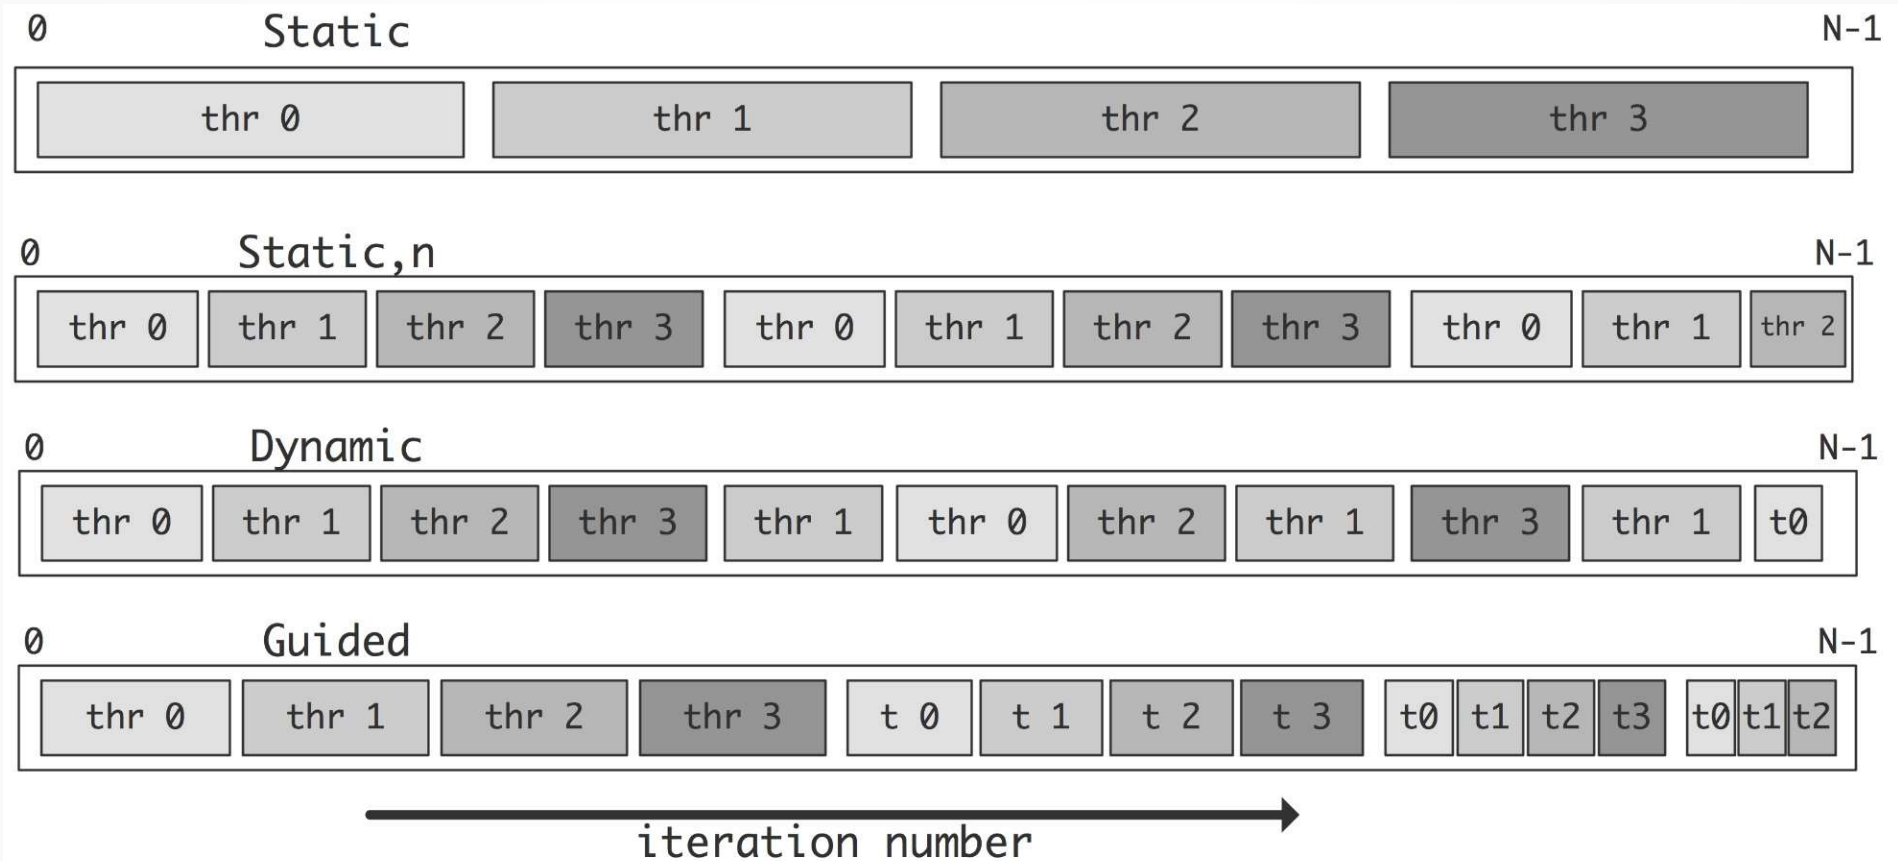
# Scheduling Loops

## 3. Guided:

- **schedule(guided, C):** dynamic allocation of chunks to tasks using guided self-scheduling heuristic. Initial chunks are bigger, later chunks are smaller, **minimum** chunk size is C.
- **schedule(guided):** guided self-scheduling with minimum chunk size 1

## 4. **schedule(runtime):** schedule chosen at run-time based on value of OMP\_SCHEDULE env variable.

# Scheduling Loops(Summary)



# No Wait Clause

- In order to avoid implicit barrier
- A thread can easily move to next after completed its assign task/iterations

**#pragma omp parallel for nowait**

# Functional / Task Parallelism in OpenMP



# Functional/Task Parallelism



**#pragma omp sections [clause list]**


# Functional/Task Parallelism

- If your code is based on different segments or sections that can be executed in parallel.
- Also known to as task parallelism

```
v = alpha();  
W = beta();  
x = gamma(v,w);  
y = delta();
```

```
printf("%6.2f\n",epsilon(x,y));
```

- Can execute alpha, beta, delta parallelly
- Remaining ones are executed sequentially according to the dependency.



# parallel sections, section pragmas

```
#pragma omp parallel sections
{
    #pragma omp section //[optional for 1st block]
    v = alpha();
    #pragma omp section
    W = beta();
    #pragma omp section
    y = delta();
}
x = gamma(v,w);
printf("%6.2f\n", epsilon(x,y));
```

- #pragma omp parallel sections creates a team of threads which executes the sections in the region parallelly
  - Sections that can be executed parallel are preceded by 'omp section' pragma.


# Functional Parallelism

## Another approach

```
v = alpha();  
W = beta();  
x = gamma(v,w);  
y = delta();
```


```
printf("%6.2f\n",epsilon(x,y));
```

- Execute alpha and beta in parallel.
- Execute gamma and delta in parallel






# *omp sections* **pragma**

- 
- Appears inside a parallel block of code
  - This pragma distributes enclosed sections among the threads in the team
  - The difference between *omp parallel sections* and *omp sections* is that,
    - ***Omp parallel sections*** generate its own team of threads
    - While simple ***omp sections*** pragma uses existing team of threads and distributes section among the threads
  - If multiple sections pragmas are inside one parallel block, may reduce fork/join costs

# sections pragma



```
#pragma omp parallel num_threads(2)
{
    #pragma omp sections
    {
        #pragma omp section //optional
        v = alpha();
        #pragma omp section
        w = beta();
    } // here an implicit barrier exists

    #pragma omp sections
    {
        x = gamma(v, w);
        #pragma omp section
        y = delta();
    }
}

printf ("%6.2f\n", epsilon(x,y));
```

# Questions



# References



1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (2017). *Introduction to parallel computing*. Redwood City, CA: Benjamin/Cummings.