

Parallel and Distributed Computing

CS3006

Lecture 14

Message Passing and MPI

11th May 2022

Dr. Rana Asif Rehman

Message Passing and MPI

Message Passing Paradigm

3

Programming Using the Message Passing Paradigm

- Oldest and most widely used approach for distributed programming.
- The logical view of a machine supporting the message-passing paradigm consists of p processes, each with its own exclusive address space.
- Most of the communication is done using simple send/receive message passing.

Characteristics

- Provides high scalability
- Complex to program
- High communication costs
- No support for incremental parallelism

Message Passing Interface (MPI)

4

- MPI defines a standard library for message-passing that can be used to develop portable message-passing programs using either C or Fortran.
- The MPI standard defines both the syntax as well as the semantics of a core set of library routines.
- It is possible to write fully-functional message-passing programs by using only the six routines.

Message Passing Interface (MPI)

5

The minimal set of MPI routines.

<code>MPI_Init</code>	Initializes MPI.
<code>MPI_Finalize</code>	Terminates MPI.
<code>MPI_Comm_size</code>	Determines the number of processes.
<code>MPI_Comm_rank</code>	Determines the label of calling process.
<code>MPI_Send</code>	Sends a message.
<code>MPI_Recv</code>	Receives a message.

Starting and Terminating the MPI Library

6

- `MPI_Init` is called prior to any calls to other MPI routines. Its purpose is to initialize the MPI environment.
- `MPI_Finalize` is called at the end of the computation, and it performs various clean-up tasks to terminate the MPI environment.
- The prototypes of these two functions are:

```
int MPI_Init(int *argc, char ***argv)
```

```
int MPI_Finalize()
```

- `MPI_Init` also strips off any MPI related command-line arguments.
- All MPI routines, data-types, and constants are prefixed by “`MPI_`”. The return code for successful completion is `MPI_SUCCESS`.

Communicators

7

- A communicator defines a *communication domain*
 - a set of processes that can communicate with each other.
- Information about communication domains is stored in variables of type `MPI_Comm`.
- Communicators are used as arguments to all message transfer MPI routines.
- A process can belong to many different (possibly overlapping) communication domains.
- MPI defines a default communicator called `MPI_COMM_WORLD` which includes all the processes.

Querying Information

8

- The `MPI_Comm_size` and `MPI_Comm_rank` functions are used to determine the number of processes and the label of the calling process, respectively.
- The calling sequences of these routines are as follows:

```
int MPI_Comm_size(MPI_Comm comm, int *size)
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- The rank of a process is an integer that ranges from zero up to the size of the communicator minus one.

Hello World Program

9

```
1. #include <mpi.h>
2. main(int argc, char *argv[])
3. {
4.     int np, myrank;
5.     MPI_Init(&argc, &argv);
6.     MPI_Comm_size(MPI_COMM_WORLD, &np);
7.     MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
8.     printf("From process %d out of %d,
           HelloWorld!\n", myrank, np);
9.     MPI_Finalize();
10. }
```

Sending and Receiving Messages

10

- The basic functions for sending and receiving messages in MPI are the `MPI_Send` and `MPI_Recv`, respectively.
- The calling sequences of these routines are as follows:

```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
int MPI_Recv(void *buf, int count, MPI_Datatype
datatype, int source, int tag, MPI_Comm comm,
MPI_Status *status)
```
- MPI provides equivalent datatypes for all C datatypes. This is done for portability reasons.
- The message-tag can take values ranging from zero up to the MPI defined constant `MPI_TAG_UB`.

MPI Datatypes

11

MPI Datatype	C Datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	

Sending and Receiving Messages

12

- MPI allows specification of **wildcard** arguments for both source and tag.
- If source is set to `MPI_ANY_SOURCE`, then any process of the communication domain can be the source of the message.
- If tag is set to `MPI_ANY_TAG`, then messages with any tag are accepted.
- On the receive side, the message must be of length equal to or less than the length field specified.

Sending and Receiving Messages

13

- On the receiving end, the status variable can be used to get information about the `MPI_Recv` operation.

- The corresponding data structure contains:

```
typedef struct MPI_Status {  
    int MPI_SOURCE;  
    int MPI_TAG;  
    int MPI_ERROR; };
```

- `MPI_Status` is usually used to take source and tag information in a 'receive' with wildcard entries on the corresponding positions.

Sending and Receiving Messages

14

Example Program

```
if(my_rank==0){
    int sendBuff=10,tag=1,dest=1;
    printf("Process:%d is sending \'%d\' to process:%d \n",my_rank, sendBuff,dest);

    MPI_Send(&sendBuff, 1, MPI_INT, dest, tag, MPI_COMM_WORLD);

}else if(my_rank==1){
    int recvBuff;int source=0,tag=1;

    MPI_Recv(&recvBuff, 1, MPI_INT, source, tag, MPI_COMM_WORLD, &status);
    printf("Process:%d is has received \'%d\' from process:%d \n",my_rank,
    recvBuff,source);

}else{
}
```

15

Questions



References

16

1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (2017). *Introduction to parallel computing*. Redwood City, CA: Benjamin/Cummings.