

# NESTED QUERIES

A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*

Retrieve the name of all employees who earn more than the Average Salary

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       Salary > (SELECT AVG(SALARY)
                     FROM      EMPLOYEE)
```

EMPLOYEE								
FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN
DEPARTMENT								
				DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE	



# NESTING OF QUERIES

A complete SELECT query, called a **nested query**, can be specified within the WHERE-clause of another query

Retrieve the name of all employees who work for the 'Research' department.

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       DNO = (SELECT DNUMBER
                   FROM   DEPARTMENT
                   WHERE  DNAME='Research' )
```

	FNAME	LNAME
1	John	Smith
2	Franklin	Wong
3	Joyce	English
4	Ramesh	Narayan

	DNUMBER
1	5

If `=` is used the inner query must return one value  
If more than one value is returned then an error msg is generated

DEPARTMENT													
EMPLOYEE		<table><tr><td>DNAME</td><td><u>DNUMBER</u></td><td>MGRSSN</td><td>MGRSTARTDATE</td></tr></table>								DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
		DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE								
FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO				



# NESTING OF QUERIES

Retrieve the name of all employees who work for the 'Research' or Administration 'department' .


```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       DNO IN (SELECT DNUMBER
                    FROM DEPARTMENT
                    WHERE DNAME='Research' OR
                        DNAME='Administration' )
```

If `=` is used the inner query must return one value.

If inner query returns more than one value then use **IN**

EMPLOYEE									
FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO

DEPARTMENT			
DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE



# NESTING OF QUERIES

Retrieve the name of all employees who do not work for the 'Research' department.

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       DNO NOT IN (SELECT DNUMBER
                        FROM DEPARTMENT
                        WHERE DNAME='Research' )
```

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

# NESTED QUERIES

You can also use:  $s > \text{ALL } R$  (means greater than every value)  
 $s > \text{ANY } R$  (means greater than any value )  
**= ANY is same as IN ,  $\neq$  ALL is same as NOT IN**

**Find name of employees whose salary is greater than the salary of all employees in department 5**

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

```
SELECT Fname
FROM Employee
WHERE Salary > ALL (SELECT Salary
                    FROM Employee
                    where Dno=5)
```

	Fname
1	James
2	Jennifer

# NESTED QUERIES & TSQL

A subquery can be nested inside the **WHERE** or **HAVING** clause of an outer SELECT, INSERT, UPDATE, or DELETE statement.

**Statements with subquery usually take one of these formats:**

- WHERE expression [NOT] IN (subquery)
- WHERE expression **comparison\_operator** [ANY | ALL] (subquery)
- WHERE [NOT] EXISTS (subquery)

**comparison\_operator** { = | <> | != | > | >= | !> | < | <= | !< }

- Up to **32 levels** of nesting is possible,
  - This limit depends on available memory and the complexity of other expressions in the query.

# Why NESTED QUERIES ?

- Many Transact-SQL statements that include subqueries can be alternatively formulated as joins.
- **Other questions can be posed only with subqueries.**
- An aggregate may not appear in the **WHERE** clause
  - unless it is in a subquery contained in a HAVING clause or a select list, and the column being aggregated is an outer reference

# CORRELATED NESTED QUERIES

- If a condition in the *nested query references* an attribute of a relation *declared in the outer query* =>
  - Then two queries are said to be correlated

Retrieve the name of each employee who has a dependent with the same first name as the employee.

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
Jones	E	Bass	888665555	1997-11-10	450 Stagg, Houston, TX	M	55000	NULL	1

Nested Correlated query is evaluated once for each tuple in outer query



# CORRELATED NESTED QUERIES

A **correlated subquery** (also called *repeating subquery*) depends on the **outer query** for its values.

- This means that the nested subquery is executed **repeatedly**, once for each row that might be selected by the **outer query**.

Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT E.FNAME, E.LNAME
FROM   EMPLOYEE AS E
WHERE  E.SSN IN (SELECT ESSN   FROM DEPENDENT
                 WHERE SSN = ESSN AND FNAME=DEPENDENT_NAME)
```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

# NESTED QUERIES

- A query written with nested SELECT... FROM... WHERE... blocks and using the = or IN comparison operators can *always* be expressed as a single block query.

**Retrieve the name of each employee who has a dependent with the same first name as the employee.**

```
SELECT      E.FNAME, E.LNAME
FROM        EMPLOYEE E, DEPENDENT D
WHERE       E.SSN=D.ESSN AND
            E.FNAME=D.DEPENDENT_NAME
```

# EXISTS FUNCTION

EXISTS Function checks whether the result of a nested query is empty or not

- Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       EXISTS (SELECT      *
                      FROM        DEPENDENT
                      WHERE       SSN=ESSN AND
                                FNAME=DEPENDENT_NAME)
```

# EXISTS FUNCTION

Retrieve the names of employees who have no dependents.

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       NOT EXISTS (SELECT *
                        FROM DEPENDENT
                        WHERE SSN=ESSN)
```

**EXISTS is necessary for the expressive power of SQL**

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4

James

The above correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple.

If *none exist* , the EMPLOYEE tuple is selected

# EXISTS FUNCTION

Find the names of managers who have at least one dependents.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE EXISTS (SELECT *
               FROM DEPARTMENT
               WHERE SSN=Mgr_SSN)

AND
EXISTS (SELECT *
        FROM DEPENDENT
        WHERE SSN=ESSN)
```

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

# EXISTS FUNCTION

Retrieve the name of each employee who works on *all* the projects controlled by department number 4.

Set theory: S1 contains S2 if  $(S2 - S1 = 0)$

S1 = set of projects of each employee

S2 = set of Dept 4 projects

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

# EXISTS FUNCTION

Retrieve the name of each employee who works on *all* the projects controlled by department number 4.

Set theory: S1 contains S2 if  $(S2 - S1 = 0)$

```
SELECT  FNAME, LNAME
FROM    EMPLOYEE
WHERE   NOT EXISTS (
        (SELECT PNUMBER
         FROM PROJECT
         WHERE  DNUM=4)
```

**EXCEPT**

```
(SELECT  PNO
 FROM    WORKS_ON
 WHERE   SSN=ESSN)
)
```

S1 = set of projects of each employee

S2 = set of dept 4 projects

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

WORKS\_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

HOW to do this in  
Relational Algebra?



## Yet another Example

Find SSN of employees who work on all the projects of Dnum= 4

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

- $PD4(Pno) \leftarrow \pi_{Pnumber} (\sigma_{DNUM=4} Project)$
- $Ssn\_Pnos \leftarrow \pi_{Essn, Pno} (Works\_on)$
- $SSNS(ssn) \leftarrow Ssn\_Pnos \text{ ??? } PD4$

WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

DIVISION

## Yet an other Example

Find SSN of employees who work on all the projects of Dnum= 4

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

PD4

Pno

10

30

SSN\_PNOS

Essn	Pno
123456789	1
123456789	2
666884444	3
453453453	1
453453453	2
333445555	2
333445555	3
333445555	10
333445555	20
999887777	30
999887777	10
987987987	10
987987987	30
987654321	30
987654321	20
888665555	20

- $PD4(Pno) \leftarrow \pi_{Pnumber} (\sigma_{DNUM=4} Project)$
- $Ssn\_Pnos \leftarrow \pi_{Essn, Pno} (Works\_on)$
- $SSNS(ssn) \leftarrow Ssn\_Pnos \div PD4$

DIVISION

# DIVISION (Binary Operation)

Division operation is applied to two relations R1 and R2

$$R1(\text{Attributes\_R1}) \div R2(\text{Attributes\_R2})$$

- where  $\text{Attributes\_R2} \subset \text{Attributes\_R1}$ .

Let **Result** =  $R1 \div R2$

$$\text{Attr\_Res} = \text{Attributes\_R1} - \text{Attributes\_R2}$$

- Attr\_Res is a set of attributes of R1 that are not the attributes of R2.

R2	
A	
a1	
a2	
a3	

Result	
B	
b1	
b4	

R1	
A	B
a1	b1
a2	b1
a3	b1
a4	b1
a1	b2
a3	b2
a2	b3
a3	b3
a4	b3
a1	b4
a2	b4
a3	b4

For a **tuple t** to appear in the result of the DIVISION, the values in t must appear in R1 in combination with *every* tuple in R2.

# Example of DIVISION

Find SSN of employees who work on all the projects that *John Smith* works on

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS\_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

SSN\_PNOS

Essn	Pno
123456789	1
123456789	2
666884444	3
453453453	1
453453453	2
333445555	2
333445555	3
333445555	10
333445555	20
999887777	30
999887777	10
987987987	10
987987987	30
987654321	30
987654321	20
888665555	20

SMITH\_PNOS

Pno
1
2

SSNS

Ssn
123456789
453453453

- $\text{Smith} \leftarrow \sigma_{\text{fname}='John' \text{ and } \text{lname}='Smith'}(\text{Employee})$
- $\text{Smith\_Pnos} \leftarrow \pi_{\text{Pno}}(\text{Works\_on} \bowtie_{\text{essn=ssn}} \text{Smith})$
- $\text{Ssn\_Pnos} \leftarrow \pi_{\text{Essn}, \text{Pno}}(\text{Works\_on})$
- $\text{SSNS}(\text{ssn}) \leftarrow \text{Ssn\_Pnos} \div \text{Smith\_Pnos}$

# Examples of Queries in RA

Find the **names** of employees who work on *all* the projects controlled by department number 5.

$$T1(Pno) \leftarrow \pi_{Pnumber} (\sigma_{Dnum=5} (Project))$$
$$T2 \leftarrow \pi_{Essn, Pno} (Work\_On)$$
$$T3 \leftarrow (T2 \div T1)$$
$$R \leftarrow \pi_{LNAME, FNAME} (T3 * Employee)$$

**PROJECT**

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

Essn	Pno
123456789	1
123456789	2
666884444	3
453453453	1
453453453	2
333445555	2
333445555	3
333445555	10
333445555	20
999887777	30
999887777	10
987987987	10
987987987	30
987654321	30
987654321	20
888665555	20