


What is DB Transactions?

Transaction is a process involving database queries and/or modification.



Why Transactions?

Database systems are normally being accessed by many users or processes at the same time

Example-ATM

You and your Friend each take Rs10000 from different ATM's at about the same time.

DBMS must make sure none of the account deduction is lost.



ACID TRANSACTIONS

Atomic

- Whole transaction or none is done.

Consistent

- Database constraints preserved.

Isolated

- It appears to the user as if only one process executes at a time.

Durable


- Effects of a process survive a crash.

Optional: weaker forms of transactions are often supported as well.

Example of *Fund Transfer*


- Transaction to transfer \$50 from account **A** to account **B**:

1. **read(A)**
2. **$A := A - 50$**
3. **write(A)**
4. **read(B)**
5. **$B := B + 50$**
6. **write(B)**




Atomicity requirement :

- if the transaction **fails** after step 3 and before step 6,
 - the system should **ensure** that :
 - its **updates** are *not reflected* in the database,
 - else an *inconsistency* will result.



Example of *Fund Transfer*

- Transaction to transfer \$50 from account **A** to account **B**:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Consistency requirement :
 - the sum of **A** and **B** is:
 - unchanged** by the execution of the transaction.



Example of *Fund Transfer*

- Transaction to transfer \$50 from account **A** to account **B**:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Isolation requirement —
 - if between steps 3 and 6, another transaction is allowed to access the **partially updated database**,
 - it will see an **inconsistent database** (the sum $A + B$ will be less than it should be).
 - Isolation can be **ensured** trivially by:
 - running transactions **serially**, that is **one** after the **other**.
 - However*, executing multiple transactions **concurrently** has significant **benefits**.




Example of *Fund Transfer*

- Transaction to transfer \$50 from account **A** to account **B**:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

- Durability requirement :

- once the user has been notified that the transaction has **completed** :
 - (i.e., the transfer of the \$50 has taken place),
 - the **updates** to the database by the transaction **must persist**
 - despite *failures*.



T-SQL AND Transactions

SQL has following transaction modes.

- **Autocommit** transactions
 - Each individual SQL statement = transaction.
- Explicit transactions
 - BEGIN TRANSACTION**
 - [SQL statements]
 - COMMIT or ROLLBACK**



Transaction Support in TSQL

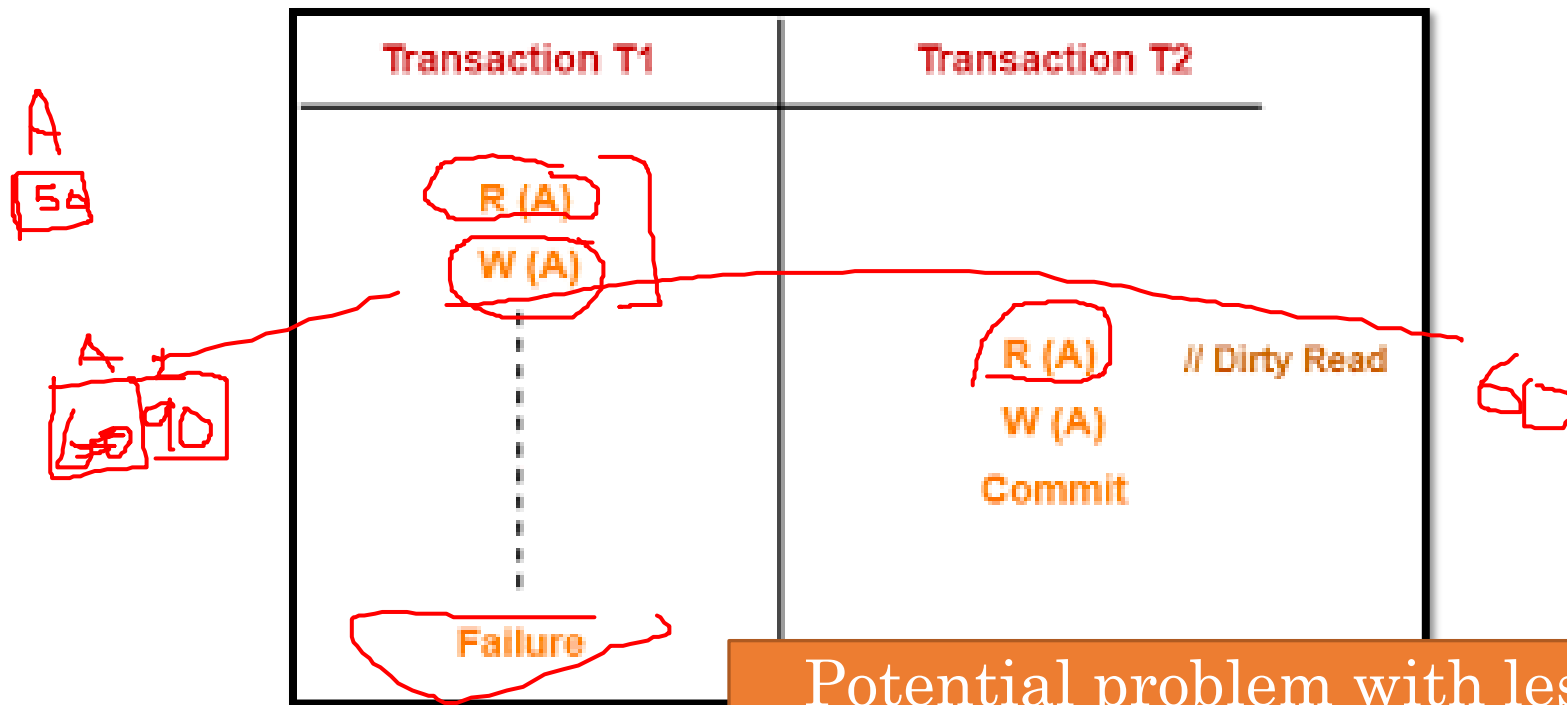
- **BEGIN TRAN**
- UPDATE Department
- SET Mgr_ssn = 123456789
- WHERE DNumber = 1
- UPDATE Department
- SET Mgr_start_date = '1981-06-19'
- WHERE Dnumber = 1
- **COMMIT TRAN**



Transaction Support in SQL

Dirty Read:

Reading a value that was written by an uncommitted transaction.

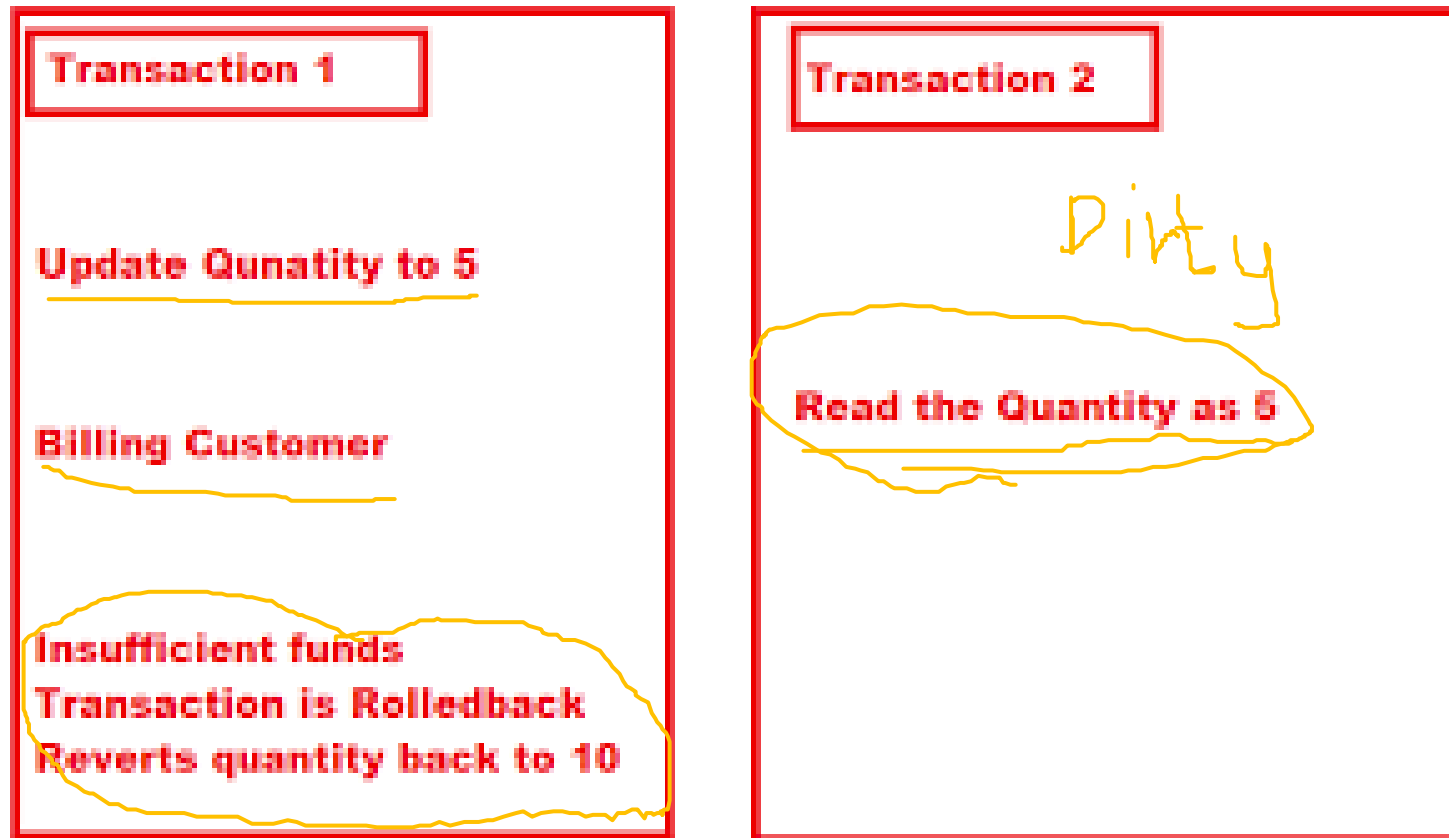


Potential problem with less restricted isolation levels

Transaction Support in SQL

Dirty Read:

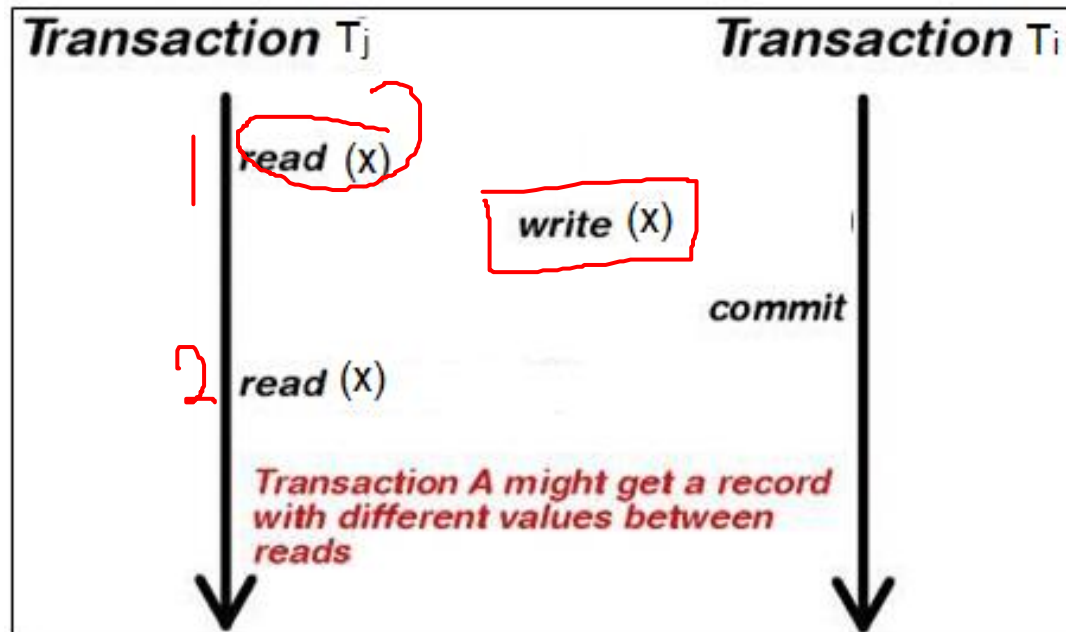
Reading a value that was written by an uncommitted transaction.




Transaction Support in SQL

○ Nonrepeatable Read

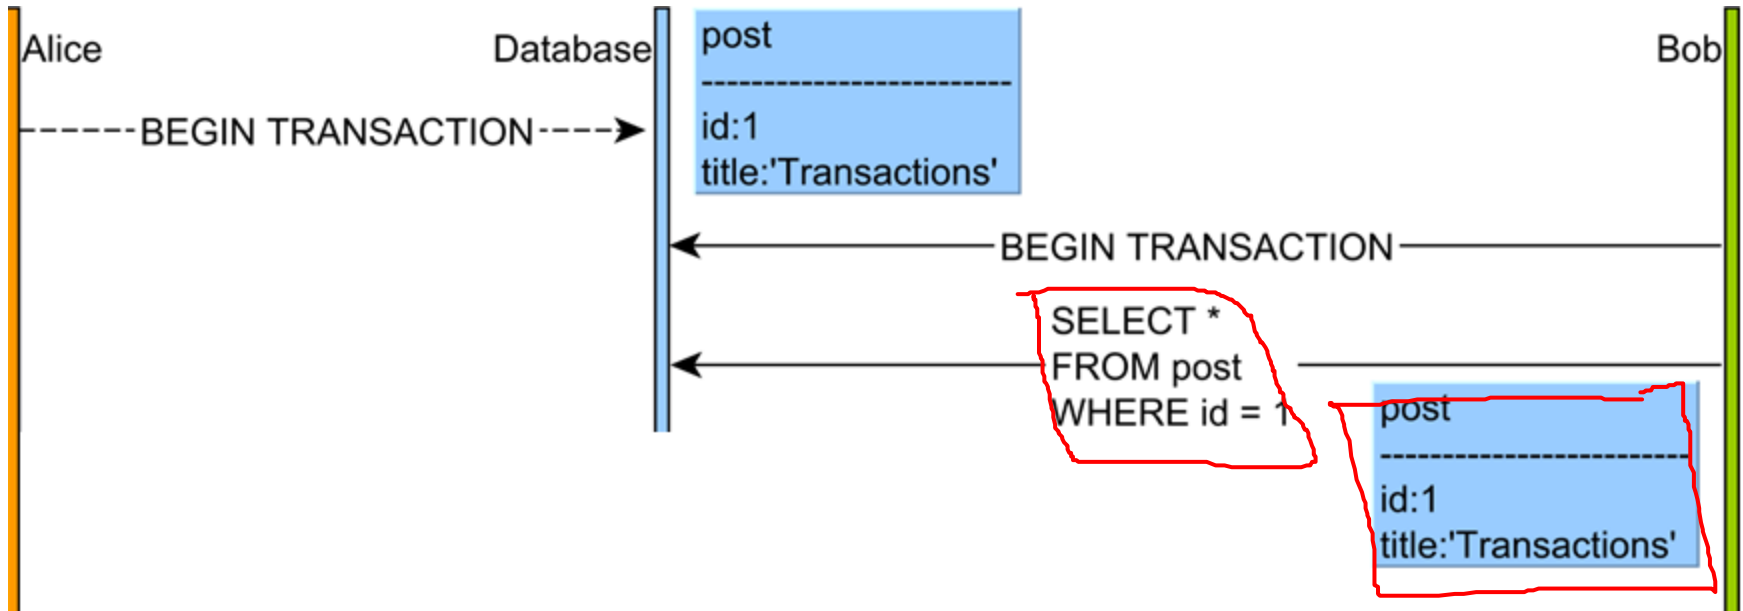
- A transaction T_i write a new value between multiple reads of transaction T_j .
 - T_j reads a given value X from a table.
 - T_i later updates X and T_j reads that value again, T_j will see a different value.



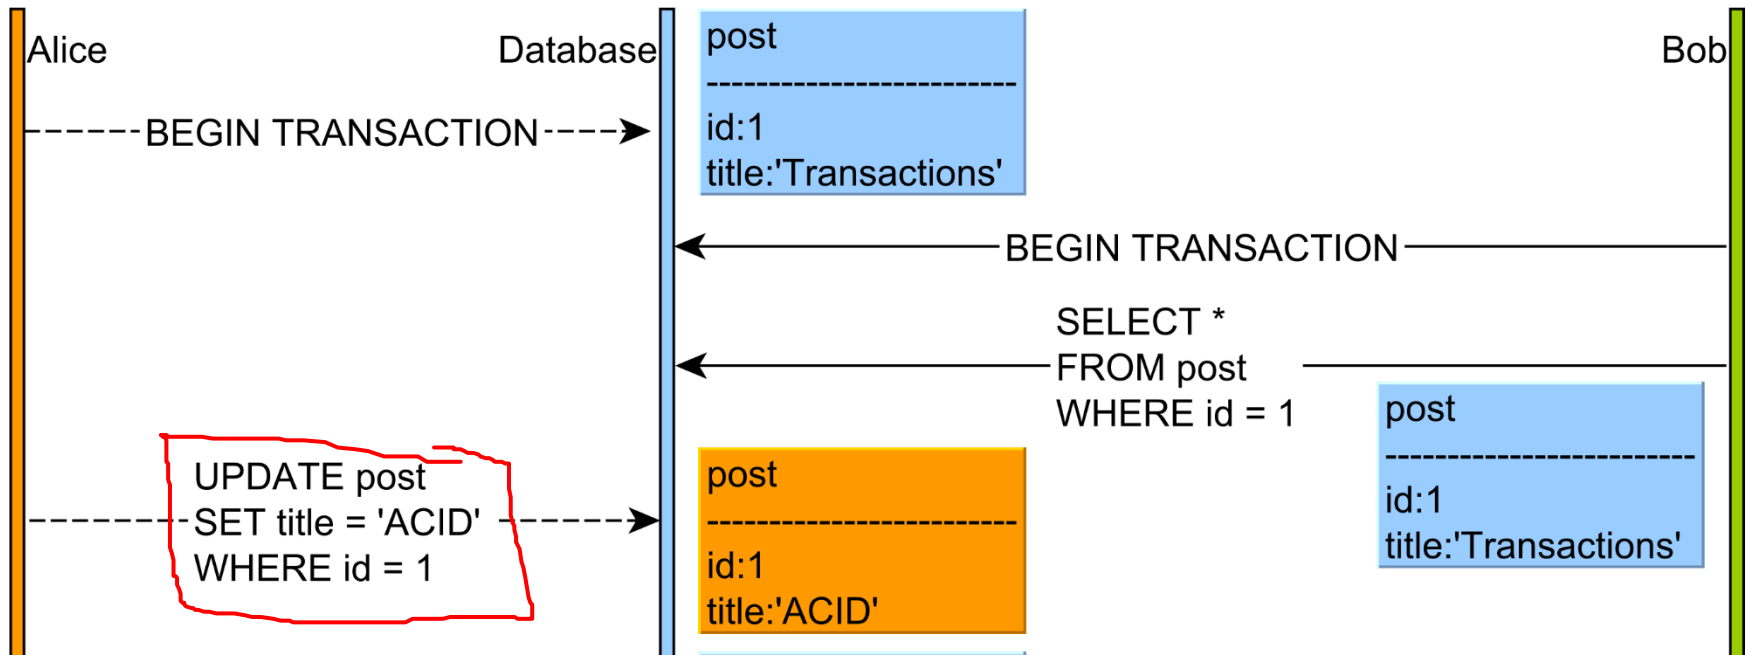
Nonrepeatable Read



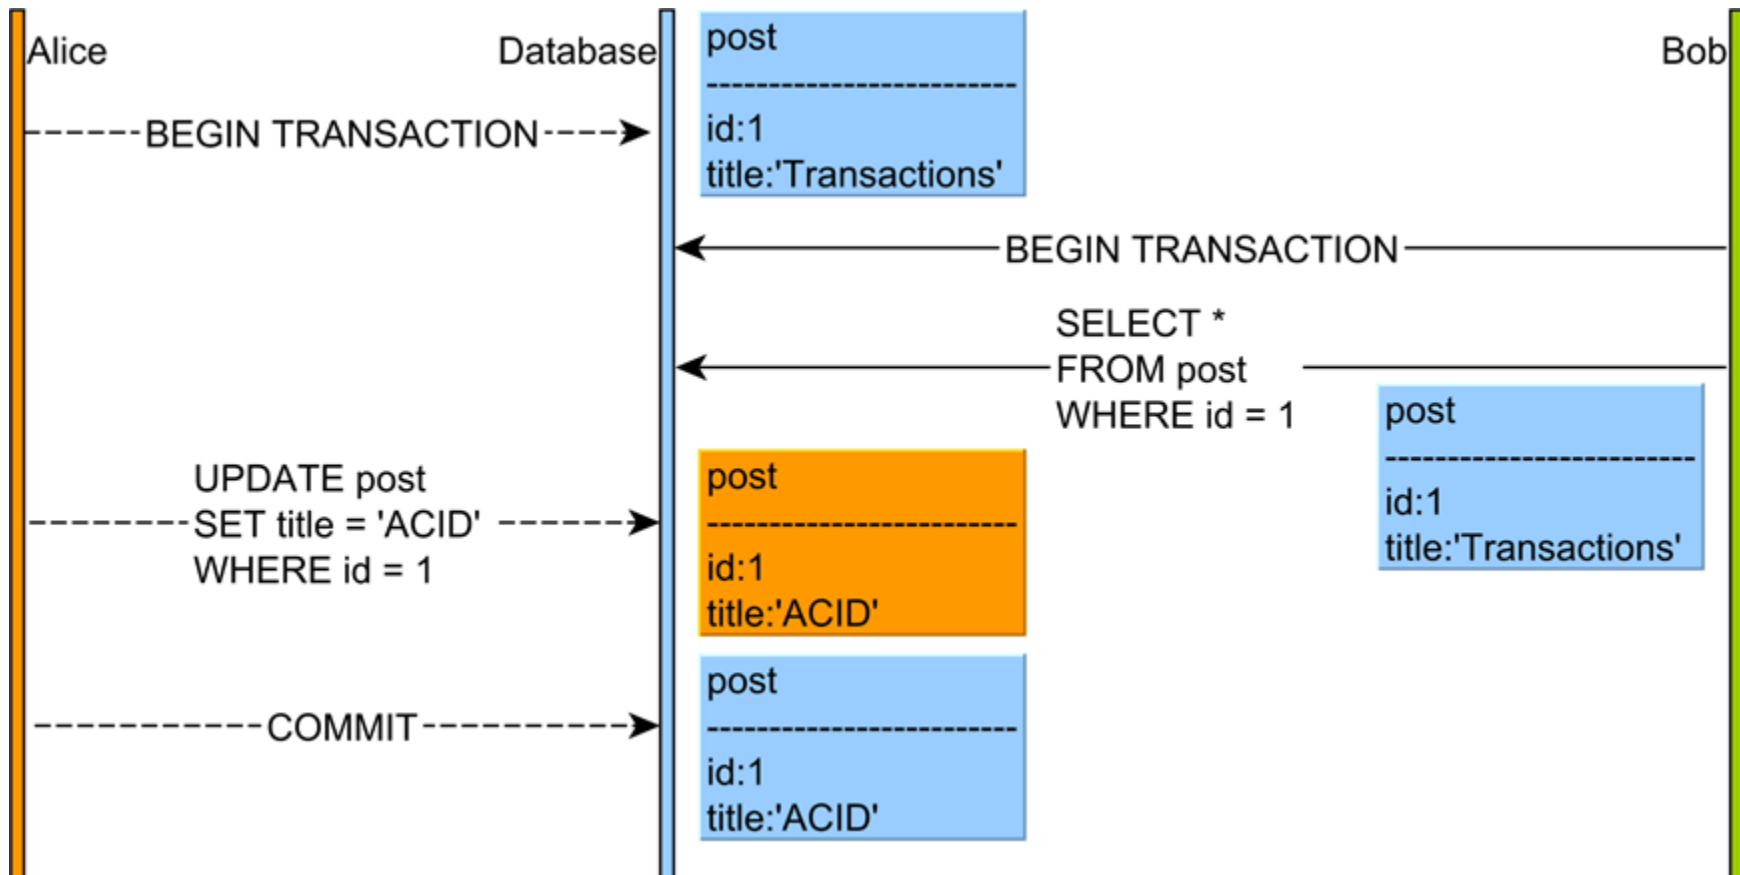
Nonrepeatable Read



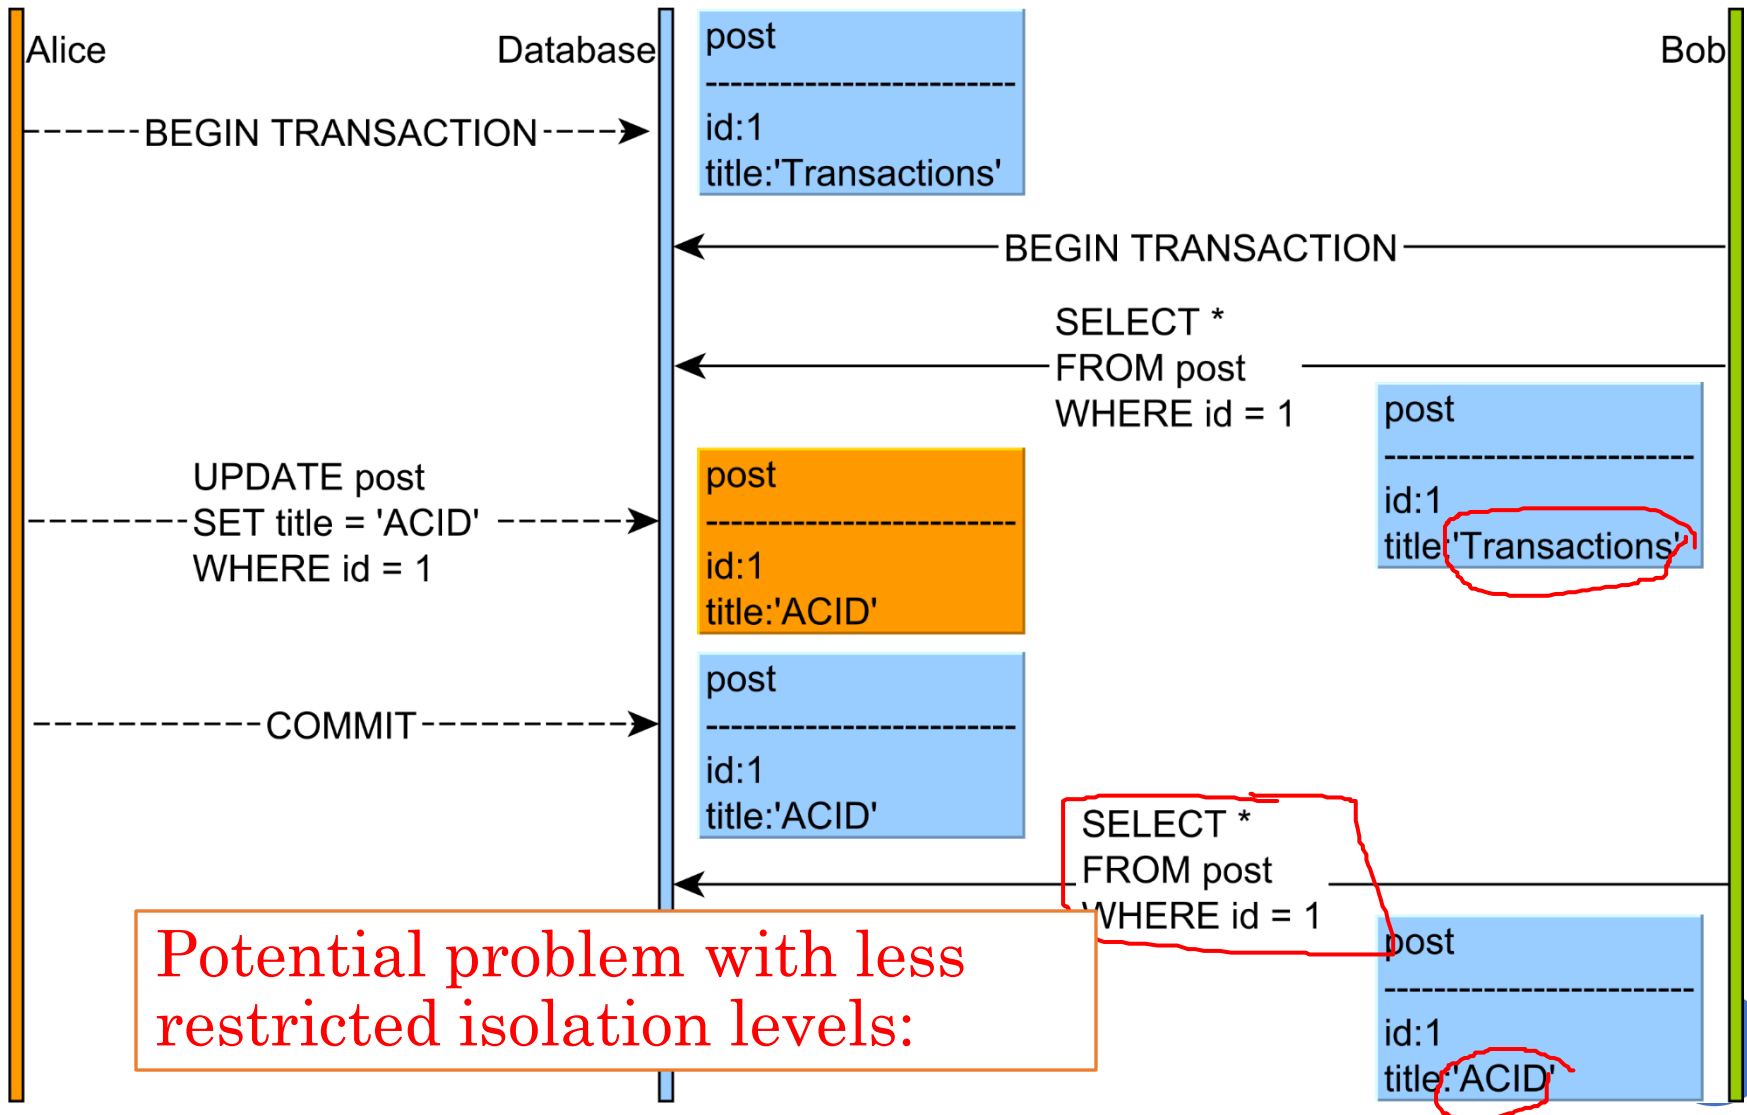
Nonrepeatable Read



Nonrepeatable Read



Nonrepeatable Read




Transaction Support in SQL

○ Phantoms

- New rows being read using the same read with a condition.

Potential problem with less restricted isolation levels



Phantom

Alice

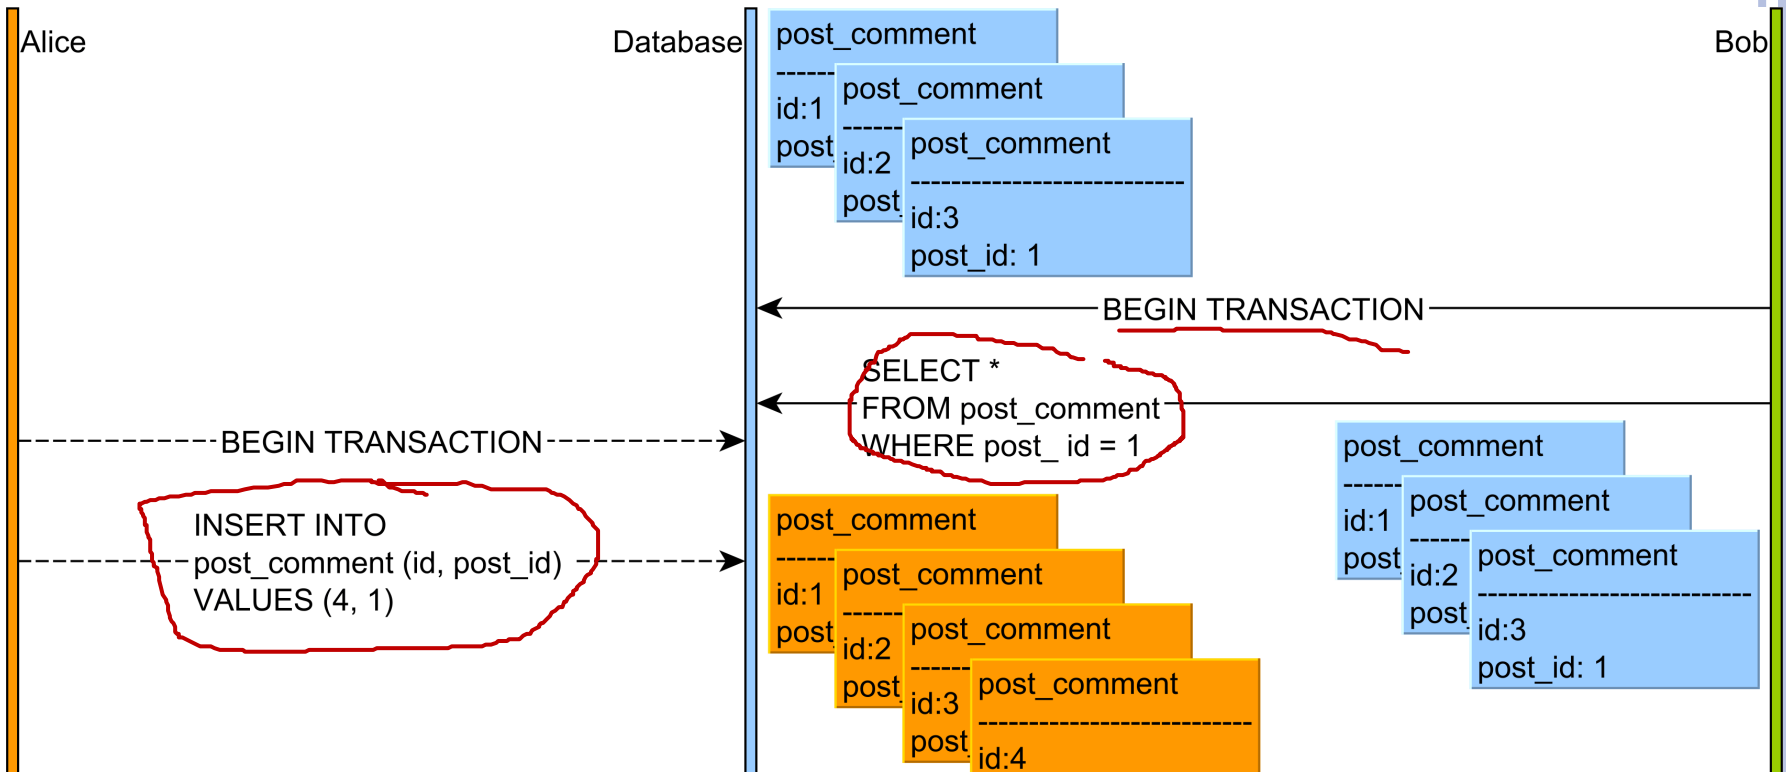
Database



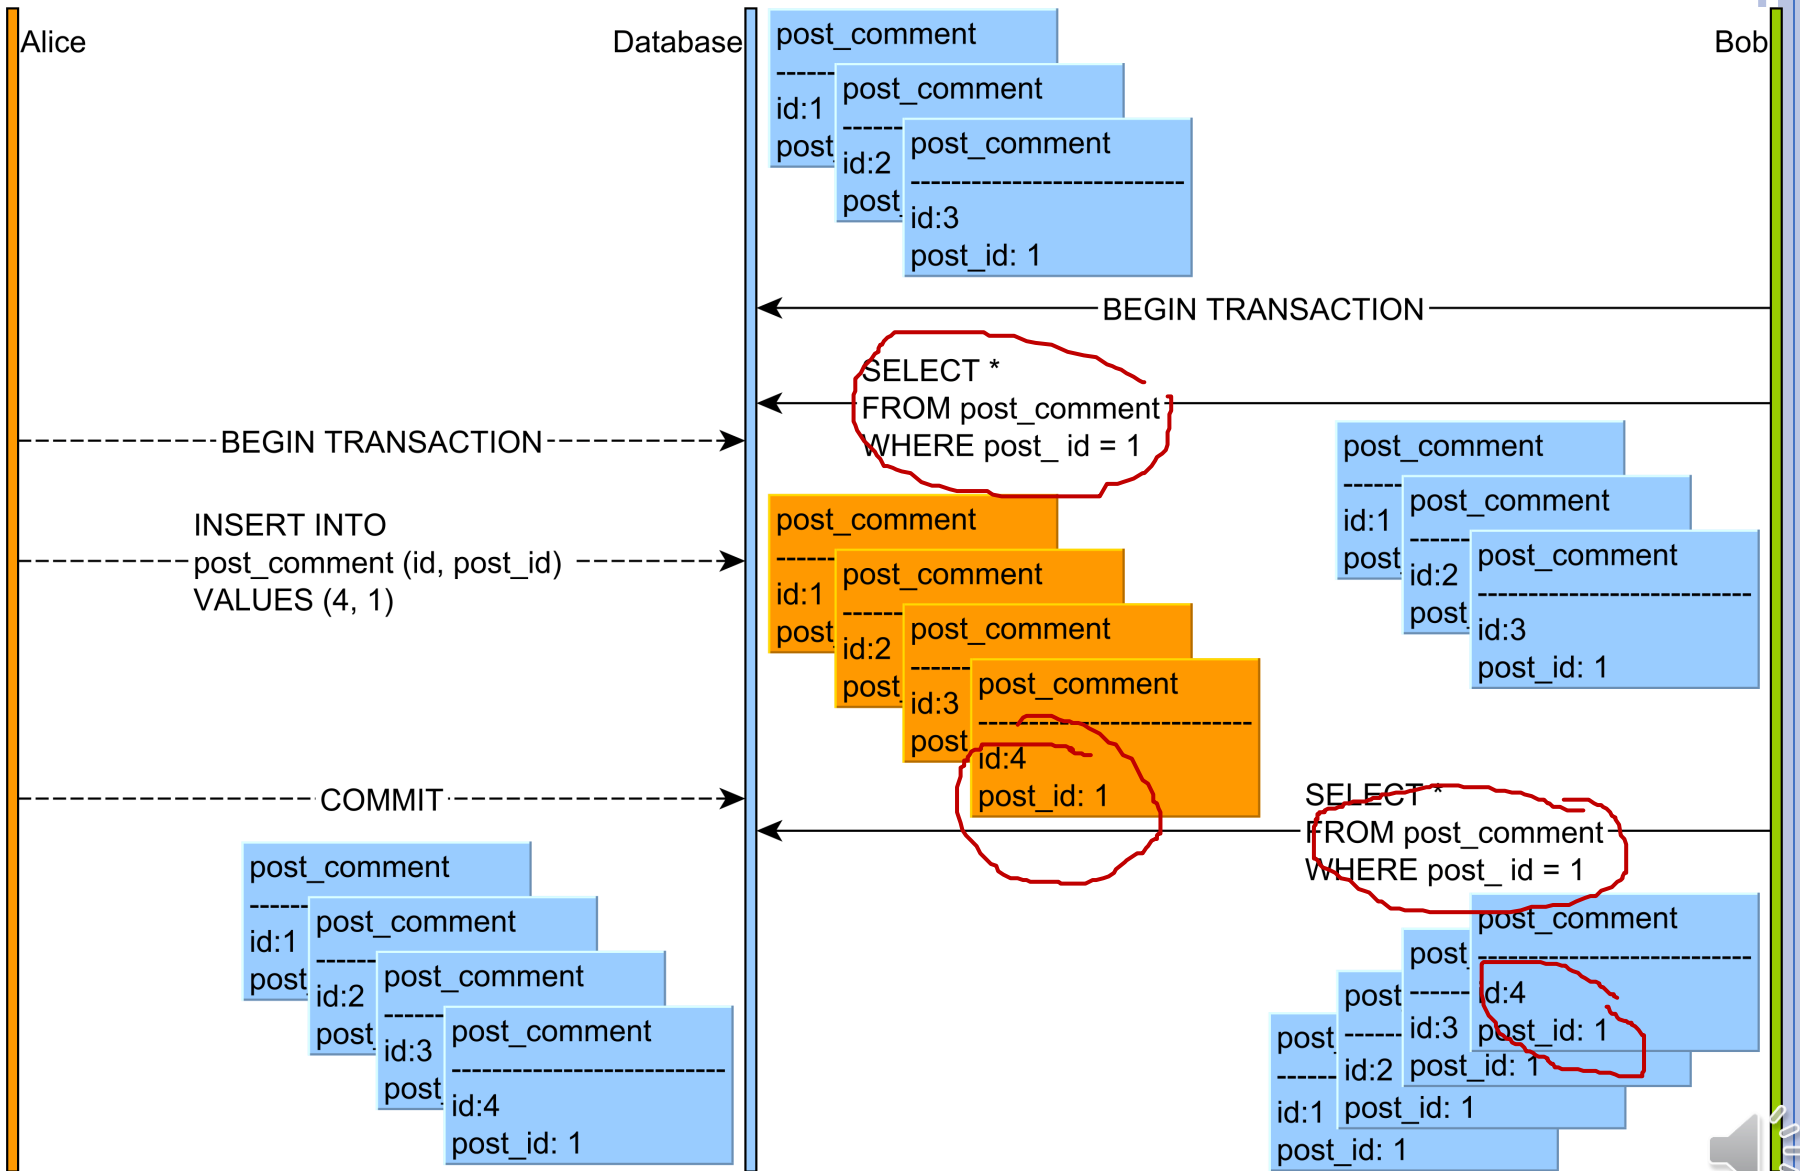
Bob



Phantom



Phantom



Transaction Support in SQL

- Potential problem with lower isolation levels (contd.):
 - **Phantoms**
 - New rows being read using the same read with a condition.
 - **Example:**
 - T_1 read a set of rows from a table,
 - on some condition specified in the SQL WHERE clause.
 - T_2 inserts a new row that also satisfies the WHERE clause condition of T_1 , into the table used by T_1 .
 - If T_1 is repeated, then T_1 will see a row that previously did not exist, called a **phantom**.



Transaction Support in TSQL

Table 21.1 Possible Violations Based on Isolation Levels as Defined in SQL

Isolation Level	Type of Violation		
	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED			
READ COMMITTED			
REPEATABLE READ			
SERIALIZABLE	No	No	No



Transaction Support in TSQL

Table 21.1 Possible Violations Based on Isolation Levels as Defined in SQL

Isolation Level	Type of Violation		
	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED			
READ COMMITTED			
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No



Transaction Support in TSQL

Table 21.1 Possible Violations Based on Isolation Levels as Defined in SQL

Isolation Level	Type of Violation		
	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED			
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No



Transaction Support in TSQL

Table 21.1 Possible Violations Based on Isolation Levels as Defined in SQL

Isolation Level	Type of Violation		
	Dirty Read	Nonrepeatable Read	Phantom
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No




TRANSACTION SUPPORT IN TSQL

1. “Dirty reads”
SET TRANSACTION ISOLATION LEVEL READ
UNCOMMITTED
2. “Committed reads”
SET TRANSACTION ISOLATION LEVEL READ
COMMITTED
3. “Repeatable reads”
SET TRANSACTION ISOLATION LEVEL
REPEATABLE READ
4. Serializable transactions (default):
SET TRANSACTION ISOLATION LEVEL
SERIALIZABLE



TRANSACTIONS EXAMPLE

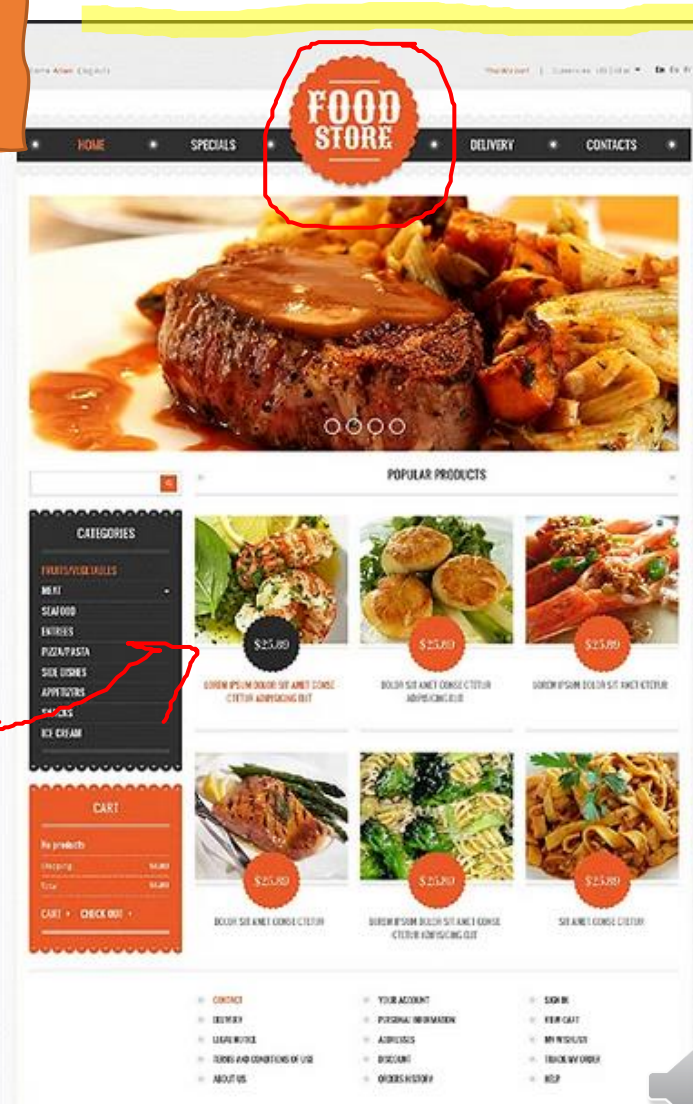


TRANSACTIONS EXAMPLE

WEB FOR FOOD COURT

Relation

Sells(shop, item, price)



TRANSACTIONS EXAMPLE

WEB FOR FOOD COURT

Relation

Sells(shop, item, price)

SCENARIO

Ahmad has a shop in
Food Store

Ahmad sells Pizza slice for Rs
200 and Sprite for Rs 50.



TRANSACTIONS EXAMPLE

WEB FOR FOOD COURT

Relation

Sells(shop, item, price)

SCENARIO

Tania is exploring
Food Store's Website

Tania visits Ahmad's Shop Page

Tania checks highest & lowest
price Ahmad charges.



TRANSACTIONS EXAMPLE

WEB FOR FOOD COURT

Relation

Sells(shop, item, price)

SCENARIO

Tania is exploring
Food Store's Website

Tania visits Ahmad's Shop Page

Tania query **Sells** for the highest &
lowest price Ahmad charges



TRANSACTIONS EXAMPLE

WEB FOR FOOD COURT

Relation

Sells(shop, item, price)

SCENARIO

Ahmad shops is not performing well

Ahmad decides to stop selling Pizza slice for Rs 200 and sell Biryani instead !!! For 250Rs



Customer's Query

Tania(customer) executes the following two SQL statements called **(min)** and **(max)**.

```
(max) SELECT MAX(price) FROM Sells  
        WHERE shop = 'Ahmad's shop'
```

```
(min) SELECT MIN(price) FROM Sells  
        WHERE shop = 'Ahmad's shop'
```



Shop Keeper's Query

At about the same time, Ahmad executes the following steps: (del) and (ins).

(del) DELETE FROM Sells
WHERE shop = 'Ahmad's shop'

(ins) INSERT INTO Sells
VALUES('Ahmad's shop', 'Biryani', 250.00)



Interleaving of Statements

- The statement (max) must come before (min), and
- The statement (del) must come before (ins),



Interleaving of Statements

- The statement (**max**) must come before (**min**), and
- The statement (**del**) must come before (**ins**),
- There are no other constraints on the order of these statements.
- Unless we group Tania's and/or Ahmad's statements into transactions.



Example: Strange Interleaving

- Suppose the steps execute in the order
(max)(del)(ins)(min).

Ahmad's Prices:	{50,200}	{50,200}	{250}	{250}
Statement:	(max)	(del)	(ins)	(min)
Result:	200			250



Example: Strange Interleaving

- Suppose the steps execute in the order
(max)(del)(ins)(min).

Ahmad's Prices: {50,200} {50,200} {250} {250}

Statement: (max) (del) (ins) (min)

Result: 200 250

- Tania sees **MAX < MIN!**



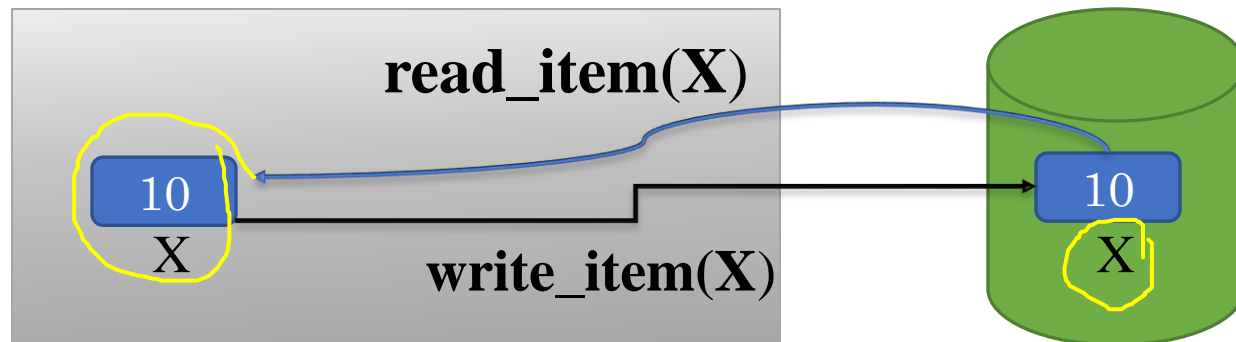
Fixing the Problem by Using Transactions

- **Solution:** Group Tania's statements (max)(min) into one transaction
 - Now, she cannot see this inconsistency.
- She sees Ahmad's prices at some fixed time.
 - Either before or after he changes prices, or in the middle, but the MAX and MIN are computed from the same prices.



Transaction Processing

- Basic operations in a DB are **read** and **write**
 - read_item(X):**
 - Reads a database item named X into a program variable.
 - To simplify our notation, we assume that the program variable is also named X.
 - write_item(X):**
 - Writes the value of program variable X into the database item named X.




Sample Transactions

(a) T_1

```
read_item (X);  
 $X := X - N$ ;  
write_item (X);  
read_item (Y);  
 $Y := Y + N$ ;  
write_item (Y);
```

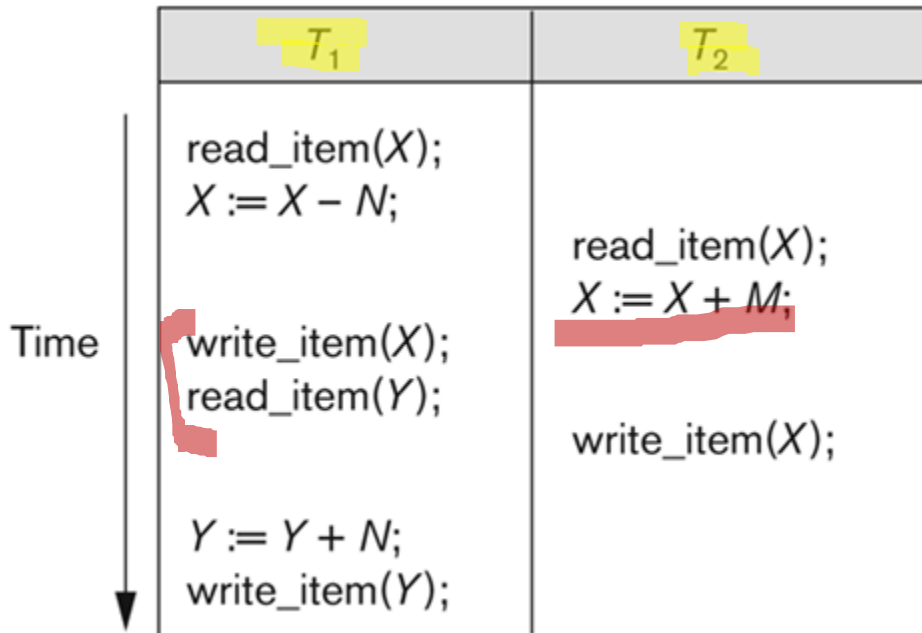
(b) T_2

```
read_item (X);  
 $X := X + M$ ;  
write_item (X);
```



Issues in Transaction Processing

Why Concurrency Control is needed?

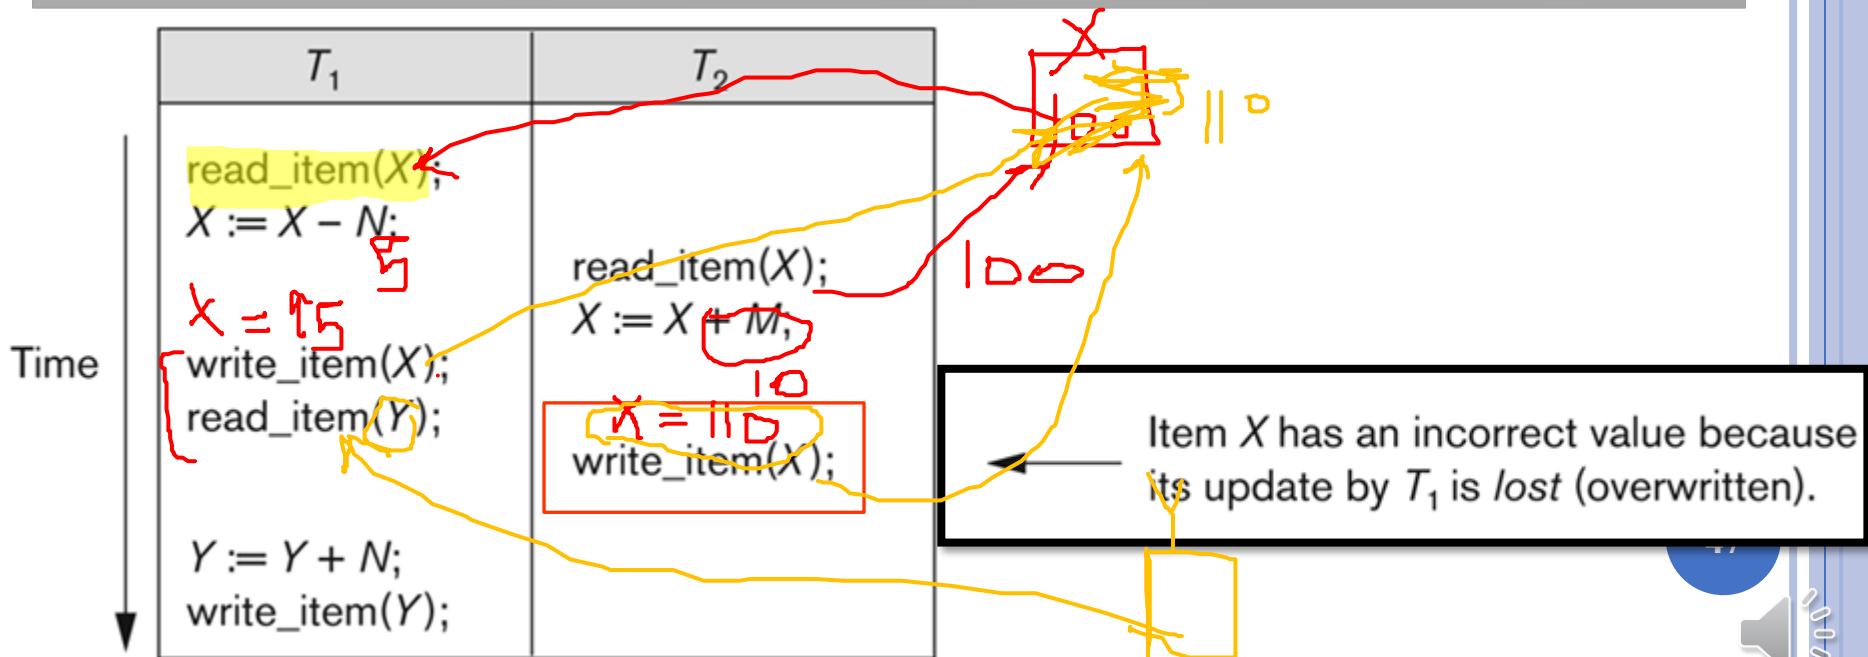


Issues in Transaction Processing

Why Concurrency Control is needed?


• The Lost Update Problem

- Two transactions (that access the same **DB items**) have their operations interleaved in a way that makes the value of some database item incorrect.



Issues in Transaction Processing


Temporary Update (or Dirty Read) Problem



Issues in Transaction Processing

The Incorrect Summary Problem

T_1	T_3
<pre> read_item(X); X := X - N; write_item(X); </pre>	<pre> sum := 0; read_item(A); sum := sum + A; ⋮ </pre>
<pre> read_item(Y); Y := Y + N; write_item(Y); </pre>	<pre> read_item(X); sum := sum + X; read_item(Y); sum := sum + Y; </pre>



← T_3 reads X after N is subtracted and reads Y before N is added; a wrong summary is the result (off by N).



STATE TRANSITION DIAGRAM

Recovery manager keeps track of the following operations

- **Begin_transaction**
- **Read or Write**
- **End_transaction**
- **Commit**
- **Rollback (or abort)**
- **Undo**
- **Redo**

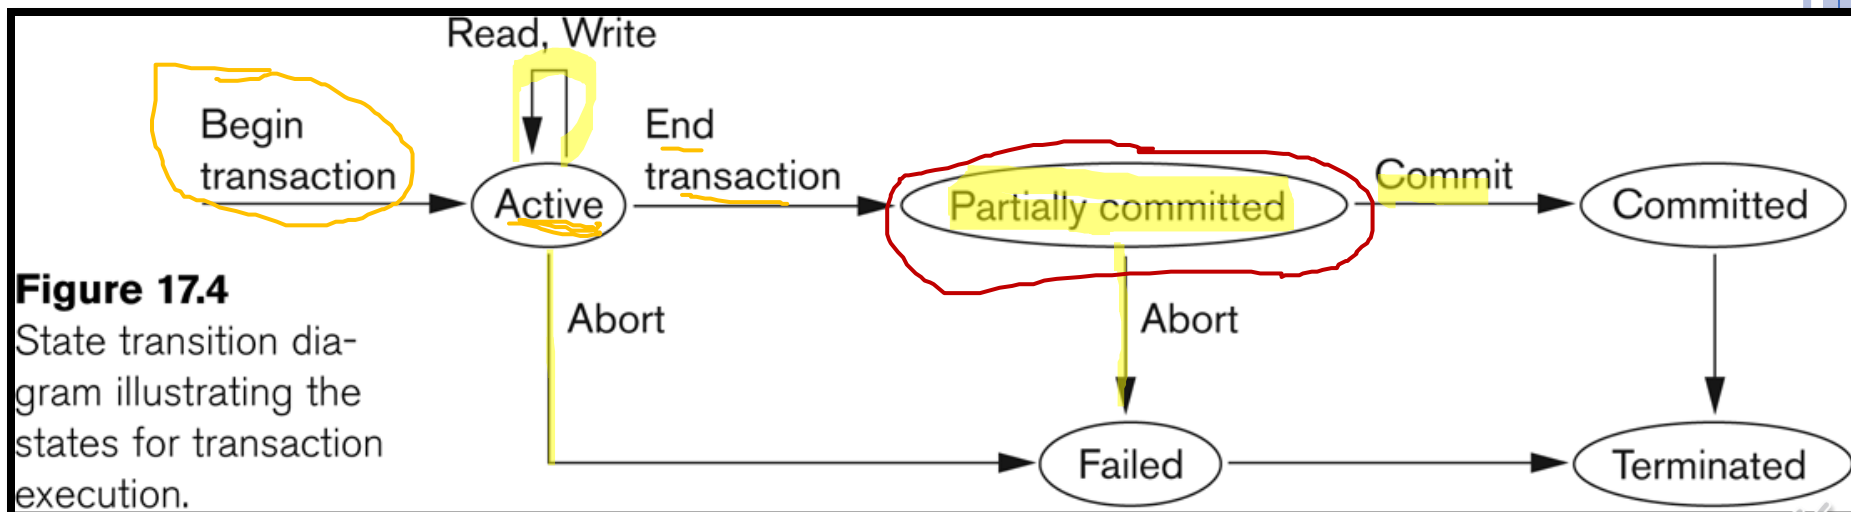


Figure 17.4

State transition diagram illustrating the states for transaction execution.



Transaction and System Concepts

A **transaction** is an atomic unit of work that is either completed in its entirety or not done at all.

For recovery purposes, the system needs to keep track of when the transaction starts, terminates, and commits or aborts

Transaction states:

- Active state
- Partially committed state
- Committed state
- Failed state
- Terminated State



Transaction Processing & Recovery

- **Why recovery is needed:**
 - A computer failure (system crash):
 - A transaction or system error:
 - Integer overflow, division by zero, erroneous parameter values or the user may interrupt the transaction
 - Local errors or exception conditions
 - Data not found or
 - insufficient account balance may cause a fund withdrawal transaction to be canceled.
 - Concurrency control enforcement
 - Transaction violates serializability or several transactions are in a state of deadlock
 - Disk failure
 - Physical problems and catastrophes



TRANSACTIONS

