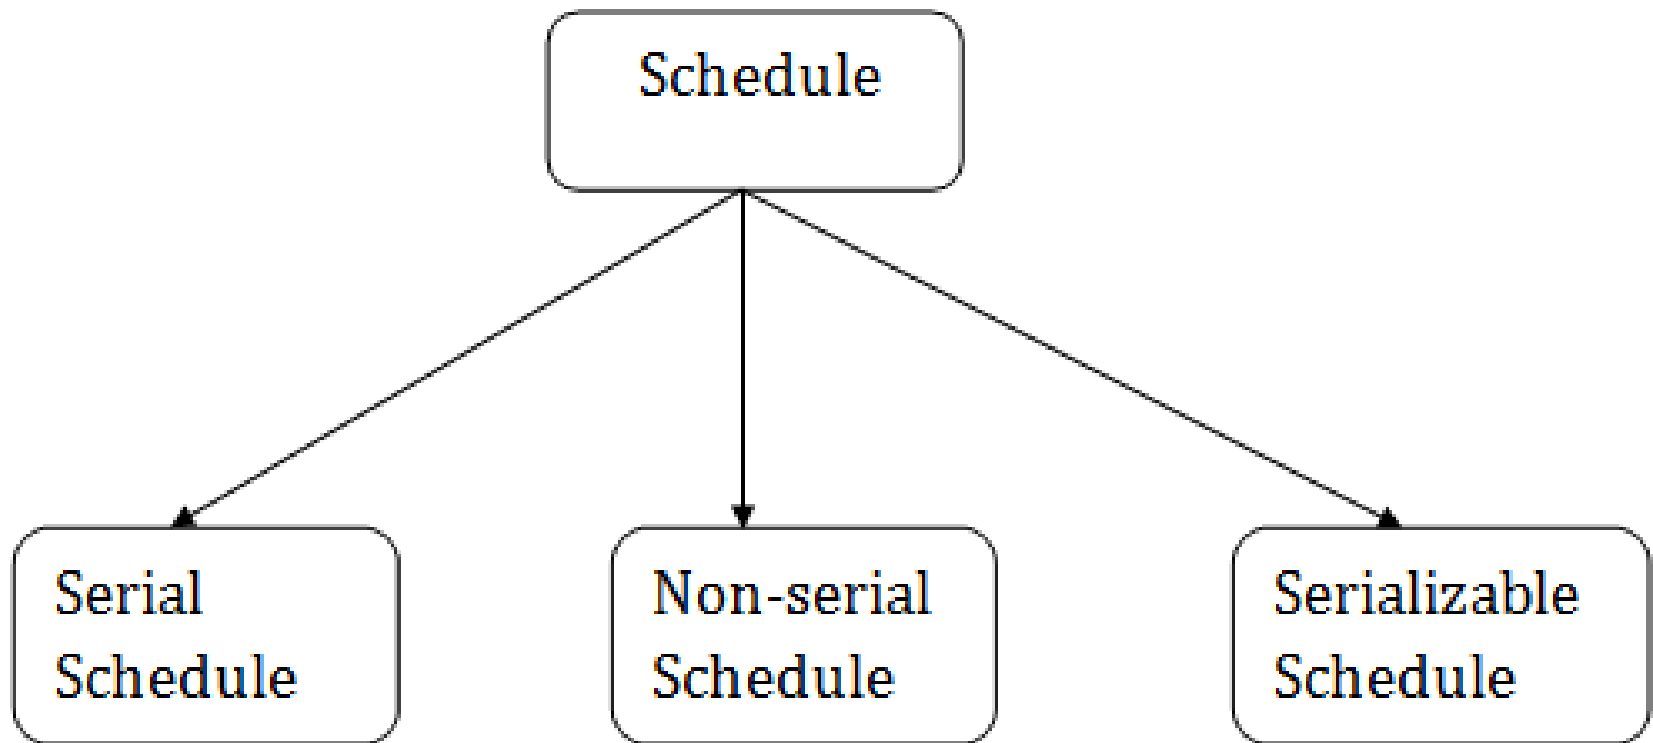




# CHARACTERIZING SCHEDULES BASED ON **SERIALIZABILITY**



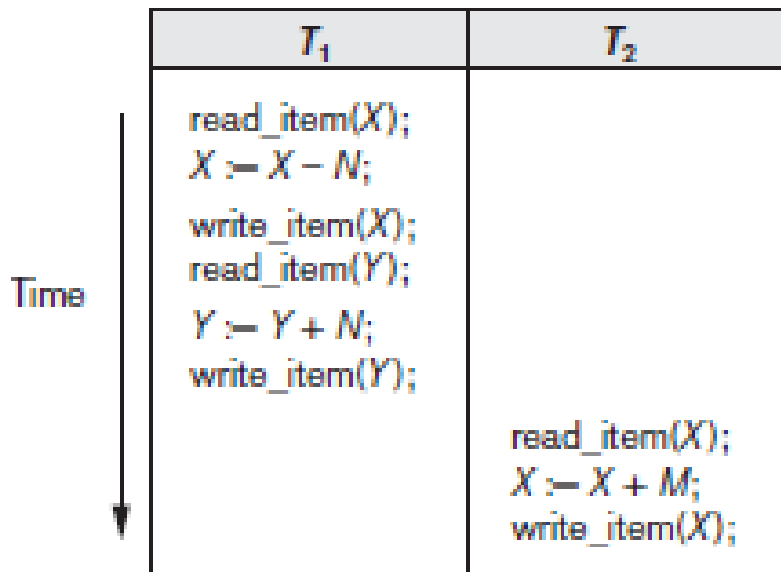
# Characterizing Schedules based on Serializability



# Schedules

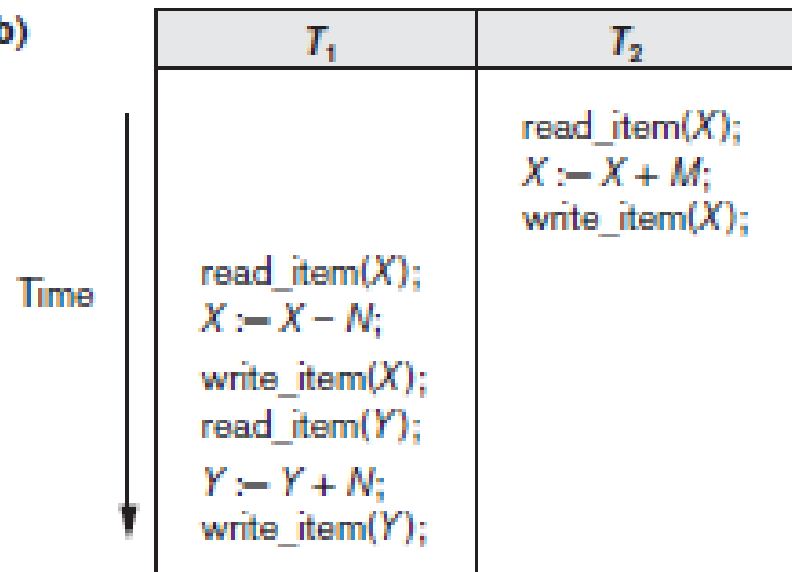
## Serial schedule

- A schedule  $S$  is serial if, for every  $T$  in  $S$ , all the operations of  $T$  are executed consecutively in the schedule.



Schedule A

(b)



Schedule B

# Schedules

## Non-Serial Schedule

- A schedule  $S$  is non-serial if the operations from different transactions are interleaved in  $S$ .

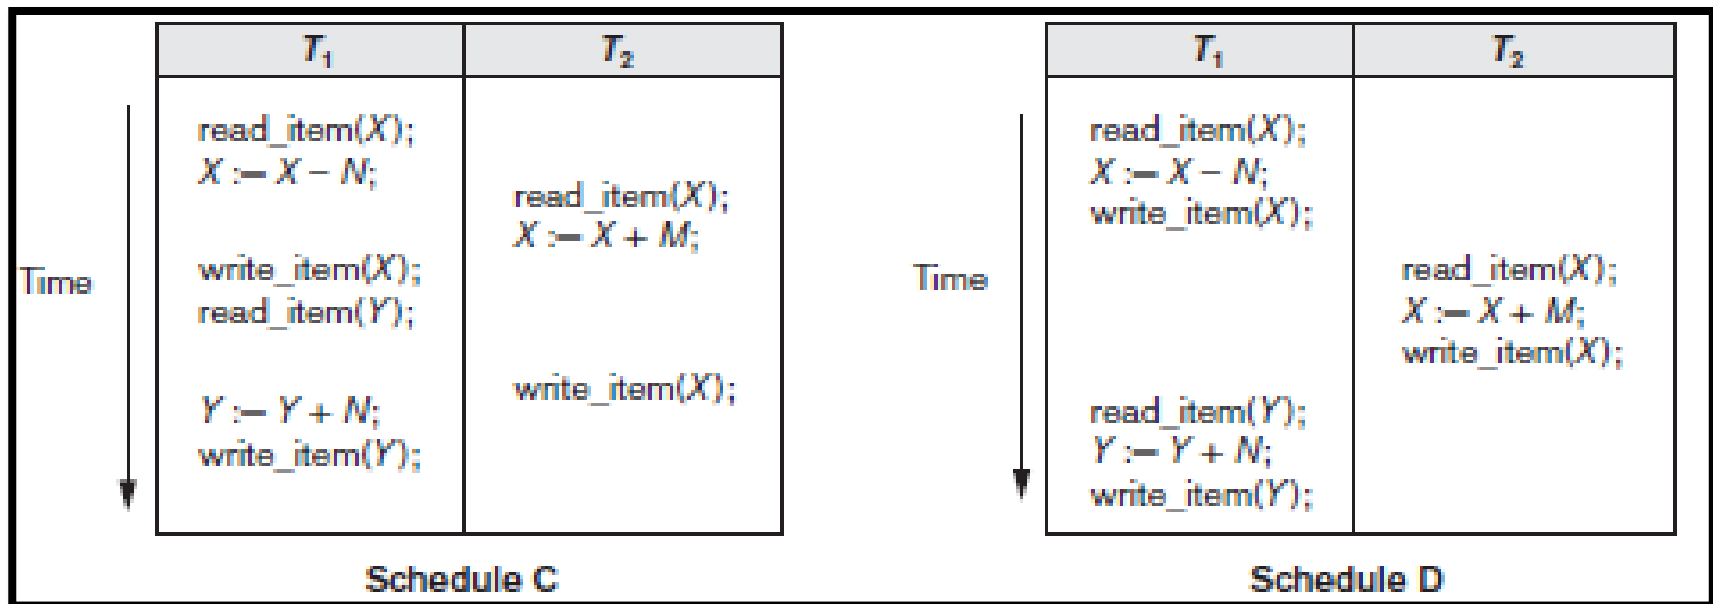
|           | $T_1$                                   | $T_2$                               |
|-----------|---|-------------------------------------|
| Time<br>↓ | read_item( $X$ );<br>$X := X - N$ ;     |                                     |
|           | write_item( $X$ );<br>read_item( $Y$ ); | read_item( $X$ );<br>$X := X + M$ ; |
|           | $Y := Y + N$ ;<br>write_item( $Y$ );    | write_item( $X$ );                  |

The operations of each  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$ .

$S_a: r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

# Schedules

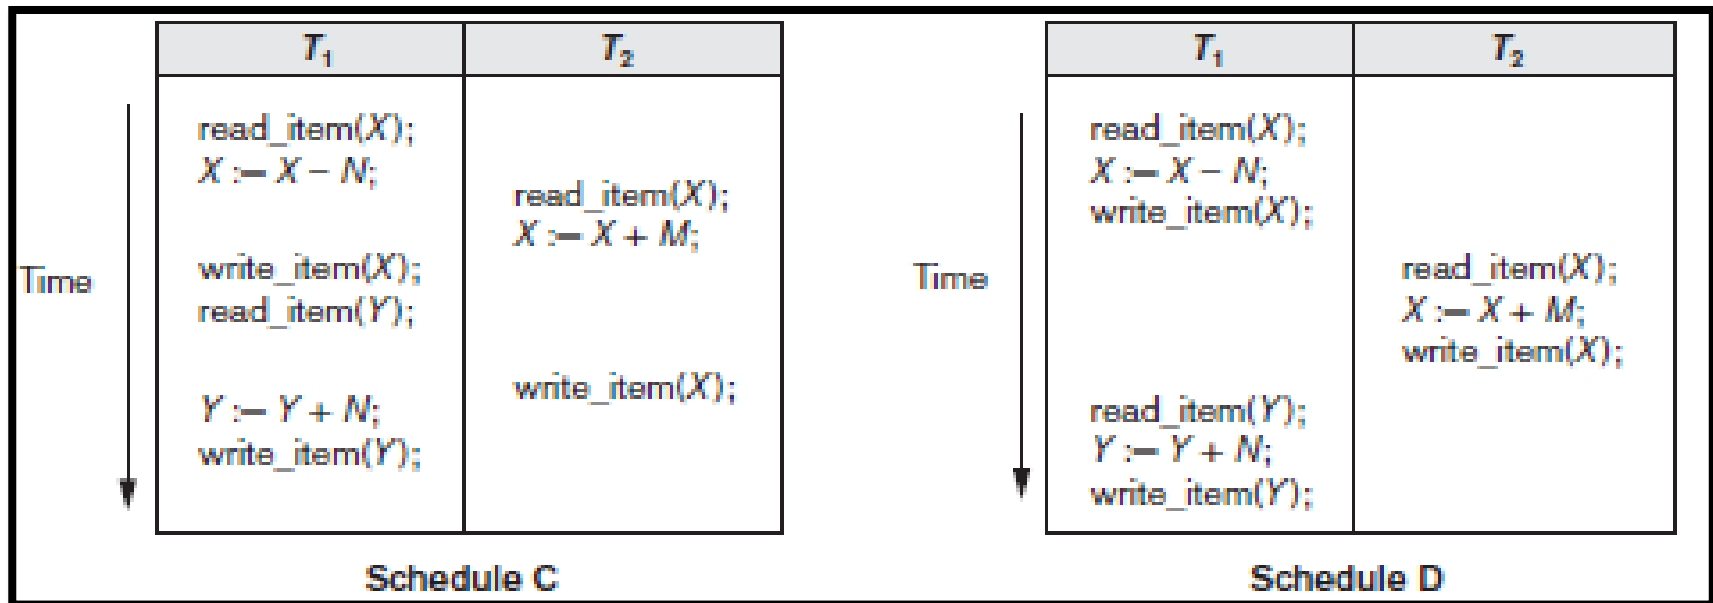
We would like to determine which of the non-serial schedules *always* give a correct result and which do not .



# Schedules

## Serializable schedule

- A schedule  $S$  is serializable if it is equivalent to some serial schedule of the same  $n$  transactions.



**A nonserial schedule  $S$  is serializable  
means it is correct**

# Schedules- conflict

- Two operations in  $S$  are in **conflict** if
  - they belong to **different transactions**
  - they **access the same item  $X$**  and
  - at least one** of the operations is a **write\_item( $X$ )**.

$S: r_1(x) \quad w_1(X) \quad r_2(X) \quad w_2(x) \quad r_1(y) \quad a_1$

## • Conflicting operations

- $r_1(X)$  and  $w_2(X)$
- $r_2(X)$  and  $w_1(X)$
- $w_1(X)$  and  $w_2(X)$

## • Non-conflicting operations

- $r_1(X)$  and  $r_2(X)$
- $w_2(X)$  and  $w_1(Y)$
- $r_1(X)$  and  $w_1(X)$

Intuitively, two operations are conflicting if changing their order can result in a different outcome.

Types of Conflict: read-write conflict and write-write conflict

# Characterizing Schedules based on Serializability

- **Conflict equivalent:** Two schedules are said to be conflict equivalent if the order of any two conflicting operations is same in both schedules.

Two operations on the same data items conflict if at least one of them is a write

- $r(X)$  and  $w(X)$
- $w(X)$  and  $r(X)$
- $w(X)$  and  $w(X)$

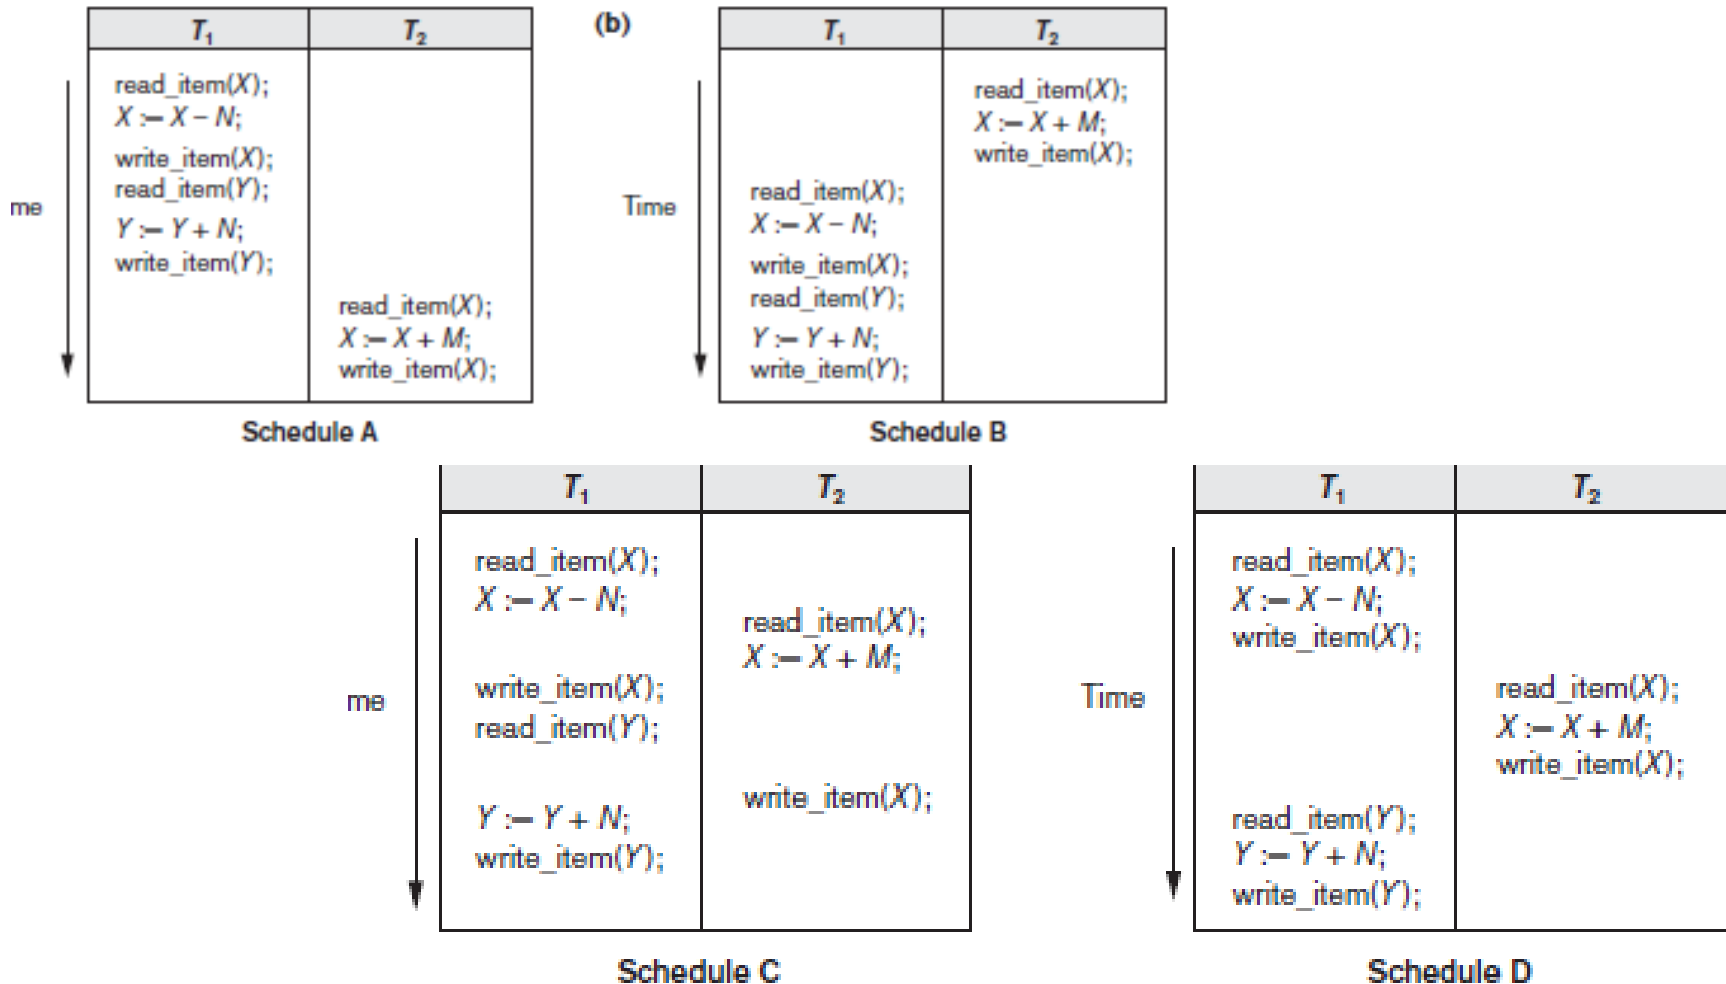
- If two conflicting operations are applied in *different orders* in two schedules, the effect can be different





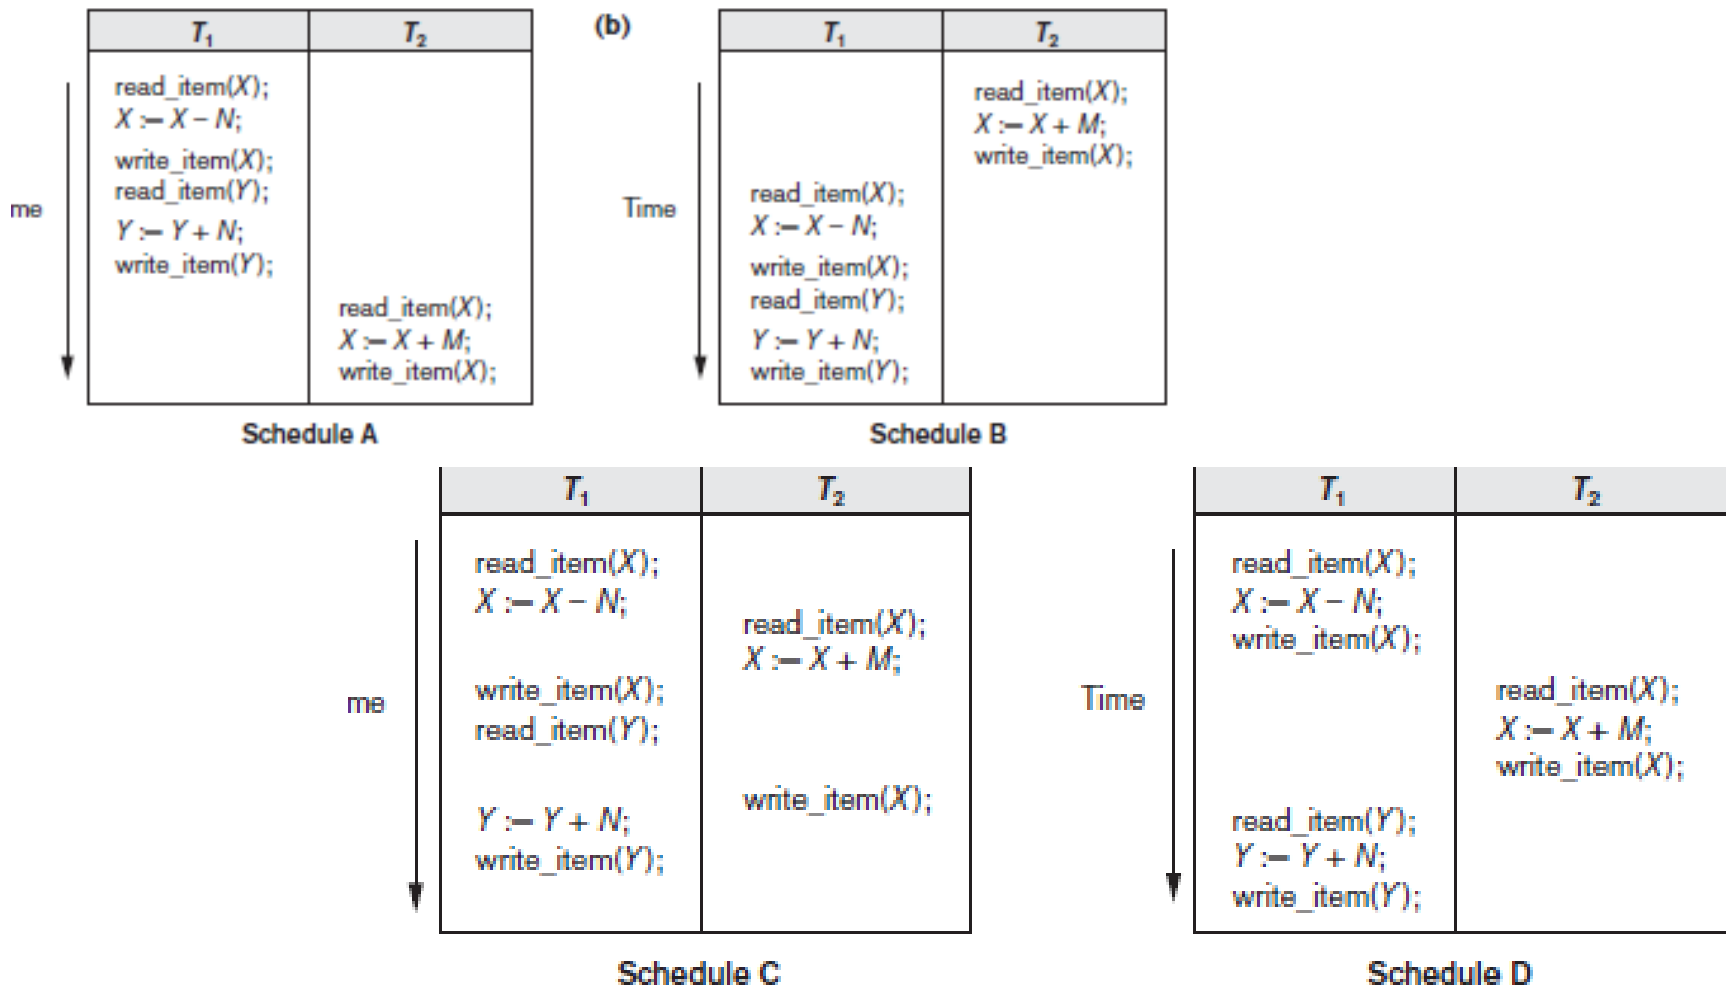
# Serializability

- A schedule S is said to be conflict serializable if it is conflict equivalent to some serial schedule S'.




# Serializability

- A schedule S is said to be conflict serializable if it is conflict equivalent to some serial schedule S'.




# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph**
  - a *directed graph* where the vertices **are** the transactions.
- Draw an arc from  $T_i$  to  $T_j$ 
  - if the two transactions conflict, and
  - $T_i$  accessed the data item on which the conflict arose **earlier**.
- We may label the arc by the **item** that was **accessed**.




# Precedence Graph

- A node for each transaction
- A directed edge from  $T_i$  to  $T_j$  if an operation of  $T_i$  precedes and conflicts with an operation of  $T_j$  in the schedule





# Precedence Graph

- A node for each transaction
- A directed edge from  $T_i$  to  $T_j$  if an operation of  $T_i$  precedes and conflicts with an operation of  $T_j$  in the schedule





# Precedence Graph

- A node for each transaction
- A directed edge from  $T_i$  to  $T_j$  if an operation of  $T_i$  precedes and conflicts with an operation of  $T_j$  in the schedule

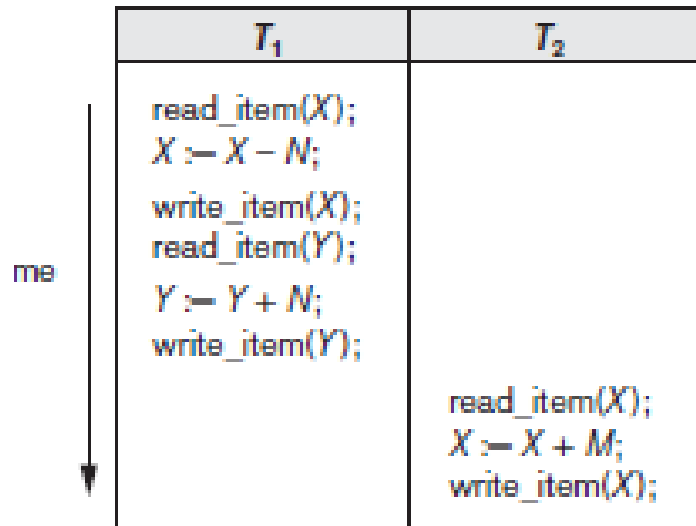


# Precedence Graph

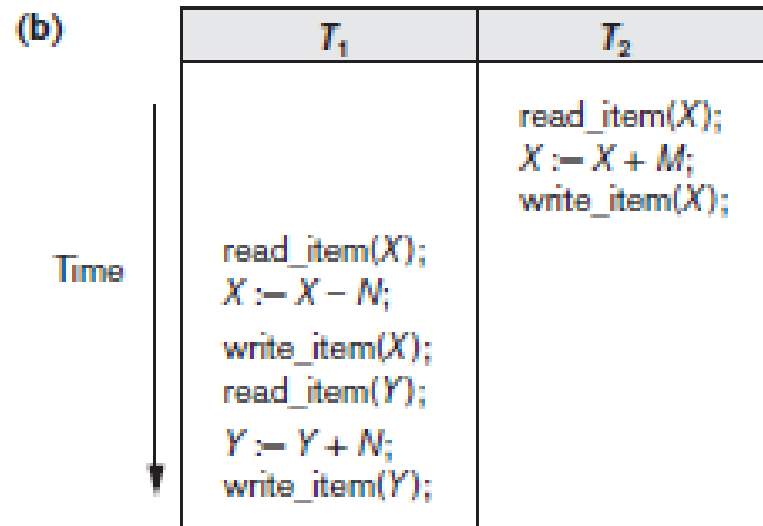
- A node for each transaction
- A directed edge from  $T_i$  to  $T_j$  if an operation of  $T_i$  precedes and conflicts with an operation of  $T_j$  in the schedule





# Constructing the Precedence Graphs



Schedule A

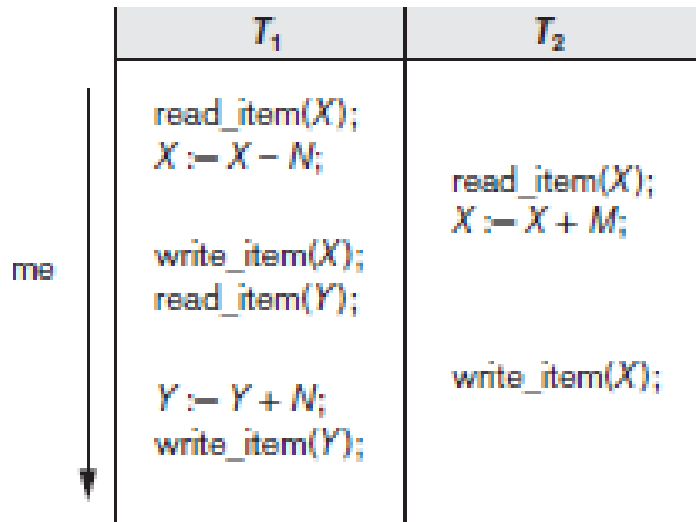


Schedule B

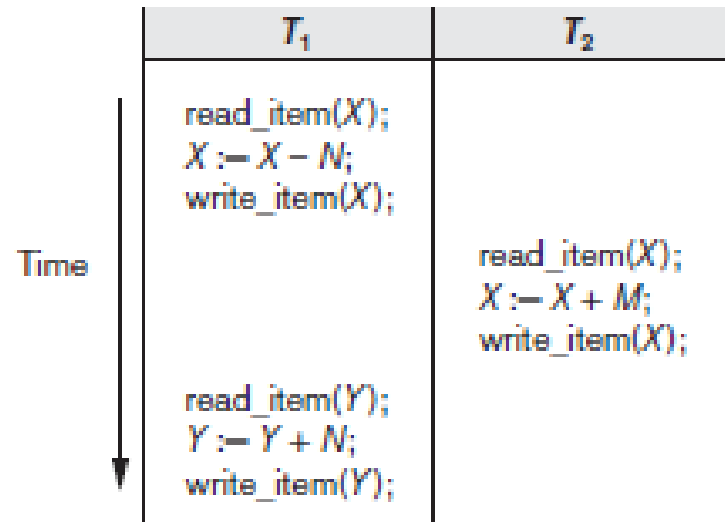






# Constructing the Precedence Graphs



Schedule C



Schedule D



# Example of serializability Testing

## Transaction $T_1$

```
read_item(X);  
write_item(X);  
read_item(Y);  
write_item(Y);
```

## Transaction $T_2$

```
read_item(Z);  
read_item(Y);  
write_item(Y);  
read_item(X);  
write_item(X);
```

## Transaction $T_3$

```
read_item(Y);  
read_item(Z);  
write_item(Y);  
write_item(Z);
```



# Example of Serializability Testing

Time ↓

| Transaction $T_1$               | Transaction $T_2$                                | Transaction $T_3$                |
|---------------------------------|--|----------------------------------|
| read_item(X);<br>write_item(X); | read_item(Z);<br>read_item(Y);<br>write_item(Y); | read_item(Y);<br>read_item(Z);   |
| read_item(Y);<br>write_item(Y); | read_item(X);<br><br>write_item(X);              | write_item(Y);<br>write_item(Z); |

**Schedule E**




# Example of Serializability Testing

Time ↓

| Transaction $T_1$               | Transaction $T_2$                                | Transaction $T_3$                |
|---------------------------------|--|----------------------------------|
| read_item(X);<br>write_item(X); | read_item(Z);<br>read_item(Y);<br>write_item(Y); | read_item(Y);<br>read_item(Z);   |
| read_item(Y);<br>write_item(Y); | read_item(X);<br><br>write_item(X);              | write_item(Y);<br>write_item(Z); |

**Schedule E**




# Example of Serializability Testing

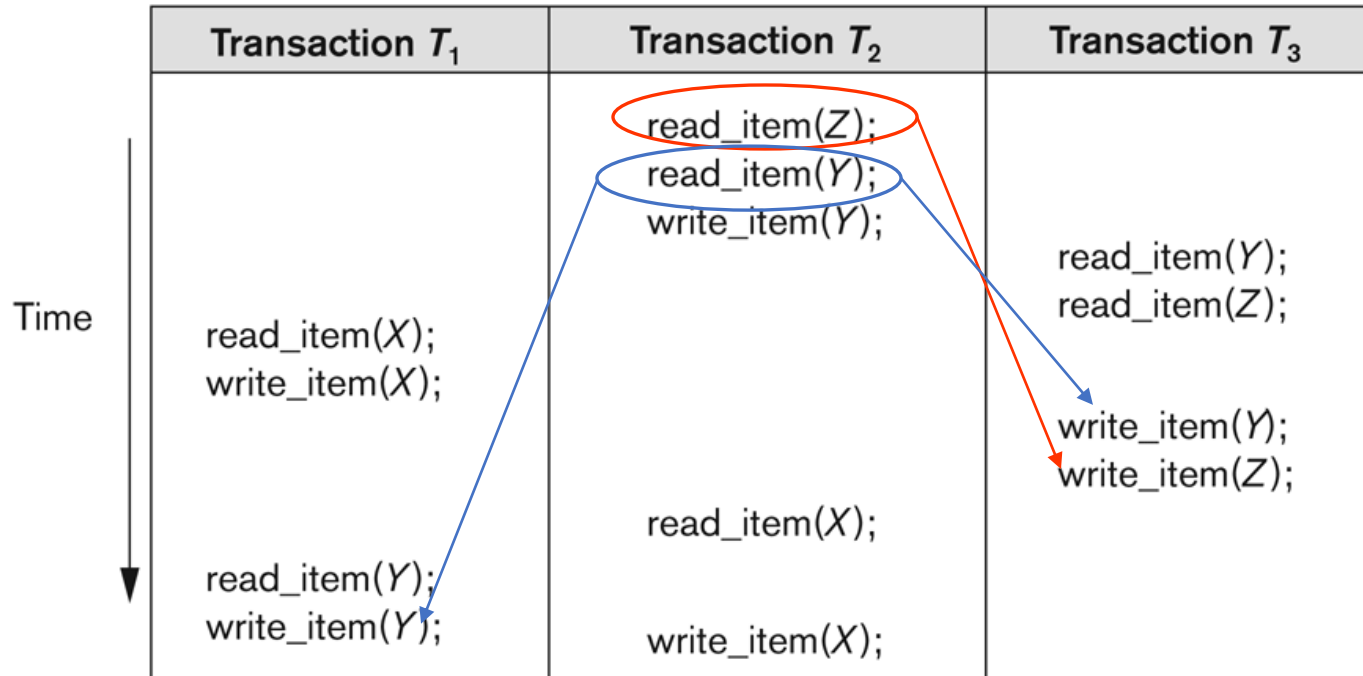
Time ↓

| Transaction $T_1$               | Transaction $T_2$                                | Transaction $T_3$                |
|---------------------------------|--|----------------------------------|
| read_item(X);<br>write_item(X); | read_item(Z);<br>read_item(Y);<br>write_item(Y); | read_item(Y);<br>read_item(Z);   |
| read_item(Y);<br>write_item(Y); | read_item(X);                                    | write_item(Y);<br>write_item(Z); |
|                                 | write_item(X);                                   |                                  |


Schedule E



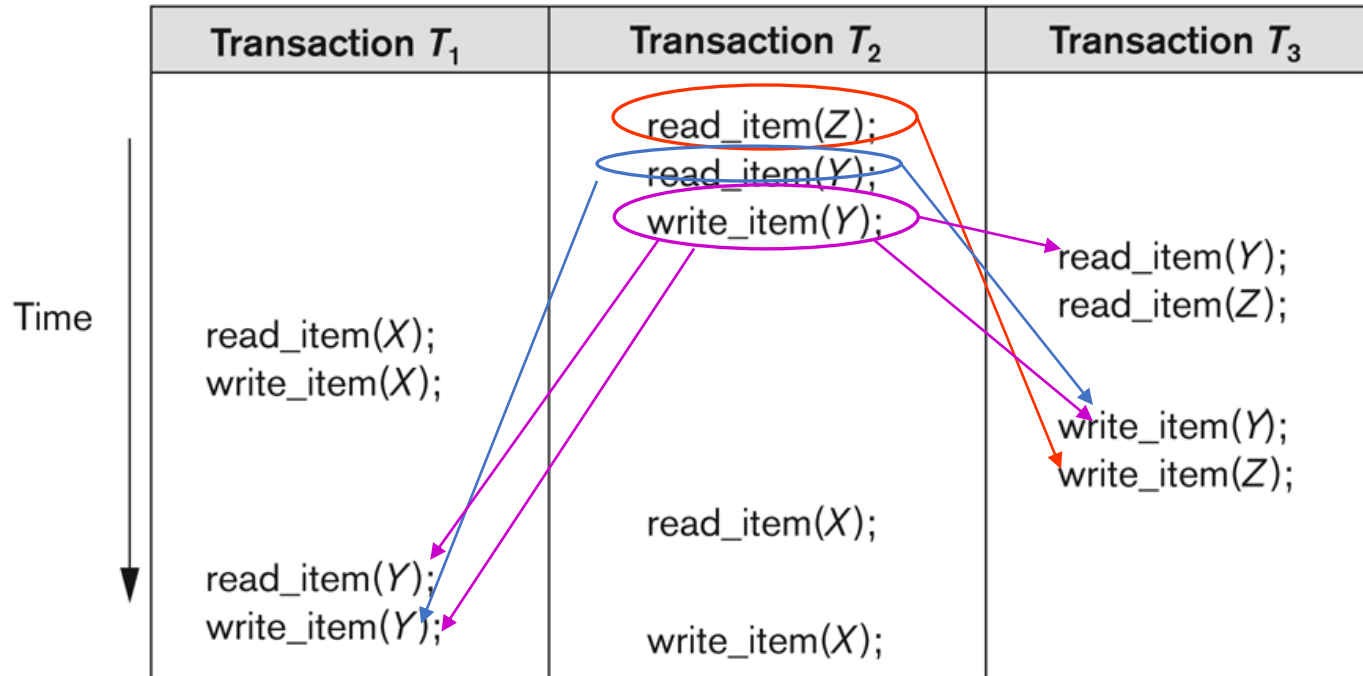
# Example of Serializability Testing




Schedule E



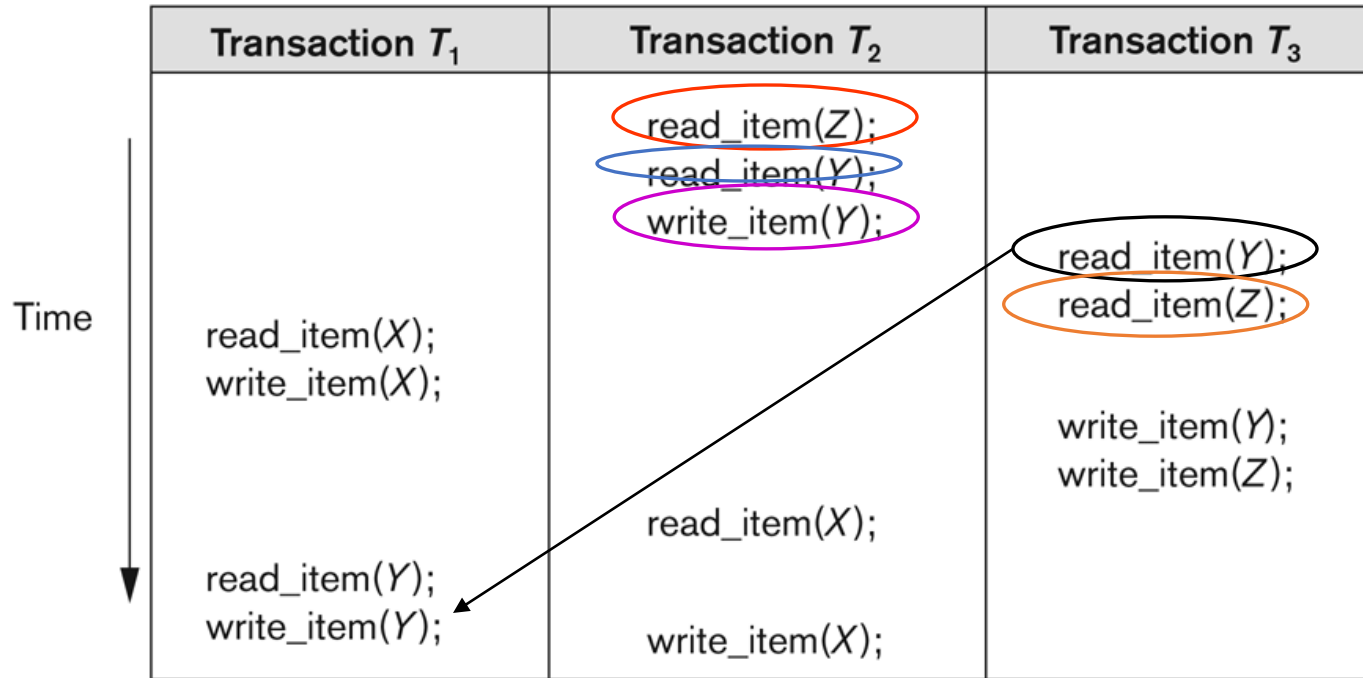
# Example of Serializability Testing




**Schedule E**



# Example of Serializability Testing

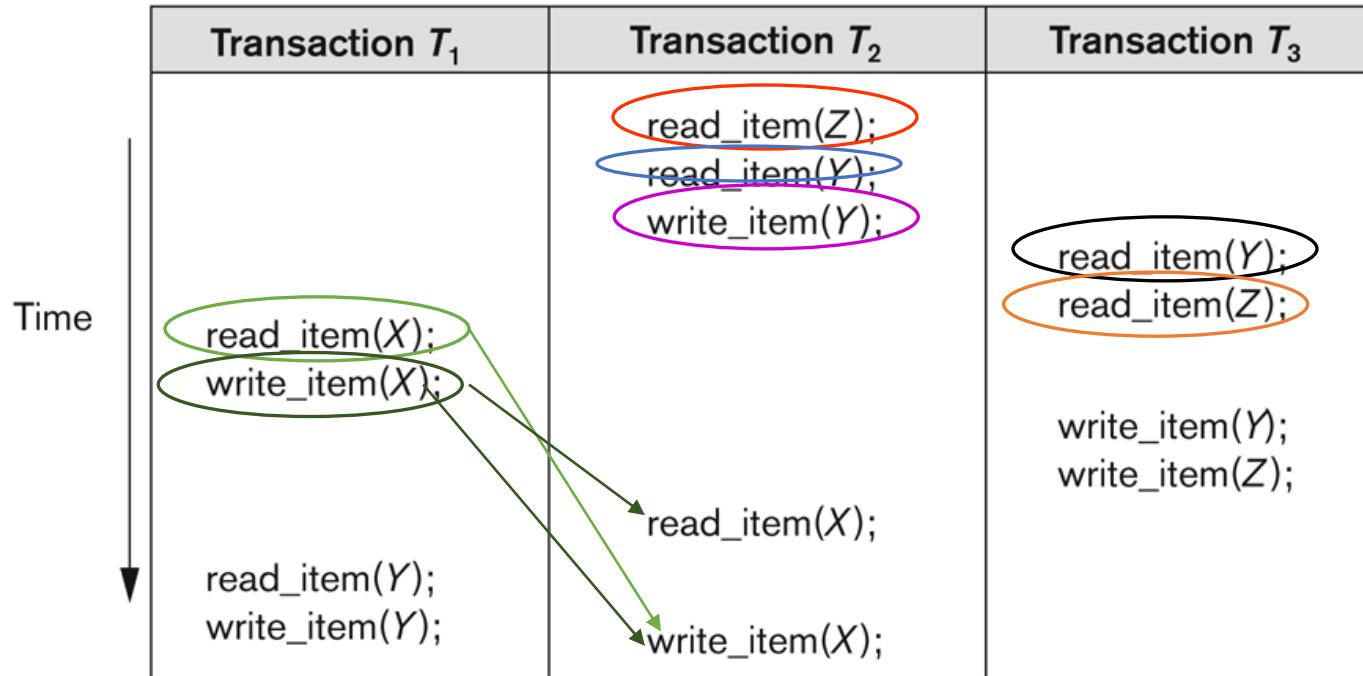


Schedule E






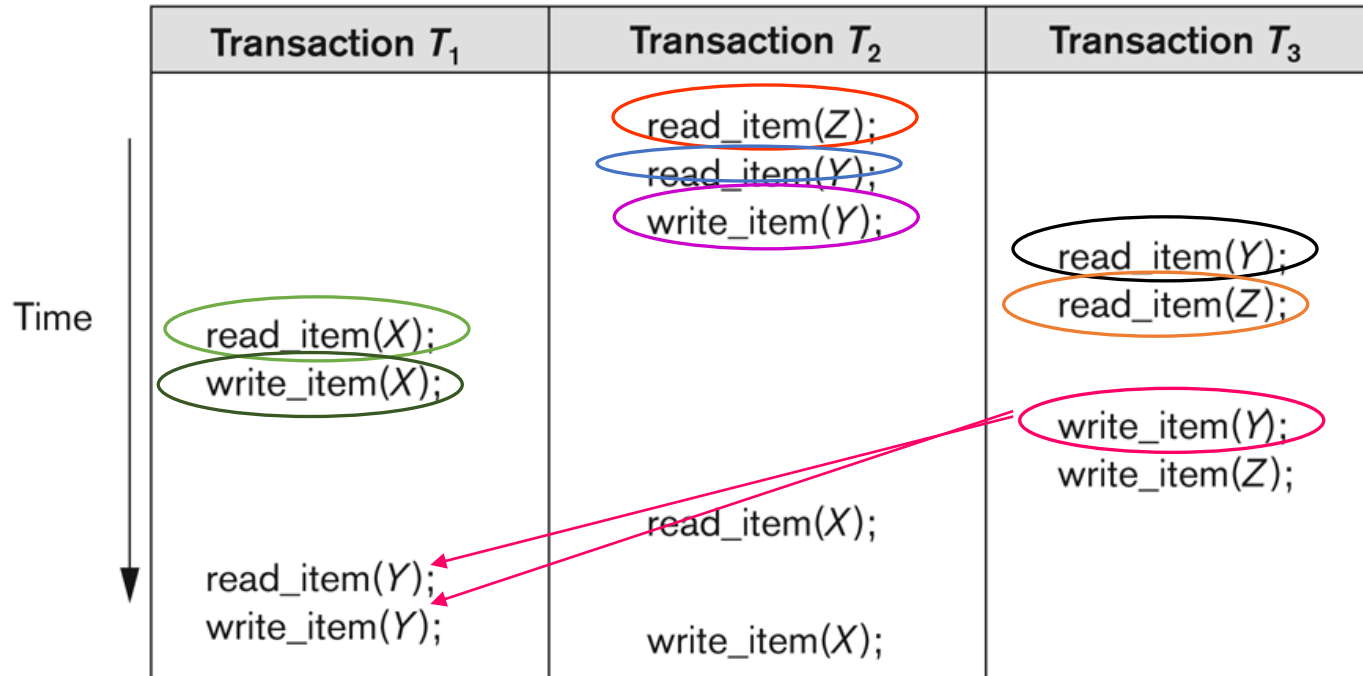
# Example of Serializability Testing




**Schedule E**



# Example of Serializability Testing



Schedule E




# Example of Serializability Testing

Time ↓

| Transaction $T_1$                                   | Transaction $T_2$  | Transaction $T_3$                                    |
|---|--|--|
| $\text{read\_item}(X);$<br>$\text{write\_item}(X);$ | $\text{read\_item}(Z);$<br>$\text{read\_item}(Y);$<br>$\text{write\_item}(Y);$ | $\text{read\_item}(Y);$<br>$\text{read\_item}(Z);$   |
| $\text{read\_item}(Y);$<br>$\text{write\_item}(Y);$ | $\text{read\_item}(X);$  | $\text{write\_item}(Y);$<br>$\text{write\_item}(Z);$ |
| $\text{read\_item}(Y);$<br>$\text{write\_item}(Y);$ | $\text{write\_item}(X);$   |  |


Schedule E



# Example of Serializability Testing

|           | Transaction $T_1$               | Transaction $T_2$                                | Transaction $T_3$                |
|-----------|---------------------------------|--|----------------------------------|
| Time<br>↓ | read_item(X);<br>write_item(X); | read_item(Z);<br>read_item(Y);<br>write_item(Y); | read_item(Y);<br>read_item(Z);   |
|           | read_item(Y);<br>write_item(Y); | read_item(X);<br><br>write_item(X);              | write_item(Y);<br>write_item(Z); |

**Schedule E**



**Equivalent serial schedules**

None

**Reason**


Cycle  $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

Cycle  $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

# Example of Serializability Testing

|           | Transaction $T_1$               | Transaction $T_2$   | Transaction $T_3$                |
|-----------|---------------------------------|---|----------------------------------|
| Time<br>↓ | read_item(X);<br>write_item(X); |   | read_item(Y);<br>read_item(Z);   |
|           | read_item(Y);<br>write_item(Y); | read_item(Z);<br><br>read_item(Y);<br>write_item(Y);<br>read_item(X);<br>write_item(X); | write_item(Y);<br>write_item(Z); |

Schedule F



Equivalent serial schedules


$T_3 \rightarrow T_1 \rightarrow T_2$

# Example of Serializability Testing

| Transaction $T_1$ |
|-------------------|
| read_item(X);     |
| write_item(X);    |
| read_item(Y);     |
| write_item(Y);    |

| Transaction $T_2$ |
|-------------------|
| read_item(Z);     |
| read_item(Y);     |
| write_item(Y);    |
| read_item(X);     |
| write_item(X);    |

| Transaction $T_3$ |
|-------------------|
| read_item(Y);     |
| read_item(Z);     |
| write_item(Y);    |
| write_item(Z);    |




To find an equivalent serial schedule:

- start with a node that does not have any incoming edges, and then
- make sure that the node order for every edge is not violated.

# Example of Serializability Testing

|           | Transaction $T_1$               | Transaction $T_2$   | Transaction $T_3$                |
|-----------|---------------------------------|---|----------------------------------|
| Time<br>↓ | read_item(X);<br>write_item(X); |   | read_item(Y);<br>read_item(Z);   |
|           | read_item(Y);<br>write_item(Y); | read_item(Z);<br><br>read_item(Y);<br>write_item(Y);<br>read_item(X);<br>write_item(X); | write_item(Y);<br>write_item(Z); |

Schedule F

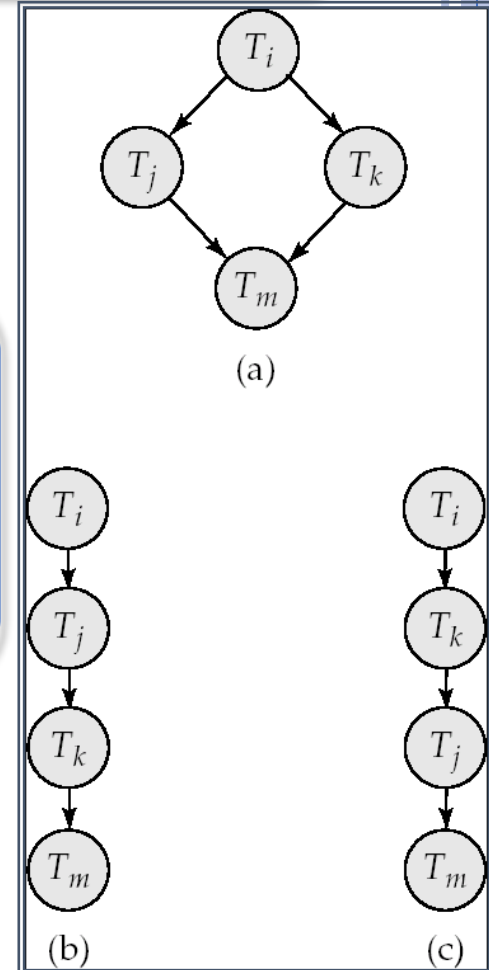


Equivalent serial schedules

$T_3 \rightarrow T_1 \rightarrow T_2$

# Test for Conflict Serializability

- A schedule is **conflict serializable**
  - if and only if its *precedence graph* is **acyclic**.
- **Cycle-detection algorithm:**
  - Use **Depth-first search** to detect cycle
    - **DFS** for a connected graph produces a tree.
    - There is a cycle in a graph only if there is a **back-edge** present in the graph
- If **precedence graph** is **acyclic**, the *serializability order* can be obtained by a **topological sorting** of the graph.



A **topological ordering** of a directed graph is a linear ordering of its vertices such that for every directed edge  $uv$  from vertex  $u$  to vertex  $v$ ,  $u$  comes before  $v$  in the ordering.



# Characterizing Schedules based on Serializability

- Serializability is **hard to check**.
  - Interleaving of operations occurs in an operating system through some scheduler
  - Difficult to determine beforehand how the operations in a schedule will be interleaved.
- Current approach used in most DBMSs:
  - Use of locks with **Two Phase locking Protocol**

